

Linux Device Drivers

This documentation is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

This program is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with this program; if not, write to the Free Software Foundation, Inc., 59 Temple Place, Suite 330, Boston, MA 02111-1307 USA

For more details see the file COPYING in the source distribution of Linux.

Table of Contents

[1. Driver Basics](#)

- [Driver Entry and Exit points](#)
- [Atomic and pointer manipulation](#)
- [Delaying, scheduling, and timer routines](#)
- [High-resolution timers](#)
- [Workqueues and Kevents](#)
- [Internal Functions](#)
- [Kernel objects manipulation](#)
- [Kernel utility functions](#)
- [Device Resource Management](#)

[2. Device drivers infrastructure](#)

- [Device Drivers Base](#)
- [Device Drivers Power Management](#)
- [Device Drivers ACPI Support](#)
- [Device drivers PnP support](#)
- [Userspace IO devices](#)

[3. Parallel Port Devices](#)

- [parport_yield](#) — relinquish a parallel port temporarily
- [parport_yield_blocking](#) — relinquish a parallel port temporarily
- [parport_wait_event](#) — wait for an event on a parallel port
- [parport_wait_peripheral](#) — wait for status lines to change in 35ms
- [parport_negotiate](#) — negotiate an IEEE 1284 mode
- [parport_write](#) — write a block of data to a parallel port
- [parport_read](#) — read a block of data from a parallel port
- [parport_set_timeout](#) — set the inactivity timeout for a device
- [parport_register_driver](#) — register a parallel port device driver

[parport_unregister_driver](#) — deregister a parallel port device driver
[parport_get_port](#) — increment a port's reference count
[parport_put_port](#) — decrement a port's reference count
[parport_register_port](#) — register a parallel port
[parport_announce_port](#) — tell device drivers about a parallel port
[parport_remove_port](#) — deregister a parallel port
[parport_register_device](#) — register a device on a parallel port
[parport_unregister_device](#) — deregister a device on a parallel port
[parport_find_number](#) — find a parallel port by number
[parport_find_base](#) — find a parallel port by base address
[parport_claim](#) — claim access to a parallel port device
[parport_claim_or_block](#) — claim access to a parallel port device
[parport_release](#) — give up access to a parallel port device
[parport_open](#) — find a device by canonical device number
[parport_close](#) — close a device opened with `parport_open`

[4. Message-based devices](#)

[Fusion message devices](#)

[I2O message devices](#)

[5. Sound Devices](#)

[snd_register_device](#) — Register the ALSA device file for the card
[snd_printk](#) — printk wrapper
[snd_printd](#) — debug printk
[snd_BUG](#) — give a BUG warning message and stack trace
[snd_BUG_ON](#) — debugging check macro
[snd_printdd](#) — debug printk
[register_sound_special_device](#) — register a special sound node
[register_sound_mixer](#) — register a mixer device
[register_sound_midi](#) — register a midi device
[register_sound_dsp](#) — register a DSP device
[unregister_sound_special](#) — unregister a special sound device
[unregister_sound_mixer](#) — unregister a mixer
[unregister_sound_midi](#) — unregister a midi device
[unregister_sound_dsp](#) — unregister a DSP device
[snd_pcm_playback_ready](#) — check whether the playback buffer is available
[snd_pcm_capture_ready](#) — check whether the capture buffer is available
[snd_pcm_playback_data](#) — check whether any data exists on the playback buffer
[snd_pcm_playback_empty](#) — check whether the playback buffer is empty
[snd_pcm_capture_empty](#) — check whether the capture buffer is empty
[snd_pcm_format_cpu_endian](#) — Check the PCM format is CPU-endian
[snd_pcm_new_stream](#) — create a new PCM stream
[snd_pcm_new](#) — create a new PCM instance
[snd_device_new](#) — create an ALSA device component
[snd_device_free](#) — release the device from the card
[snd_device_register](#) — register the device
[snd_iprintf](#) — printf on the procfs buffer
[snd_info_get_line](#) — read one line from the procfs buffer
[snd_info_get_str](#) — parse a string token

[snd_info_create_module_entry](#) — create an info entry for the given module
[snd_info_create_card_entry](#) — create an info entry for the given card
[snd_card_proc_new](#) — create an info entry for the given card
[snd_info_free_entry](#) — release the info entry
[snd_info_register](#) — register the info entry
[snd_rawmidi_receive](#) — receive the input data from the device
[snd_rawmidi_transmit_empty](#) — check whether the output buffer is empty
[snd_rawmidi_transmit_peek](#) — copy data from the internal buffer
[snd_rawmidi_transmit_ack](#) — acknowledge the transmission
[snd_rawmidi_transmit](#) — copy from the buffer to the device
[snd_rawmidi_new](#) — create a rawmidi instance
[snd_rawmidi_set_ops](#) — set the rawmidi operators
[snd_request_card](#) — try to load the card module
[snd_lookup_minor_data](#) — get user data of a registered device
[snd_register_device_for_dev](#) — Register the ALSA device file for the card
[snd_unregister_device](#) — unregister the device on the given card
[copy_to_user_fromio](#) — copy data from mmio-space to user-space
[copy_from_user_toio](#) — copy data from user-space to mmio-space
[snd_pcm_lib_preallocate_free_for_all](#) — release all pre-allocated buffers on the pcm
[snd_pcm_lib_preallocate_pages](#) — pre-allocation for the given DMA type
[snd_pcm_lib_preallocate_pages_for_all](#) — pre-allocation for continous memory type (all substreams)
[snd_pcm_sgbuf_ops_page](#) — get the page struct at the given offset
[snd_pcm_lib_malloc_pages](#) — allocate the DMA buffer
[snd_pcm_lib_free_pages](#) — release the allocated DMA buffer.
[snd_card_create](#) — create and initialize a soundcard structure
[snd_card_disconnect](#) — disconnect all APIs from the file-operations (user space)
[snd_card_set_id](#) — set card identification name
[snd_card_register](#) — register the soundcard
[snd_component_add](#) — add a component string
[snd_card_file_add](#) — add the file to the file list of the card
[snd_card_file_remove](#) — remove the file from the file list
[snd_power_wait](#) — wait until the power-state is changed.
[snd_dma_program](#) — program an ISA DMA transfer
[snd_dma_disable](#) — stop the ISA DMA transfer
[snd_dma_pointer](#) — return the current pointer to DMA transfer buffer in bytes
[snd_ctl_new1](#) — create a control instance from the template
[snd_ctl_free_one](#) — release the control instance
[snd_ctl_add](#) — add the control instance to the card
[snd_ctl_remove](#) — remove the control from the card and release it
[snd_ctl_remove_id](#) — remove the control of the given id and release it
[snd_ctl_rename_id](#) — replace the id of a control on the card
[snd_ctl_find_numid](#) — find the control instance with the given number-id
[snd_ctl_find_id](#) — find the control instance with the given id
[snd_pcm_set_ops](#) — set the PCM operators
[snd_pcm_set_sync](#) — set the PCM sync id
[snd_interval_refine](#) — refine the interval value of configurator
[snd_interval_ratnum](#) — refine the interval value
[snd_interval_list](#) — refine the interval value from the list
[snd_pcm_hw_rule_add](#) — add the hw-constraint rule
[snd_pcm_hw_constraint_integer](#) — apply an integer constraint to an interval

[snd_pcm_hw_constraint_minmax](#) — apply a min/max range constraint to an interval
[snd_pcm_hw_constraint_list](#) — apply a list of constraints to a parameter
[snd_pcm_hw_constraint_ratnums](#) — apply ratnums constraint to a parameter
[snd_pcm_hw_constraint_ratdens](#) — apply ratdens constraint to a parameter
[snd_pcm_hw_constraint_msbits](#) — add a hw constraint msbits rule
[snd_pcm_hw_constraint_step](#) — add a hw constraint step rule
[snd_pcm_hw_constraint_pow2](#) — add a hw constraint power-of-2 rule
[snd_pcm_hw_param_value](#) — return *params* field *var* value
[snd_pcm_hw_param_first](#) — refine config space and return minimum value
[snd_pcm_hw_param_last](#) — refine config space and return maximum value
[snd_pcm_lib_ioctl](#) — a generic PCM ioctl callback
[snd_pcm_period_elapsed](#) — update the pcm status for the next period
[snd_hwdep_new](#) — create a new hwdep instance
[snd_pcm_stop](#) — try to stop all running streams in the substream group
[snd_pcm_suspend](#) — trigger SUSPEND to all linked streams
[snd_pcm_suspend_all](#) — trigger SUSPEND to all substreams in the given pcm
[snd_malloc_pages](#) — allocate pages with the given size
[snd_free_pages](#) — release the pages
[snd_dma_alloc_pages](#) — allocate the buffer area according to the given type
[snd_dma_alloc_pages_fallback](#) — allocate the buffer area according to the given type with fallback
[snd_dma_free_pages](#) — release the allocated buffer
[snd_dma_get_reserved_buf](#) — get the reserved buffer for the given device
[snd_dma_reserve_buf](#) — reserve the buffer

[6. 16x50 UART Driver](#)

[uart_handle_dcd_change](#) — handle a change of carrier detect state
[uart_handle_cts_change](#) — handle a change of clear-to-send state
[uart_update_timeout](#) — update per-port FIFO timeout.
[uart_get_baud_rate](#) — return baud rate for a particular port
[uart_get_divisor](#) — return uart clock divisor
[uart_parse_options](#) — Parse serial port baud/parity/bits/flow contro.
[uart_set_options](#) — setup the serial console parameters
[uart_register_driver](#) — register a driver with the uart core layer
[uart_unregister_driver](#) — remove a driver from the uart core layer
[uart_add_one_port](#) — attach a driver-defined port structure
[uart_remove_one_port](#) — detach a driver defined port structure
[serial8250_suspend_port](#) — suspend one serial port
[serial8250_resume_port](#) — resume one serial port
[serial8250_register_port](#) — register a serial port
[serial8250_unregister_port](#) — remove a 16x50 serial port at runtime

[7. Frame Buffer Library](#)

[Frame Buffer Memory](#)
[Frame Buffer Colormap](#)
[Frame Buffer Video Mode Database](#)
[Frame Buffer Macintosh Video Mode Database](#)
[Frame Buffer Fonts](#)

8. Input Subsystem

[struct ff_replay](#) — defines scheduling of the force-feedback effect
[struct ff_trigger](#) — defines what triggers the force-feedback effect
[struct ff_envelope](#) — generic force-feedback effect envelope
[struct ff_constant_effect](#) — defines parameters of a constant force-feedback effect
[struct ff_ramp_effect](#) — defines parameters of a ramp force-feedback effect
[struct ff_condition_effect](#) — defines a spring or friction force-feedback effect
[struct ff_periodic_effect](#) — defines parameters of a periodic force-feedback effect
[struct ff_rumble_effect](#) — defines parameters of a periodic force-feedback effect
[struct ff_effect](#) — defines force feedback effect
[struct input_dev](#) — represents an input device
[struct input_handler](#) — implements one of interfaces for input devices
[struct input_handle](#) — links input device with an input handler
[struct ff_device](#) — force-feedback part of an input device
[input_event](#) — report new input event
[input_inject_event](#) — send input event from input handler
[input_grab_device](#) — grabs device for exclusive use
[input_release_device](#) — release previously grabbed device
[input_open_device](#) — open input device
[input_close_device](#) — close input device
[input_get_keycode](#) — retrieve keycode currently mapped to a given scancode
[input_set_keycode](#) — assign new keycode to a given scancode
[input_allocate_device](#) — allocate memory for new input device
[input_free_device](#) — free memory occupied by input_dev structure
[input_set_capability](#) — mark device as capable of a certain event
[input_register_device](#) — register device with input core
[input_unregister_device](#) — unregister previously registered device
[input_register_handler](#) — register a new input handler
[input_unregister_handler](#) — unregisters an input handler
[input_register_handle](#) — register a new input handle
[input_unregister_handle](#) — unregister an input handle
[input_ff_upload](#) — upload effect into force-feedback device
[input_ff_erase](#) — erase a force-feedback effect from device
[input_ff_event](#) — generic handler for force-feedback events
[input_ff_create](#) — create force-feedback device
[input_ff_destroy](#) — frees force feedback portion of input device
[input_ff_create_memless](#) — create memoryless force-feedback device

9. Serial Peripheral Interface (SPI)

[struct spi_device](#) — Master side proxy for an SPI slave device
[struct spi_driver](#) — Host side “protocol” driver
[spi_unregister_driver](#) — reverse effect of spi_register_driver
[struct spi_master](#) — interface to SPI master controller
[struct spi_transfer](#) — a read/write buffer pair
[struct spi_message](#) — one multi-segment SPI transaction
[spi_write](#) — SPI synchronous write
[spi_read](#) — SPI synchronous read
[spi_w8r8](#) — SPI synchronous 8 bit write followed by 8 bit read
[spi_w8r16](#) — SPI synchronous 8 bit write followed by 16 bit read

[struct spi_board_info](#) — board-specific template for a SPI device
[spi_register_board_info](#) — register SPI devices for a given board
[spi_register_driver](#) — register a SPI driver
[spi_alloc_device](#) — Allocate a new SPI device
[spi_add_device](#) — Add spi_device allocated with spi_alloc_device
[spi_new_device](#) — instantiate one new SPI device
[spi_alloc_master](#) — allocate SPI master controller
[spi_register_master](#) — register SPI master controller
[spi_unregister_master](#) — unregister SPI master controller
[spi_busnum_to_master](#) — look up master associated with bus_num
[spi_setup](#) — setup SPI mode and clock rate
[spi_async](#) — asynchronous SPI transfer
[spi_sync](#) — blocking/synchronous SPI data transfers
[spi_write_then_read](#) — SPI synchronous write followed by read

10. I²C and SMBus Subsystem

[struct i2c_driver](#) — represent an I2C device driver
[struct i2c_client](#) — represent an I2C slave device
[struct i2c_board_info](#) — template for device creation
[I2C_BOARD_INFO](#) — macro used to list an i2c device and its address
[struct i2c_msg](#) — an I2C transaction segment beginning with START
[i2c_register_board_info](#) — statically declare I2C devices
[i2c_verify_client](#) — return parameter as i2c_client, or NULL
[i2c_new_device](#) — instantiate an i2c device
[i2c_unregister_device](#) — reverse effect of i2c_new_device
[i2c_new_dummy](#) — return a new i2c device bound to a dummy driver
[i2c_add_adapter](#) — declare i2c adapter, use dynamic bus number
[i2c_add_numbered_adapter](#) — declare i2c adapter, use static bus number
[i2c_del_adapter](#) — unregister I2C adapter
[i2c_del_driver](#) — unregister I2C driver
[i2c_use_client](#) — increments the reference count of the i2c client structure
[i2c_release_client](#) — release a use of the i2c client structure
[i2c_transfer](#) — execute a single or combined I2C message
[i2c_master_send](#) — issue a single I2C message in master transmit mode
[i2c_master_recv](#) — issue a single I2C message in master receive mode
[i2c_smbus_read_byte](#) — SMBus “receive byte” protocol
[i2c_smbus_write_byte](#) — SMBus “send byte” protocol
[i2c_smbus_read_byte_data](#) — SMBus “read byte” protocol
[i2c_smbus_write_byte_data](#) — SMBus “write byte” protocol
[i2c_smbus_read_word_data](#) — SMBus “read word” protocol
[i2c_smbus_write_word_data](#) — SMBus “write word” protocol
[i2c_smbus_process_call](#) — SMBus “process call” protocol
[i2c_smbus_read_block_data](#) — SMBus “block read” protocol
[i2c_smbus_write_block_data](#) — SMBus “block write” protocol
[i2c_smbus_xfer](#) — execute SMBus protocol operations

Chapter 1. Driver Basics

Table of Contents

[Driver Entry and Exit points](#)
[Atomic and pointer manipulation](#)
[Delaying, scheduling, and timer routines](#)
[High-resolution timers](#)
[Workqueues and Kevents](#)
[Internal Functions](#)
[Kernel objects manipulation](#)
[Kernel utility functions](#)
[Device Resource Management](#)

Driver Entry and Exit points

Name

`module_init` — driver initialization entry point

Synopsis

```
module_init (x);  
  
x;
```

Arguments

x
function to be run at kernel boot time or module insertion

Description

`module_init` will either be called during `do_initcalls` (if builtin) or at module insertion time (if a module). There can only be one per module.

Name

`module_exit` — driver exit entry point

Synopsis

```
module_exit (x);  
  
x;
```

Arguments

x

function to be run when driver is removed

Description

`module_exit` will wrap the driver clean-up code with `cleanup_module` when used with `rmmod` when the driver is a module. If the driver is statically compiled into the kernel, `module_exit` has no effect. There can only be one per module.

Atomic and pointer manipulation

Name

`atomic_read` — read atomic variable

Synopsis

```
int atomic_read (v);

const atomic_t * v;
```

Arguments

`v`

pointer of type `atomic_t`

Description

Atomically reads the value of `v`.

Name

`atomic_set` — set atomic variable

Synopsis

```
void atomic_set (v,
                i);

atomic_t * v;
int      i;
```


Arguments

v

pointer of type `atomic_t`

i

required value

Description

Atomically sets the value of *v* to *i*.

Name

`atomic_add` — add integer to atomic variable

Synopsis

```
void atomic_add (i,  
                 v);
```

```
int          i;  
atomic_t * v;
```

Arguments

i

integer value to add

v

pointer of type `atomic_t`

Description

Atomically adds *i* to *v*.

Name

`atomic_sub` — subtract integer from atomic variable

Synopsis

```
void atomic_sub (i,  
                v);  
  
int      i;  
atomic_t * v;
```

Arguments

i
integer value to subtract

v
pointer of type `atomic_t`

Description

Atomically subtracts *i* from *v*.

Name

`atomic_sub_and_test` — subtract value from variable and test result

Synopsis

```
int atomic_sub_and_test (i,  
                        v);  
  
int      i;  
atomic_t * v;
```

Arguments

i
integer value to subtract

v
pointer of type `atomic_t`

Description

Atomically subtracts *i* from *v* and returns true if the result is zero, or false for all other cases.

Name

`atomic_inc` — increment atomic variable

Synopsis

```
void atomic_inc (v);
```

```
atomic_t * v;
```

Arguments

`v`

pointer of type `atomic_t`

Description

Atomically increments `v` by 1.

Name

`atomic_dec` — decrement atomic variable

Synopsis

```
void atomic_dec (v);
```

```
atomic_t * v;
```

Arguments

`v`

pointer of type `atomic_t`

Description

Atomically decrements `v` by 1.

Name

`atomic_dec_and_test` — decrement and test

Synopsis

```
int atomic_dec_and_test (v);
```

```
atomic_t * v;
```

Arguments

v

pointer of type `atomic_t`

Description

Atomically decrements *v* by 1 and returns true if the result is 0, or false for all other cases.

Name

`atomic_inc_and_test` — increment and test

Synopsis

```
int atomic_inc_and_test (v);
```

```
atomic_t * v;
```

Arguments

v

pointer of type `atomic_t`

Description

Atomically increments *v* by 1 and returns true if the result is zero, or false for all other cases.

Name

`atomic_add_negative` — add and test if negative

Synopsis

```
int atomic_add_negative (i,
```

```
v);
```

```
int      i;
atomic_t * v;
```

Arguments

i

integer value to add

v

pointer of type `atomic_t`

Description

Atomically adds *i* to *v* and returns true if the result is negative, or false when result is greater than or equal to zero.

Name

`atomic_add_return` — add integer and return

Synopsis

```
int atomic_add_return (i,
                      v);
```

```
int      i;
atomic_t * v;
```

Arguments

i

integer value to add

v

pointer of type `atomic_t`

Description

Atomically adds *i* to *v* and returns *i + v*

Name

`atomic_sub_return` — subtract integer and return

Synopsis

```
int atomic_sub_return (i,  
                       v);
```

```
int          i;  
atomic_t *   v;
```

Arguments

i

integer value to subtract

v

pointer of type `atomic_t`

Description

Atomically subtracts *i* from *v* and returns *v* - *i*

Name

`atomic_add_unless` — add unless the number is already a given value

Synopsis

```
int atomic_add_unless (v,  
                       a,  
                       u);
```

```
atomic_t *   v;  
int          a;  
int          u;
```

Arguments

v

pointer of type `atomic_t`

*a*the amount to add to *v*...*u*...unless *v* is equal to *u*.

Description

Atomically adds *a* to *v*, so long as *v* was not already *u*. Returns non-zero if *v* was not *u*, and zero otherwise.

Name

`atomic64_xchg` — xchg atomic64 variable

Synopsis

```
u64 atomic64_xchg (ptr,
                  new_val);
```

```
atomic64_t * ptr;
u64         new_val;
```

Arguments

*ptr*pointer to type `atomic64_t`*new_val*

value to assign

Description

Atomically xchgs the value of *ptr* to *new_val* and returns the old value.

Name

`atomic64_set` — set atomic64 variable

Synopsis


```
void atomic64_set (ptr,  
                  new_val);  
  
atomic64_t * ptr;  
u64         new_val;
```

Arguments

ptr

pointer to type `atomic64_t`

new_val

value to assign

Description

Atomically sets the value of *ptr* to *new_val*.

Name

`atomic64_read` — read `atomic64` variable

Synopsis

```
u64 atomic64_read (ptr);  
  
atomic64_t * ptr;
```

Arguments

ptr

pointer to type `atomic64_t`

Description

Atomically reads the value of *ptr* and returns it.

Name

`atomic64_add_return` — add and return

Synopsis

```
u64 atomic64_add_return (delta,  
                          ptr);
```

```
u64          delta;  
atomic64_t * ptr;
```

Arguments

delta

integer value to add

ptr

pointer to type atomic64_t

Description

Atomically adds *delta* to *ptr* and returns $delta + *ptr$

Name

atomic64_add — add integer to atomic64 variable

Synopsis

```
void atomic64_add (delta,  
                  ptr);
```

```
u64          delta;  
atomic64_t * ptr;
```

Arguments

delta

integer value to add

ptr

pointer to type atomic64_t

Description

Atomically adds *delta* to *ptr*.

Name

`atomic64_sub` — subtract the `atomic64` variable

Synopsis

```
void atomic64_sub (delta,  
                  ptr);
```

```
u64          delta;  
atomic64_t * ptr;
```

Arguments

delta

integer value to subtract

ptr

pointer to type `atomic64_t`

Description

Atomically subtracts *delta* from *ptr*.

Name

`atomic64_sub_and_test` — subtract value from variable and test result

Synopsis

```
int atomic64_sub_and_test (delta,  
                           ptr);
```

```
u64          delta;  
atomic64_t * ptr;
```

Arguments

delta

integer value to subtract

ptr

pointer to type `atomic64_t`

Description

Atomically subtracts *delta* from *ptr* and returns true if the result is zero, or false for all other cases.

Name

`atomic64_inc` — increment `atomic64` variable

Synopsis

```
void atomic64_inc (ptr);
```

```
atomic64_t * ptr;
```

Arguments

ptr

pointer to type `atomic64_t`

Description

Atomically increments *ptr* by 1.

Name

`atomic64_dec` — decrement `atomic64` variable

Synopsis

```
void atomic64_dec (ptr);
```

```
atomic64_t * ptr;
```

Arguments

ptr

pointer to type `atomic64_t`

Description

Atomically decrements *ptr* by 1.

Name

`atomic64_dec_and_test` — decrement and test

Synopsis

```
int atomic64_dec_and_test (ptr);
```

```
atomic64_t * ptr;
```

Arguments

ptr

pointer to type `atomic64_t`

Description

Atomically decrements *ptr* by 1 and returns true if the result is 0, or false for all other cases.

Name

`atomic64_inc_and_test` — increment and test

Synopsis

```
int atomic64_inc_and_test (ptr);
```

```
atomic64_t * ptr;
```

Arguments

ptr

pointer to type `atomic64_t`

Description

Atomically increments *ptr* by 1 and returns true if the result is zero, or false for all other cases.

Name

atomic64_add_negative — add and test if negative

Synopsis

```
int atomic64_add_negative (delta,  
                           ptr);
```

```
u64          delta;  
atomic64_t * ptr;
```

Arguments

delta

integer value to add

ptr

pointer to type atomic64_t

Description

Atomically adds *delta* to *ptr* and returns true if the result is negative, or false when result is greater than or equal to zero.

Name

/home/landley/linux/temp//arch/x86/include/asm/unaligned.h — Document generation inconsistency

Oops

Warning

The template for this document tried to insert the structured comment from the file /home/landley/linux/temp//arch/x86/include/asm/unaligned.h at this point, but none was found. This dummy section is inserted to allow generation to continue.

Delaying, scheduling, and timer routines

Name

struct task_cputime — collected CPU time counts

Synopsis

```
struct task_cputime {
    cputime_t utime;
    cputime_t stime;
    unsigned long long sum_exec_runtime;
};
```

Members

utime

time spent in user mode, in cputime_t units

stime

time spent in kernel mode, in cputime_t units

sum_exec_runtime

total time spent on the CPU, in nanoseconds

Description

This structure groups together three kinds of CPU time that are tracked for threads and thread groups. Most things considering CPU time want to group these counts together and treat all three of them in parallel.

Name

struct thread_group_cputimer — thread group interval timer counts

Synopsis

```
struct thread_group_cputimer {
    struct task_cputime cputime;
    int running;
    spinlock_t lock;
};
```

Members

cputime

thread group interval timers.

running

non-zero when there are timers running and *cputime* receives updates.

lock

lock for fields in this struct.

Description

This structure contains the version of `task_cputime`, above, that is used for thread group CPU timer calculations.

Name

`pid_alive` — check that a task structure is not stale

Synopsis

```
int pid_alive (p);  
  
struct task_struct * p;
```

Arguments

p

Task structure to be checked.

Description

Test if a process is not yet dead (at most zombie state) If `pid_alive` fails, then pointers within the task structure can be stale and must not be dereferenced.

Name

`is_global_init` — check if a task structure is init

Synopsis

```
int is_global_init (tsk);  
  
struct task_struct * tsk;
```

Arguments

tsk

Task structure to be checked.

Description

Check if a task structure is the first user space task the kernel created.

Name

`wake_up_process` — Wake up a specific process

Synopsis

```
int wake_up_process (p);  
  
struct task_struct * p;
```

Arguments

p

The process to be woken up.

Description

Attempt to wake up the nominated process and move it to the set of runnable processes. Returns 1 if the process was woken up, 0 if it was already running.

It may be assumed that this function implies a write memory barrier before changing the task state if and only if any tasks are woken up.

Name

`preempt_notifier_register` — tell me when current is being preempted & rescheduled

Synopsis

```
void preempt_notifier_register (notifier);  
  
struct preempt_notifier * notifier;
```

Arguments

notifier

notifier struct to register

Name

preempt_notifier_unregister — no longer interested in preemption notifications

Synopsis

```
void preempt_notifier_unregister (notifier);
```

```
struct preempt_notifier * notifier;
```

Arguments

notifier

notifier struct to unregister

Description

This is safe to call from within a preemption notifier.

Name

__wake_up — wake up threads blocked on a waitqueue.

Synopsis

```
void __wake_up (q,
               mode,
               nr_exclusive,
               key);
```

```
wait_queue_head_t * q;
unsigned int        mode;
int                 nr_exclusive;
void *              key;
```

Arguments

q

the waitqueue

mode

which threads

nr_exclusive

how many wake-one or wake-many threads to wake up

key

is directly passed to the wakeup function

Description

It may be assumed that this function implies a write memory barrier before changing the task state if and only if any tasks are woken up.

Name

`__wake_up_sync_key` — wake up threads blocked on a waitqueue.

Synopsis

```
void __wake_up_sync_key (q,  
                        mode,  
                        nr_exclusive,  
                        key);
```

```
wait_queue_head_t * q;  
unsigned int      mode;  
int              nr_exclusive;  
void *           key;
```

Arguments

q

the waitqueue

mode

which threads

nr_exclusive

how many wake-one or wake-many threads to wake up

key

opaque value to be passed to wakeup targets

Description

The sync wakeup differs that the waker knows that it will schedule away soon, so while the target thread will be woken up, it will not be migrated to another CPU - ie. the two threads are 'synchronized' with each other. This can prevent needless bouncing between CPUs.

On UP it can prevent extra preemption.

It may be assumed that this function implies a write memory barrier before changing the task state if and only if any tasks are woken up.

Name

`complete` — signals a single thread waiting on this completion

Synopsis

```
void complete (x);  
  
struct completion * x;
```

Arguments

`x`

holds the state of this particular completion

Description

This will wake up a single thread waiting on this completion. Threads will be awakened in the same order in which they were queued.

See also `complete_all`, `wait_for_completion` and related routines.

It may be assumed that this function implies a write memory barrier before changing the task state if and only if any tasks are woken up.

Name

`complete_all` — signals all threads waiting on this completion

Synopsis

```
void complete_all (x);

struct completion * x;
```

Arguments

x

holds the state of this particular completion

Description

This will wake up all threads waiting on this particular completion event.

It may be assumed that this function implies a write memory barrier before changing the task state if and only if any tasks are woken up.

Name

`wait_for_completion` — waits for completion of a task

Synopsis

```
void __sched wait_for_completion (x);

struct completion * x;
```

Arguments

x

holds the state of this particular completion

Description

This waits to be signaled for completion of a specific task. It is NOT interruptible and there is no timeout.

See also similar routines (i.e. `wait_for_completion_timeout`) with timeout and interrupt capability. Also see `complete`.

Name

`wait_for_completion_timeout` — waits for completion of a task (w/timeout)

Synopsis

```
unsigned long __sched wait_for_completion_timeout (x,
                                                    timeout);

struct completion * x;
unsigned long      timeout;
```

Arguments

x

holds the state of this particular completion

timeout

timeout value in jiffies

Description

This waits for either a completion of a specific task to be signaled or for a specified timeout to expire. The timeout is in jiffies. It is not interruptible.

Name

`wait_for_completion_interruptible` — waits for completion of a task (w/intr)

Synopsis

```
int __sched wait_for_completion_interruptible (x);

struct completion * x;
```

Arguments

x

holds the state of this particular completion

Description

This waits for completion of a specific task to be signaled. It is interruptible.

Name

`wait_for_completion_interruptible_timeout` — waits for completion (w/(to,intr))

Synopsis

```
unsigned long __sched wait_for_completion_interruptible_timeout (x,
                                                                timeout);
```

```
struct completion * x;
unsigned long      timeout;
```

Arguments

x

holds the state of this particular completion

timeout

timeout value in jiffies

Description

This waits for either a completion of a specific task to be signaled or for a specified timeout to expire. It is interruptible. The timeout is in jiffies.

Name

`wait_for_completion_killable` — waits for completion of a task (killable)

Synopsis

```
int __sched wait_for_completion_killable (x);
```

```
struct completion * x;
```

Arguments

x

holds the state of this particular completion

Description

This waits to be signaled for completion of a specific task. It can be interrupted by a kill signal.

Name

`try_wait_for_completion` — try to decrement a completion without blocking

Synopsis

```
bool try_wait_for_completion (x);
```

```
struct completion * x;
```

Arguments

`x`

completion structure

Returns

0 if a decrement cannot be done without blocking 1 if a decrement succeeded.

If a completion is being used as a counting completion, attempt to decrement the counter without blocking. This enables us to avoid waiting if the resource the completion is protecting is not available.

Name

`completion_done` — Test to see if a completion has any waiters

Synopsis

```
bool completion_done (x);
```

```
struct completion * x;
```

Arguments

`x`

completion structure

Returns

0 if there are waiters (`wait_for_completion` in progress) 1 if there are no waiters.

Name

`task_nice` — return the nice value of a given task.

Synopsis

```
int task_nice (p);

const struct task_struct * p;
```

Arguments

p
the task in question.

Name

`sched_setscheduler` — change the scheduling policy and/or RT priority of a thread.

Synopsis

```
int sched_setscheduler (p,  
                        policy,  
                        param);

struct task_struct * p;  
int policy;  
struct sched_param * param;
```

Arguments

p
the task in question.

policy
new policy.

param
structure containing the new RT priority.

Description

NOTE that the task may be already dead.

Name

`yield` — yield the current processor to other threads.

Synopsis

```
void __sched yield (void);  
  
void;
```

Arguments

```
void  
  
no arguments
```

Description

This is a shortcut for kernel-space yielding - it marks the thread runnable and calls `sys_sched_yield`.

Name

`__round_jiffies` — function to round jiffies to a full second

Synopsis

```
unsigned long __round_jiffies (j,  
                                cpu);  
  
unsigned long j;  
int           cpu;
```

Arguments

```
j  
  
the time in (absolute) jiffies that should be rounded  
  
cpu  
  
the processor number on which the timeout will happen
```

Description

`__round_jiffies` rounds an absolute time in the future (in jiffies) up or down to (approximately) full

seconds. This is useful for timers for which the exact time they fire does not matter too much, as long as they fire approximately every X seconds.

By rounding these timers to whole seconds, all such timers will fire at the same time, rather than at various times spread out. The goal of this is to have the CPU wake up less, which saves power.

The exact rounding is skewed for each processor to avoid all processors firing at the exact same time, which could lead to lock contention or spurious cache line bouncing.

The return value is the rounded version of the *j* parameter.

Name

`__round_jiffies_relative` — function to round jiffies to a full second

Synopsis

```
unsigned long __round_jiffies_relative (j,  
                                         cpu);
```

```
unsigned long j;  
int          cpu;
```

Arguments

j
the time in (relative) jiffies that should be rounded

cpu
the processor number on which the timeout will happen

Description

`__round_jiffies_relative` rounds a time delta in the future (in jiffies) up or down to (approximately) full seconds. This is useful for timers for which the exact time they fire does not matter too much, as long as they fire approximately every X seconds.

By rounding these timers to whole seconds, all such timers will fire at the same time, rather than at various times spread out. The goal of this is to have the CPU wake up less, which saves power.

The exact rounding is skewed for each processor to avoid all processors firing at the exact same time, which could lead to lock contention or spurious cache line bouncing.

The return value is the rounded version of the *j* parameter.

Name

`round_jiffies` — function to round jiffies to a full second

Synopsis

```
unsigned long round_jiffies (j);
```

```
unsigned long j;
```

Arguments

j
the time in (absolute) jiffies that should be rounded

Description

`round_jiffies` rounds an absolute time in the future (in jiffies) up or down to (approximately) full seconds. This is useful for timers for which the exact time they fire does not matter too much, as long as they fire approximately every X seconds.

By rounding these timers to whole seconds, all such timers will fire at the same time, rather than at various times spread out. The goal of this is to have the CPU wake up less, which saves power.

The return value is the rounded version of the *j* parameter.

Name

`round_jiffies_relative` — function to round jiffies to a full second

Synopsis

```
unsigned long round_jiffies_relative (j);
```

```
unsigned long j;
```

Arguments

j
the time in (relative) jiffies that should be rounded

Description

`round_jiffies_relative` rounds a time delta in the future (in jiffies) up or down to (approximately) full seconds. This is useful for timers for which the exact time they fire does not matter too much, as long as they fire approximately every X seconds.

By rounding these timers to whole seconds, all such timers will fire at the same time, rather than at various times spread out. The goal of this is to have the CPU wake up less, which saves power.

The return value is the rounded version of the *j* parameter.

Name

`__round_jiffies_up` — function to round jiffies up to a full second

Synopsis

```
unsigned long __round_jiffies_up (j,  
                                cpu);
```

```
unsigned long j;  
int          cpu;
```

Arguments

j
the time in (absolute) jiffies that should be rounded

cpu
the processor number on which the timeout will happen

Description

This is the same as `__round_jiffies` except that it will never round down. This is useful for timeouts for which the exact time of firing does not matter too much, as long as they don't fire too early.

Name

`__round_jiffies_up_relative` — function to round jiffies up to a full second

Synopsis

```
unsigned long __round_jiffies_up_relative (j,  
                                          cpu);
```

```
unsigned long j;
```



```
int          cpu;
```

Arguments

j

the time in (relative) jiffies that should be rounded

cpu

the processor number on which the timeout will happen

Description

This is the same as `__round_jiffies_relative` except that it will never round down. This is useful for timeouts for which the exact time of firing does not matter too much, as long as they don't fire too early.

Name

`round_jiffies_up` — function to round jiffies up to a full second

Synopsis

```
unsigned long round_jiffies_up (j);
```

```
unsigned long j;
```

Arguments

j

the time in (absolute) jiffies that should be rounded

Description

This is the same as `round_jiffies` except that it will never round down. This is useful for timeouts for which the exact time of firing does not matter too much, as long as they don't fire too early.

Name

`round_jiffies_up_relative` — function to round jiffies up to a full second

Synopsis

```
unsigned long round_jiffies_up_relative (j);
```

```
unsigned long j;
```

Arguments

j

the time in (relative) jiffies that should be rounded

Description

This is the same as `round_jiffies_relative` except that it will never round down. This is useful for timeouts for which the exact time of firing does not matter too much, as long as they don't fire too early.

Name

`init_timer_key` — initialize a timer

Synopsis

```
void init_timer_key (timer,  
                     name,  
                     key);  
  
struct timer_list *      timer;  
const char *             name;  
struct lock_class_key *  key;
```

Arguments

timer

the timer to be initialized

name

name of the timer

key

lockdep class key of the fake lock used for tracking timer sync lock dependencies

Description

`init_timer_key` must be done to a timer prior calling *any* of the other timer functions.

Name

`mod_timer_pending` — modify a pending timer's timeout

Synopsis

```
int mod_timer_pending (timer,  
                        expires);
```

```
struct timer_list * timer;  
unsigned long      expires;
```

Arguments

timer

the pending timer to be modified

expires

new timeout in jiffies

Description

`mod_timer_pending` is the same for pending timers as `mod_timer`, but will not re-activate and modify already deleted timers.

It is useful for unserialized use of timers.

Name

`mod_timer` — modify a timer's timeout

Synopsis

```
int mod_timer (timer,  
               expires);
```

```
struct timer_list * timer;  
unsigned long      expires;
```

Arguments

timer

the timer to be modified

expires

new timeout in jiffies

Description

`mod_timer` is a more efficient way to update the `expire` field of an active timer (if the timer is inactive it will be activated)

`mod_timer(timer, expires)` is equivalent to:

```
del_timer(timer); timer->expires = expires; add_timer(timer);
```

Note that if there are multiple unserialized concurrent users of the same timer, then `mod_timer` is the only safe way to modify the timeout, since `add_timer` cannot modify an already running timer.

The function returns whether it has modified a pending timer or not. (ie. `mod_timer` of an inactive timer returns 0, `mod_timer` of an active timer returns 1.)

Name

`mod_timer_pinned` — modify a timer's timeout

Synopsis

```
int mod_timer_pinned (timer,  
                     expires);
```

```
struct timer_list * timer;  
unsigned long      expires;
```

Arguments

timer

the timer to be modified

expires

new timeout in jiffies

Description

`mod_timer_pinned` is a way to update the `expire` field of an active timer (if the timer is inactive it will be activated) and not allow the timer to be migrated to a different CPU.

`mod_timer_pinned(timer, expires)` is equivalent to:

```
del_timer(timer); timer->expires = expires; add_timer(timer);
```

Name

add_timer — start a timer

Synopsis

```
void add_timer (timer);  
  
struct timer_list * timer;
```

Arguments

timer

the timer to be added

Description

The kernel will do a `->function(->data)` callback from the timer interrupt at the `->expires` point in the future. The current time is 'jiffies'.

The timer's `->expires`, `->function` (and if the handler uses it, `->data`) fields must be set prior calling this function.

Timers with an `->expires` field in the past will be executed in the next timer tick.

Name

add_timer_on — start a timer on a particular CPU

Synopsis

```
void add_timer_on (timer,  
                  cpu);  
  
struct timer_list * timer;  
int                cpu;
```

Arguments

timer

the timer to be added

cpu

the CPU to start it on

Description

This is not very scalable on SMP. Double adds are not possible.

Name

`del_timer` — deactivate a timer.

Synopsis

```
int del_timer (timer);
```

```
struct timer_list * timer;
```

Arguments

timer

the timer to be deactivated

Description

`del_timer` deactivates a timer - this works on both active and inactive timers.

The function returns whether it has deactivated a pending timer or not. (ie. `del_timer` of an inactive timer returns 0, `del_timer` of an active timer returns 1.)

Name

`try_to_del_timer_sync` — Try to deactivate a timer

Synopsis

```
int try_to_del_timer_sync (timer);
```

```
struct timer_list * timer;
```

Arguments

timer

timer do del

Description

This function tries to deactivate a timer. Upon successful (`ret >= 0`) exit the timer is not queued and the handler is not running on any CPU.

It must not be called from interrupt contexts.

Name

`del_timer_sync` — deactivate a timer and wait for the handler to finish.

Synopsis

```
int del_timer_sync (timer);  
  
struct timer_list * timer;
```

Arguments

timer

the timer to be deactivated

Description

This function only differs from `del_timer` on SMP: besides deactivating the timer it also makes sure the handler has finished executing on other CPUs.

Synchronization rules

Callers must prevent restarting of the timer, otherwise this function is meaningless. It must not be called from interrupt contexts. The caller must not hold locks which would prevent completion of the timer's handler. The timer's handler must not call `add_timer_on`. Upon exit the timer is not queued and the handler is not running on any CPU.

The function returns whether it has deactivated a pending timer or not.

Name

`schedule_timeout` — sleep until timeout

Synopsis

```
signed long __sched schedule_timeout (timeout);
```

```
signed long timeout;
```

Arguments

timeout

timeout value in jiffies

Description

Make the current task sleep until *timeout* jiffies have elapsed. The routine will return immediately unless the current task state has been set (see `set_current_state`).

You can set the task state as follows -

`TASK_UNINTERRUPTIBLE` - at least *timeout* jiffies are guaranteed to pass before the routine returns. The routine will return 0

`TASK_INTERRUPTIBLE` - the routine may return early if a signal is delivered to the current task. In this case the remaining time in jiffies will be returned, or 0 if the timer expired in time

The current task state is guaranteed to be `TASK_RUNNING` when this routine returns.

Specifying a *timeout* value of `MAX_SCHEDULE_TIMEOUT` will schedule the CPU away without a bound on the timeout. In this case the return value will be `MAX_SCHEDULE_TIMEOUT`.

In all cases the return value is guaranteed to be non-negative.

Name

`msleep` — sleep safely even with waitqueue interruptions

Synopsis

```
void msleep (msecs);
```

```
unsigned int msecs;
```

Arguments

msecs

Time in milliseconds to sleep for

Name

`msleep_interruptible` — sleep waiting for signals

Synopsis

```
unsigned long msleep_interruptible (msecs);
```

```
unsigned int msecs;
```

Arguments

msecs

Time in milliseconds to sleep for

High-resolution timers

Name

`ktime_set` — Set a `ktime_t` variable from a seconds/nanoseconds value

Synopsis

```
ktime_t ktime_set (secs,  
                   nsecs);
```

```
const long          secs;  
const unsigned long nsecs;
```

Arguments

secs

seconds to set

nsecs

nanoseconds to set

Description

Return the `ktime_t` representation of the value

Name

`ktime_sub` — subtract two `ktime_t` variables

Synopsis

```
ktime_t ktime_sub (lhs,  
                  rhs);
```

```
const ktime_t lhs;  
const ktime_t rhs;
```

Arguments

lhs

minuend

rhs

subtrahend

Description

Returns the remainder of the subtraction

Name

`ktime_add` — add two `ktime_t` variables

Synopsis

```
ktime_t ktime_add (add1,  
                  add2);
```

```
const ktime_t add1;  
const ktime_t add2;
```

Arguments

add1

addend1

add2

addend2

Description

Returns the sum of *add1* and *add2*.

Name

`timespec_to_ktime` — convert a `timespec` to `ktime_t` format

Synopsis

```
ktime_t timespec_to_ktime (ts);
```

```
const struct timespec ts;
```

Arguments

ts

the `timespec` variable to convert

Description

Returns a `ktime_t` variable with the converted `timespec` value

Name

`timeval_to_ktime` — convert a `timeval` to `ktime_t` format

Synopsis

```
ktime_t timeval_to_ktime (tv);
```

```
const struct timeval tv;
```

Arguments

tv

the `timeval` variable to convert

Description

Returns a `ktime_t` variable with the converted `timeval` value

Name

`ktime_to_timespec` — convert a `ktime_t` variable to `timespec` format

Synopsis

```
struct timespec ktime_to_timespec (kt);  
  
const ktime_t kt;
```

Arguments

kt
the `ktime_t` variable to convert

Description

Returns the `timespec` representation of the `ktime` value

Name

`ktime_to_timeval` — convert a `ktime_t` variable to `timeval` format

Synopsis

```
struct timeval ktime_to_timeval (kt);  
  
const ktime_t kt;
```

Arguments

kt
the `ktime_t` variable to convert

Description

Returns the `timeval` representation of the `ktime` value

Name

`ktime_to_ns` — convert a `ktime_t` variable to scalar nanoseconds

Synopsis

```
s64 ktime_to_ns (kt);
```

```
const ktime_t kt;
```

Arguments

kt

the `ktime_t` variable to convert

Description

Returns the scalar nanoseconds representation of *kt*

Name

`ktime_equal` — Compares two `ktime_t` variables to see if they are equal

Synopsis

```
int ktime_equal (cmp1,  
                 cmp2);
```

```
const ktime_t cmp1;  
const ktime_t cmp2;
```

Arguments

cmp1

comparable1

cmp2

comparable2

Description

Compare two `ktime_t` variables, returns 1 if equal

Name

struct hrtimer — the basic hrtimer structure

Synopsis

```
struct hrtimer {
    struct rb_node node;
    ktime_t _expires;
    ktime_t _softexpires;
    enum hrtimer_restart (* function) (struct hrtimer *);
    struct hrtimer_clock_base * base;
    unsigned long state;
#ifdef CONFIG_TIMER_STATS
    int start_pid;
    void * start_site;
    char start_comm[16];
#endif
};
```

Members

node

red black tree node for time ordered insertion

_expires

the absolute expiry time in the hrtimers internal representation. The time is related to the clock on which the timer is based. Is setup by adding slack to the _softexpires value. For non range timers identical to _softexpires.

_softexpires

the absolute earliest expiry time of the hrtimer. The time which was given as expiry time when the timer was armed.

function

timer expiry callback function

base

pointer to the timer base (per cpu and per clock)

state

state information (See bit values above)

start_pid

timer statistics field to store the pid of the task which started the timer

start_site

timer statistics field to store the site where the timer was started

start_comm[16]

timer statistics field to store the name of the process which started the timer

Description

The hrtimer structure must be initialized by `hrtimer_init`

Name

`struct hrtimer_sleeper` — simple sleeper structure

Synopsis

```
struct hrtimer_sleeper {
    struct hrtimer timer;
    struct task_struct * task;
};
```

Members

timer

embedded timer structure

task

task to wake up

Description

task is set to NULL, when the timer expires.

Name

`struct hrtimer_clock_base` — the timer base for a specific clock

Synopsis

```
struct hrtimer_clock_base {
```

```
struct hrtimer_cpu_base * cpu_base;
clockid_t index;
struct rb_root active;
struct rb_node * first;
ktime_t resolution;
ktime_t (* get_time) (void);
ktime_t softirq_time;
#ifdef CONFIG_HIGH_RES_TIMERS
    ktime_t offset;
#endif
};
```

Members

cpu_base

per cpu clock base

index

clock type index for per_cpu support when moving a timer to a base on another cpu.

active

red black tree root node for the active timers

first

pointer to the timer node which expires first

resolution

the resolution of the clock, in nanoseconds

get_time

function to retrieve the current time of the clock

softirq_time

the time when running the hrtimer queue in the softirq

offset

offset of this clock to the monotonic base

Name

ktime_add_ns — Add a scalar nanoseconds value to a ktime_t variable

Synopsis


```
ktime_t ktime_add_ns (kt,  
                      nsec);
```

```
const ktime_t kt;  
u64          nsec;
```

Arguments

kt

addend

nsec

the scalar nsec value to add

Description

Returns the sum of *kt* and *nsec* in *ktime_t* format

Name

ktime_sub_ns — Subtract a scalar nanoseconds value from a *ktime_t* variable

Synopsis

```
ktime_t ktime_sub_ns (kt,  
                      nsec);
```

```
const ktime_t kt;  
u64          nsec;
```

Arguments

kt

minuend

nsec

the scalar nsec value to subtract

Description

Returns the subtraction of *nsec* from *kt* in *ktime_t* format

Name

`hrtimer_forward` — forward the timer expiry

Synopsis

```
u64 hrtimer_forward (timer,
                    now,
                    interval);
```

```
struct hrtimer * timer;
ktime_t          now;
ktime_t          interval;
```

Arguments

timer

hrtimer to forward

now

forward past this time

interval

the interval to forward

Description

Forward the timer expiry so it will expire in the future. Returns the number of overruns.

Name

`hrtimer_start_range_ns` — (re)start an hrtimer on the current CPU

Synopsis

```
int hrtimer_start_range_ns (timer,
                           tim,
                           delta_ns,
                           mode);
```

```
struct hrtimer *      timer;
ktime_t               tim;
unsigned long         delta_ns;
const enum hrtimer_mode mode;
```

Arguments

timer

the timer to be added

tim

expiry time

delta_ns

"slack" range for the timer

mode

expiry mode: absolute (HRTIMER_ABS) or relative (HRTIMER_REL)

Returns

0 on success 1 when the timer was active

Name

hrtimer_start — (re)start an hrtimer on the current CPU

Synopsis

```
int hrtimer_start (timer,  
                  tim,  
                  mode);  
  
struct hrtimer *      timer;  
ktime_t               tim;  
const enum hrtimer_mode mode;
```

Arguments

timer

the timer to be added

tim

expiry time

mode

expiry mode: absolute (HRTIMER_ABS) or relative (HRTIMER_REL)

Returns

0 on success 1 when the timer was active

Name

hrtimer_try_to_cancel — try to deactivate a timer

Synopsis

```
int hrtimer_try_to_cancel (timer);
```

```
struct hrtimer * timer;
```

Arguments

timer

hrtimer to stop

Returns

0 when the timer was not active 1 when the timer was active -1 when the timer is currently excuting the callback function and cannot be stopped

Name

hrtimer_cancel — cancel a timer and wait for the handler to finish.

Synopsis

```
int hrtimer_cancel (timer);
```

```
struct hrtimer * timer;
```

Arguments

timer

the timer to be cancelled

Returns

0 when the timer was not active 1 when the timer was active

Name

`hrtimer_get_remaining` — get remaining time for the timer

Synopsis

```
ktime_t hrtimer_get_remaining (timer);
```

```
const struct hrtimer * timer;
```

Arguments

timer

the timer to read

Name

`hrtimer_init` — initialize a timer to the given clock

Synopsis

```
void hrtimer_init (timer,  
                  clock_id,  
                  mode);
```

```
struct hrtimer *   timer;  
clockid_t          clock_id;  
enum hrtimer_mode  mode;
```

Arguments

timer

the timer to be initialized

clock_id

the clock to be used

mode

timer mode abs/rel

Name

hrtimer_get_res — get the timer resolution for a clock

Synopsis

```
int hrtimer_get_res (which_clock,
                    tp);
```

```
const clockid_t    which_clock;
struct timespec *  tp;
```

Arguments

which_clock

which clock to query

tp

pointer to timespec variable to store the resolution

Description

Store the resolution of the clock selected by *which_clock* in the variable pointed to by *tp*.

Name

schedule_hrtimeout_range — sleep until timeout

Synopsis

```
int __sched schedule_hrtimeout_range (expires,
                                     delta,
                                     mode);
```

```
ktime_t *          expires;
unsigned long      delta;
const enum hrtimer_mode mode;
```

Arguments

expires

timeout value (ktime_t)

delta

slack in expires timeout (ktime_t)

mode

timer mode, HRTIMER_MODE_ABS or HRTIMER_MODE_REL

Description

Make the current task sleep until the given expiry time has elapsed. The routine will return immediately unless the current task state has been set (see `set_current_state`).

The *delta* argument gives the kernel the freedom to schedule the actual wakeup to a time that is both power and performance friendly. The kernel give the normal best effort behavior for “*expires+delta*”, but may decide to fire the timer earlier, but no earlier than *expires*.

You can set the task state as follows -

`TASK_UNINTERRUPTIBLE` - at least *timeout* time is guaranteed to pass before the routine returns.

`TASK_INTERRUPTIBLE` - the routine may return early if a signal is delivered to the current task.

The current task state is guaranteed to be `TASK_RUNNING` when this routine returns.

Returns 0 when the timer has expired otherwise `-EINTR`

Name

`schedule_hrtimeout` — sleep until timeout

Synopsis

```
int __sched schedule_hrtimeout (expires,
                                mode);
```

```
ktime_t *                expires;
const enum hrtimer_mode mode;
```

Arguments

expires

timeout value (ktime_t)

mode

timer mode, `HRTIMER_MODE_ABS` or `HRTIMER_MODE_REL`

Description

Make the current task sleep until the given expiry time has elapsed. The routine will return immediately unless the current task state has been set (see `set_current_state`).

You can set the task state as follows -

`TASK_UNINTERRUPTIBLE` - at least *timeout* time is guaranteed to pass before the routine returns.

`TASK_INTERRUPTIBLE` - the routine may return early if a signal is delivered to the current task.

The current task state is guaranteed to be `TASK_RUNNING` when this routine returns.

Returns 0 when the timer has expired otherwise `-EINTR`

Workqueues and Kevents

Name

`queue_work` — queue work on a workqueue

Synopsis

```
int queue_work (wq,  
               work);
```

```
struct workqueue_struct * wq;  
struct work_struct *      work;
```

Arguments

wq

workqueue to use

work

work to queue

Description

Returns 0 if *work* was already on a queue, non-zero otherwise.

We queue the work to the CPU on which it was submitted, but if the CPU dies it can be processed by another CPU.

Name

`queue_work_on` — queue work on specific cpu

Synopsis

```
int queue_work_on (cpu,
                  wq,
                  work);

int                cpu;
struct workqueue_struct * wq;
struct work_struct * work;
```

Arguments

cpu

CPU number to execute work on

wq

workqueue to use

work

work to queue

Description

Returns 0 if *work* was already on a queue, non-zero otherwise.

We queue the work to a specific CPU, the caller must ensure it can't go away.

Name

`queue_delayed_work` — queue work on a workqueue after delay

Synopsis

```
int queue_delayed_work (wq,
                       dwork,
                       delay);

struct workqueue_struct * wq;
struct delayed_work * dwork;
unsigned long          delay;
```

Arguments

wq

workqueue to use

dwork

delayable work to queue

delay

number of jiffies to wait before queueing

Description

Returns 0 if *work* was already on a queue, non-zero otherwise.

Name

queue_delayed_work_on — queue work on specific CPU after delay

Synopsis

```
int queue_delayed_work_on (cpu,  
                           wq,  
                           dwork,  
                           delay);
```

```
int                cpu;  
struct workqueue_struct * wq;  
struct delayed_work * dwork;  
unsigned long      delay;
```

Arguments

cpu

CPU number to execute work on

wq

workqueue to use

dwork

work to queue

delay

number of jiffies to wait before queueing

Description

Returns 0 if *work* was already on a queue, non-zero otherwise.

Name

`flush_workqueue` — ensure that any scheduled work has run to completion.

Synopsis

```
void flush_workqueue (wq);  
  
struct workqueue_struct * wq;
```

Arguments

wq

workqueue to flush

Description

Forces execution of the workqueue and blocks until its completion. This is typically used in driver shutdown handlers.

We sleep until all works which were queued on entry have been handled, but we are not livelocked by new incoming ones.

This function used to run the workqueues itself. Now we just wait for the helper threads to do it.

Name

`flush_work` — block until a `work_struct`'s callback has terminated

Synopsis

```
int flush_work (work);  
  
struct work_struct * work;
```

Arguments

work

the work which is to be flushed

Description

Returns false if *work* has already terminated.

It is expected that, prior to calling `flush_work`, the caller has arranged for the work to not be requeued, otherwise it doesn't make sense to use this function.

Name

`cancel_work_sync` — block until a `work_struct`'s callback has terminated

Synopsis

```
int cancel_work_sync (work);
```

```
struct work_struct * work;
```

Arguments

work

the work which is to be flushed

Description

Returns true if *work* was pending.

`cancel_work_sync` will cancel the work if it is queued. If the work's callback appears to be running, `cancel_work_sync` will block until it has completed.

It is possible to use this function if the work re-queues itself. It can cancel the work even if it migrates to another workqueue, however in that case it only guarantees that `work->func` has completed on the last queued workqueue.

`cancel_work_sync(delayed_work->work)` should be used only if `->timer` is not pending, otherwise it goes into a busy-wait loop until the timer expires.

The caller must ensure that `workqueue_struct` on which this work was last queued can't be destroyed before this function returns.

Name

`cancel_delayed_work_sync` — reliably kill off a delayed work.

Synopsis

```
int cancel_delayed_work_sync (dwork);
```

```
struct delayed_work * dwork;
```

Arguments

dwork

the delayed work struct

Description

Returns true if *dwork* was pending.

It is possible to use this function if *dwork* rearms itself via `queue_work` or `queue_delayed_work`. See also the comment for `cancel_work_sync`.

Name

`schedule_work` — put work task in global workqueue

Synopsis

```
int schedule_work (work);
```

```
struct work_struct * work;
```

Arguments

work

job to be done

Description

Returns zero if *work* was already on the kernel-global workqueue and non-zero otherwise.

This puts a job in the kernel-global workqueue if it was not already queued and leaves it in the same position on the kernel-global workqueue otherwise.

Name

`schedule_delayed_work` — put work task in global workqueue after delay

Synopsis

```
int schedule_delayed_work (dwork,  
                           delay);
```

```
struct delayed_work * dwork;  
unsigned long        delay;
```

Arguments

dwork

job to be done

delay

number of jiffies to wait or 0 for immediate execution

Description

After waiting for a given time this puts a job in the kernel-global workqueue.

Name

`flush_delayed_work` — block until a `dwork_struct`'s callback has terminated

Synopsis

```
void flush_delayed_work (dwork);
```

```
struct delayed_work * dwork;
```

Arguments

dwork

the delayed work which is to be flushed

Description

Any timeout is cancelled, and any pending work is run immediately.

Name

`schedule_delayed_work_on` — queue work in global workqueue on CPU after delay

Synopsis

```
int schedule_delayed_work_on (cpu,  
                               dwork,  
                               delay);
```

```
int                cpu;  
struct delayed_work * dwork;  
unsigned long      delay;
```

Arguments

cpu

cpu to use

dwork

job to be done

delay

number of jiffies to wait

Description

After waiting for a given time this puts a job in the kernel-global workqueue on the specified CPU.

Name

`execute_in_process_context` — reliably execute the routine with user context

Synopsis

```
int execute_in_process_context (fn,  
                                ew);
```

```
work_func_t        fn;  
struct execute_work * ew;
```

Arguments

fn

the function to execute

ew

guaranteed storage for the execute work structure (must be available when the work executes)

Description

Executes the function immediately if process context is available, otherwise schedules the function for delayed execution.

Returns

0 - function was executed 1 - function was scheduled for execution

Name

`destroy_workqueue` — safely terminate a workqueue

Synopsis

```
void destroy_workqueue (wq);  
  
struct workqueue_struct * wq;
```

Arguments

wq

target workqueue

Description

Safely destroy a workqueue. All work currently pending will be done first.

Name

`work_on_cpu` — run a function in user context on a particular cpu

Synopsis


```
long work_on_cpu (cpu,  
                  fn,  
                  arg);  
  
unsigned int cpu;  
long (*      fn(void *));  
void *      arg;
```

Arguments

cpu

the cpu to run on

fn

the function to run

arg

the function arg

Description

This will return the value *fn* returns. It is up to the caller to ensure that the cpu doesn't go offline. The caller must not hold any locks which would prevent *fn* from completing.

Internal Functions

Name

reparent_to_kthreadd — Reparent the calling kernel thread to kthreadd

Synopsis

```
void reparent_to_kthreadd (void);  
  
void;
```

Arguments

void

no arguments

Description

If a kernel thread is launched as a result of a system call, or if it ever exits, it should generally reparent itself to `kthreadd` so it isn't in the way of other processes and is correctly cleaned up on exit.

The various task state such as scheduling policy and priority may have been inherited from a user process, so we reset them to sane values here.

NOTE that `reparent_to_kthreadd` gives the caller full capabilities.

Name

`sys_tgkill` — send signal to one specific thread

Synopsis

```
long sys_tgkill (tgid,  
                pid,  
                sig);
```

```
pid_t tgid;  
pid_t pid;  
int sig;
```

Arguments

tgid

the thread group ID of the thread

pid

the PID of the thread

sig

signal to be sent

Description

This syscall also checks the *tgid* and returns `-ESRCH` even if the PID exists but it's not belonging to the target process anymore. This method solves the problem of threads exiting and PIDs getting reused.

Name

`kthread_run` — create and wake a thread.

Synopsis

```
kthread_run (threadfn,  
             data,  
             namefmt,  
             ...);
```

```
threadfn;  
data;  
namefmt;  
...;
```

Arguments

threadfn

the function to run until `signal_pending(current)`.

data

data ptr for *threadfn*.

namefmt

printf-style name for the thread.

...

variable arguments

Description

Convenient wrapper for `kthread_create` followed by `wake_up_process`. Returns the kthread or `ERR_PTR(-ENOMEM)`.

Name

`kthread_should_stop` — should this kthread return now?

Synopsis

```
int kthread_should_stop (void);
```

```
void;
```

Arguments

void

no arguments

Description

When someone calls `kthread_stop` on your `kthread`, it will be woken and this will return true. You should then return, and your return value will be passed through to `kthread_stop`.

Name

`kthread_create` — create a `kthread`.

Synopsis

```
struct task_struct * kthread_create (threadfn,
                                       data,
                                       namefmt[],
                                       ...);
```

```
int (*      threadfn(void *data);
void *      data;
const char namefmt[];
          ...;
```

Arguments

threadfn

the function to run until `signal_pending(current)`.

data

data ptr for *threadfn*.

namefmt[]

printf-style name for the thread.

...

variable arguments

Description

This helper function creates and names a kernel thread. The thread will be stopped: use `wake_up_process` to start it. See also `kthread_run`, `kthread_create_on_cpu`.

When woken, the thread will run `threadfn()` with *data* as its argument. `threadfn()` can either call `do_exit` directly if it is a standalone thread for which noone will call `kthread_stop`, or return when 'kthread_should_stop' is true (which means `kthread_stop` has been called). The return value should

be zero or a negative error number; it will be passed to `kthread_stop`.

Returns a `task_struct` or `ERR_PTR(-ENOMEM)`.

Name

`kthread_bind` — bind a just-created `kthread` to a `cpu`.

Synopsis

```
void kthread_bind (k,  
                  cpu);  
  
struct task_struct * k;  
unsigned int         cpu;
```

Arguments

k

thread created by `kthread_create`.

cpu

`cpu` (might not be online, must be possible) for *k* to run on.

Description

This function is equivalent to `set_cpus_allowed`, except that *cpu* doesn't need to be online, and the thread must be stopped (i.e., just returned from `kthread_create`).

Name

`kthread_stop` — stop a thread created by `kthread_create`.

Synopsis

```
int kthread_stop (k);  
  
struct task_struct * k;
```

Arguments

k

thread created by `kthread_create`.

Description

Sets `kthread_should_stop` for *k* to return true, wakes it, and waits for it to exit. This can also be called after `kthread_create` instead of calling `wake_up_process`: the thread will exit without calling `threadfn`.

If `threadfn` may call `do_exit` itself, the caller must ensure `task_struct` can't go away.

Returns the result of `threadfn`, or `-EINTR` if `wake_up_process` was never called.

Kernel objects manipulation

Name

`kobject_get_path` — generate and return the path associated with a given `kobj` and `kset` pair.

Synopsis

```
char * kobject_get_path (kobj,
                        gfp_mask);
```

```
struct kobject * kobj;
gfp_t           gfp_mask;
```

Arguments

kobj

kobject in question, with which to build the path

gfp_mask

the allocation type used to allocate the path

Description

The result must be freed by the caller with `kfree`.

Name

`kobject_set_name` — Set the name of a `kobject`

Synopsis

```
int kobject_set_name (kobj,  
                     fmt,  
                     ...);  
  
struct kobject * kobj;  
const char *    fmt;  
...;
```

Arguments

kobj

struct kobject to set the name of

fmt

format string used to build the name

...

variable arguments

Description

This sets the name of the kobject. If you have already added the kobject to the system, you must call `kobject_rename` in order to change the name of the kobject.

Name

`kobject_init` — initialize a kobject structure

Synopsis

```
void kobject_init (kobj,  
                  ktype);  
  
struct kobject *    kobj;  
struct kobj_type *  ktype;
```

Arguments

kobj

pointer to the kobject to initialize

ktype

pointer to the ktype for this kobject.

Description

This function will properly initialize a kobject such that it can then be passed to the `kobject_add` call.

After this function is called, the kobject **MUST** be cleaned up by a call to `kobject_put`, not by a call to `kfree` directly to ensure that all of the memory is cleaned up properly.

Name

`kobject_add` — the main kobject add function

Synopsis

```
int kobject_add (kobj,  
                 parent,  
                 fmt,  
                 ...);  
  
struct kobject * kobj;  
struct kobject * parent;  
const char *    fmt;  
               ...;
```

Arguments

kobj

the kobject to add

parent

pointer to the parent of the kobject.

fmt

format to name the kobject with.

...

variable arguments

Description

The kobject name is set and added to the kobject hierarchy in this function.

If *parent* is set, then the parent of the *kobj* will be set to it. If *parent* is NULL, then the parent of the *kobj* will be set to the kobject associated with the kset assigned to this kobject. If no kset is assigned to

the `kobject`, then the `kobject` will be located in the root of the `sysfs` tree.

If this function returns an error, `kobject_put` must be called to properly clean up the memory associated with the object. Under no instance should the `kobject` that is passed to this function be directly freed with a call to `kfree`, that can leak memory.

Note, no “add” uevent will be created with this call, the caller should set up all of the necessary `sysfs` files for the object and then call `kobject_uevent` with the `UEVENT_ADD` parameter to ensure that userspace is properly notified of this `kobject`'s creation.

Name

`kobject_init_and_add` — initialize a `kobject` structure and add it to the `kobject` hierarchy

Synopsis

```
int kobject_init_and_add (kobj,
                        ktype,
                        parent,
                        fmt,
                        ...);
```

```
struct kobject *    kobj;
struct kobj_type *  ktype;
struct kobject *    parent;
const char *        fmt;
...;
```

Arguments

kobj

pointer to the `kobject` to initialize

ktype

pointer to the `ktype` for this `kobject`.

parent

pointer to the parent of this `kobject`.

fmt

the name of the `kobject`.

...

variable arguments

Description

This function combines the call to `kobject_init` and `kobject_add`. The same type of error handling after a call to `kobject_add` and `kobject` lifetime rules are the same here.

Name

`kobject_rename` — change the name of an object

Synopsis

```
int kobject_rename (kobj,  
                    new_name);
```

```
struct kobject * kobj;  
const char *    new_name;
```

Arguments

kobj

object in question.

new_name

object's new name

Description

It is the responsibility of the caller to provide mutual exclusion between two different calls of `kobject_rename` on the same `kobject` and to ensure that `new_name` is valid and won't conflict with other `kobjects`.

Name

`kobject_del` — unlink `kobject` from hierarchy.

Synopsis

```
void kobject_del (kobj);
```

```
struct kobject * kobj;
```

Arguments

kobj

object.

Name

`kobject_get` — increment refcount for object.

Synopsis

```
struct kobject * kobject_get (kobj);
```

```
struct kobject * kobj;
```

Arguments

kobj

object.

Name

`kobject_put` — decrement refcount for object.

Synopsis

```
void kobject_put (kobj);
```

```
struct kobject * kobj;
```

Arguments

kobj

object.

Description

Decrement the refcount, and if 0, call `kobject_cleanup`.

Name

`kobject_create_and_add` — create a struct kobject dynamically and register it with sysfs

Synopsis

```
struct kobject * kobject_create_and_add (name,  
                                           parent);  
  
const char *      name;  
struct kobject * parent;
```

Arguments

name

the name for the kset

parent

the parent kobject of this kobject, if any.

Description

This function creates a kobject structure dynamically and registers it with sysfs. When you are finished with this structure, call `kobject_put` and the structure will be dynamically freed when it is no longer being used.

If the kobject was not able to be created, NULL will be returned.

Name

`kset_register` — initialize and add a kset.

Synopsis

```
int kset_register (k);  
  
struct kset * k;
```

Arguments

k

kset.

Name

`kset_unregister` — remove a kset.

Synopsis

```
void kset_unregister (k);
```

```
struct kset * k;
```

Arguments

k

kset.

Name

kset_create_and_add — create a struct kset dynamically and add it to sysfs

Synopsis

```
struct kset * kset_create_and_add (name,  
                                   uevent_ops,  
                                   parent_kobj);
```

```
const char *      name;  
struct kset_uevent_ops * uevent_ops;  
struct kobject *  parent_kobj;
```

Arguments

name

the name for the kset

uevent_ops

a struct kset_uevent_ops for the kset

parent_kobj

the parent kobject of this kset, if any.

Description

This function creates a kset structure dynamically and registers it with sysfs. When you are finished with this structure, call kset_unregister and the structure will be dynamically freed when it is no longer being used.

If the kset was not able to be created, NULL will be returned.

Kernel utility functions

Name

`upper_32_bits` — return bits 32-63 of a number

Synopsis

```
upper_32_bits (n);
```

```
n;
```

Arguments

n

the number we're accessing

Description

A basic shift-right of a 64- or 32-bit quantity. Use this to suppress the “right shift count \geq width of type” warning when that quantity is 32-bits.

Name

`lower_32_bits` — return bits 0-31 of a number

Synopsis

```
lower_32_bits (n);
```

```
n;
```

Arguments

n

the number we're accessing

Name

`might_sleep` — annotation for functions that can sleep

Synopsis

```
might_sleep ();
```

Arguments

None

Description

this macro will print a stack trace if it is executed in an atomic context (spinlock, irq-handler, ...).

This is a useful debugging help to be able to catch problems early and not be bitten later when the calling function happens to sleep when it is not supposed to.

Name

`trace_printk` — printf formatting in the ftrace buffer

Synopsis

```
trace_printk (fmt,  
               args...);
```

```
fmt;  
args...;
```

Arguments

fmt

the printf format for printing

args...

variable arguments

Note

`__trace_printk` is an internal function for `trace_printk` and the *ip* is passed in via the `trace_printk` macro.

This function allows a kernel developer to debug fast path sections that `printk` is not appropriate for. By scattering in various `printk` like tracing in the code, a developer can quickly see where problems are

occurring.

This is intended as a debugging tool for the developer only. Please refrain from leaving `trace_printks` scattered around in your code.

Name

`clamp` — return a value clamped to a given range with strict typechecking

Synopsis

```
clamp (val,  
        min,  
        max);
```

```
val;  
min;  
max;
```

Arguments

val

current value

min

minimum allowable value

max

maximum allowable value

Description

This macro does strict typechecking of `min/max` to make sure they are of the same type as `val`. See the unnecessary pointer comparisons.

Name

`clamp_t` — return a value clamped to a given range using a given type

Synopsis

```
clamp_t (type,  
         val,
```



```
        min,  
        max);  
  
type;  
val;  
min;  
max;
```

Arguments

type

the type of variable to use

val

current value

min

minimum allowable value

max

maximum allowable value

Description

This macro does no typechecking and uses temporary variables of type 'type' to make all the comparisons.

Name

`clamp_val` — return a value clamped to a given range using val's type

Synopsis

```
clamp_val (val,  
           min,  
           max);  
  
val;  
min;  
max;
```

Arguments

val

current value

min

minimum allowable value

max

maximum allowable value

Description

This macro does no typechecking and uses temporary variables of whatever type the input argument 'val' is. This is useful when val is an unsigned type and min and max are literals that will otherwise be assigned a signed integer type.

Name

`container_of` — cast a member of a structure out to the containing structure

Synopsis

```
container_of (ptr,  
              type,  
              member);
```

```
ptr;  
type;  
member;
```

Arguments

ptr

the pointer to the member.

type

the type of the container struct this is embedded in.

member

the name of the member within the struct.

Name

`printk` — print a kernel message

Synopsis

```
int printk (fmt,
            ...);

const char * fmt;
            ...;
```

Arguments

fmt

format string

...

variable arguments

Description

This is `printk`. It can be called from any context. We want it to work.

We try to grab the `console_sem`. If we succeed, it's easy - we log the output and call the console drivers. If we fail to get the semaphore we place the output into the log buffer and return. The current holder of the `console_sem` will notice the new output in `release_console_sem` and will send it to the consoles before releasing the semaphore.

One effect of this deferred printing is that code which calls `printk` and then changes `console_loglevel` may break. This is because `console_loglevel` is inspected when the actual printing occurs.

See also

`printf(3)`

See the `vsnprintf` documentation for format string extensions over C99.

Name

`acquire_console_sem` — lock the console system for exclusive use.

Synopsis

```
void acquire_console_sem (void);

void;
```

Arguments

void

no arguments

Description

Acquires a semaphore which guarantees that the caller has exclusive access to the console system and the `console_drivers` list.

Can sleep, returns nothing.

Name

`release_console_sem` — unlock the console system

Synopsis

```
void release_console_sem (void);
```

```
void;
```

Arguments

void

no arguments

Description

Releases the semaphore which the caller holds on the console system and the console driver list.

While the semaphore was held, console output may have been buffered by `printk`. If this is the case, `release_console_sem` emits the output prior to releasing the semaphore.

If there is output waiting for `klogd`, we wake it up.

`release_console_sem` may be called from any context.

Name

`console_conditional_schedule` — yield the CPU if required

Synopsis

```
void __sched console_conditional_schedule (void);  
  
void;
```

Arguments

void

no arguments

Description

If the console code is currently allowed to sleep, and if this CPU should yield the CPU to another task, do so here.

Must be called within `acquire_console_sem`.

Name

`printk_timed_ratelimit` — caller-controlled printk ratelimiting

Synopsis

```
bool printk_timed_ratelimit (caller_jiffies,  
                             interval_msecs);  
  
unsigned long * caller_jiffies;  
unsigned int   interval_msecs;
```

Arguments

caller_jiffies

pointer to caller's state

interval_msecs

minimum interval between prints

Description

`printk_timed_ratelimit` returns true if more than *interval_msecs* milliseconds have elapsed since the last time `printk_timed_ratelimit` returned true.

Name

panic — halt the system

Synopsis

```
NORET_TYPE void panic (fmt,  
                        ...);  
  
const char * fmt;  
            ...;
```

Arguments

fmt
The text string to print
...
variable arguments

Description

Display a message, then perform cleanups.

This function never returns.

Name

emergency_restart — reboot the system

Synopsis

```
void emergency_restart (void);  
  
void;
```

Arguments

void
no arguments

Description

Without shutting down any hardware or taking any locks reboot the system. This is called when we know we are in trouble so this is our best effort to reboot. This is safe to call in interrupt context.

Name

kernel_restart — reboot the system

Synopsis

```
void kernel_restart (cmd);
```

```
char * cmd;
```

Arguments

cmd

pointer to buffer containing command to execute for restart or NULL

Description

Shutdown everything and perform a clean reboot. This is not safe to call in interrupt context.

Name

kernel_halt — halt the system

Synopsis

```
void kernel_halt (void);
```

```
void;
```

Arguments

void

no arguments

Description

Shutdown everything and perform a clean system halt.

Name

`kernel_power_off` — `power_off` the system

Synopsis

```
void kernel_power_off (void);  
  
void;
```

Arguments

void
no arguments

Description

Shutdown everything and perform a clean system `power_off`.

Name

`orderly_poweroff` — Trigger an orderly system poweroff

Synopsis

```
int orderly_poweroff (force);  
  
bool force;
```

Arguments

force
force poweroff if command execution fails

Description

This may be called from any context to trigger a system shutdown. If the orderly shutdown fails, it will force an immediate shutdown.

Name

`synchronize_rcu` — wait until a grace period has elapsed.

Synopsis

```
void synchronize_rcu (void);  
  
void;
```

Arguments

void

no arguments

Description

Control will return to the caller some time after a full grace period has elapsed, in other words after all currently executing RCU read-side critical sections have completed. RCU read-side critical sections are delimited by `rcu_read_lock` and `rcu_read_unlock`, and may be nested.

Name

`synchronize_sched` — wait until an rcu-sched grace period has elapsed.

Synopsis

```
void synchronize_sched (void);  
  
void;
```

Arguments

void

no arguments

Description

Control will return to the caller some time after a full rcu-sched grace period has elapsed, in other words after all currently executing rcu-sched read-side critical sections have completed. These read-side critical sections are delimited by `rcu_read_lock_sched` and `rcu_read_unlock_sched`, and may be nested. Note that `preempt_disable`, `local_irq_disable`, and so on may be used in place of `rcu_read_lock_sched`.

This means that all `preempt_disable` code sequences, including NMI and hardware-interrupt handlers, in progress on entry will have completed before this primitive returns. However, this does not guarantee

that softirq handlers will have completed, since in some kernels, these handlers can run in process context, and can block.

This primitive provides the guarantees made by the (now removed) `synchronize_kernel` API. In contrast, `synchronize_rcu` only guarantees that `rcu_read_lock` sections will have completed. In “classic RCU”, these two guarantees happen to be one and the same, but can differ in realtime RCU implementations.

Name

`synchronize_rcu_bh` — wait until an `rcu_bh` grace period has elapsed.

Synopsis

```
void synchronize_rcu_bh (void);  
  
void;
```

Arguments

```
void  
  
no arguments
```

Description

Control will return to the caller some time after a full `rcu_bh` grace period has elapsed, in other words after all currently executing `rcu_bh` read-side critical sections have completed. RCU read-side critical sections are delimited by `rcu_read_lock_bh` and `rcu_read_unlock_bh`, and may be nested.

Device Resource Management

Name

`devres_alloc` — Allocate device resource data

Synopsis

```
void * devres_alloc (release,  
                    size,  
                    gfp);  
  
dr_release_t  release;  
size_t        size;  
gfp_t         gfp;
```

Arguments

release

Release function devres will be associated with

size

Allocation size

gfp

Allocation flags

Description

Allocate devres of *size* bytes. The allocated area is zeroed, then associated with *release*. The returned pointer can be passed to other devres_*() functions.

RETURNS

Pointer to allocated devres on success, NULL on failure.

Name

devres_free — Free device resource data

Synopsis

```
void devres_free (res);
```

```
void * res;
```

Arguments

res

Pointer to devres data to free

Description

Free devres created with devres_alloc.

Name

devres_add — Register device resource

Synopsis

```
void devres_add (dev,  
                res);
```

```
struct device * dev;  
void *         res;
```

Arguments

dev

Device to add resource to

res

Resource to register

Description

Register devres *res* to *dev*. *res* should have been allocated using `devres_alloc`. On driver detach, the associated release function will be invoked and devres will be freed automatically.

Name

devres_find — Find device resource

Synopsis

```
void * devres_find (dev,  
                   release,  
                   match,  
                   match_data);
```

```
struct device * dev;  
dr_release_t   release;  
dr_match_t     match;  
void *         match_data;
```

Arguments

dev

Device to lookup resource from

release

Look for resources associated with this release function

match

Match function (optional)

match_data

Data for the match function

Description

Find the latest devres of *dev* which is associated with *release* and for which *match* returns 1. If *match* is NULL, it's considered to match all.

RETURNS

Pointer to found devres, NULL if not found.

Name

`devres_get` — Find devres, if non-existent, add one atomically

Synopsis

```
void * devres_get (dev,  
                  new_res,  
                  match,  
                  match_data);
```

```
struct device * dev;  
void *          new_res;  
dr_match_t      match;  
void *          match_data;
```

Arguments

dev

Device to lookup or add devres for

new_res

Pointer to new initialized devres to add if not found

match

Match function (optional)

match_data

Data for the match function

Description

Find the latest devres of *dev* which has the same release function as *new_res* and for which *match* return 1. If found, *new_res* is freed; otherwise, *new_res* is added atomically.

RETURNS

Pointer to found or added devres.

Name

`devres_remove` — Find a device resource and remove it

Synopsis

```
void * devres_remove (dev,  
                      release,  
                      match,  
                      match_data);
```

```
struct device * dev;  
dr_release_t   release;  
dr_match_t     match;  
void *         match_data;
```

Arguments

dev

Device to find resource from

release

Look for resources associated with this release function

match

Match function (optional)

match_data

Data for the match function

Description

Find the latest devres of *dev* associated with *release* and for which *match* returns 1. If *match* is NULL, it's considered to match all. If found, the resource is removed atomically and returned.

RETURNS

Pointer to removed devres on success, NULL if not found.

Name

`devres_destroy` — Find a device resource and destroy it

Synopsis

```
int devres_destroy (dev,
                   release,
                   match,
                   match_data);
```

```
struct device * dev;
dr_release_t   release;
dr_match_t     match;
void *         match_data;
```

Arguments

dev

Device to find resource from

release

Look for resources associated with this release function

match

Match function (optional)

match_data

Data for the match function

Description

Find the latest devres of *dev* associated with *release* and for which *match* returns 1. If *match* is NULL, it's considered to match all. If found, the resource is removed atomically and freed.

RETURNS

0 if devres is found and freed, -ENOENT if not found.

Name

devres_open_group — Open a new devres group

Synopsis

```
void * devres_open_group (dev,  
                           id,  
                           gfp);
```

```
struct device * dev;  
void *          id;  
gfp_t          gfp;
```

Arguments

dev

Device to open devres group for

id

Separator ID

gfp

Allocation flags

Description

Open a new devres group for *dev* with *id*. For *id*, using a pointer to an object which won't be used for another group is recommended. If *id* is NULL, address-wise unique ID is created.

RETURNS

ID of the new group, NULL on failure.

Name

devres_close_group — Close a devres group

Synopsis

```
void devres_close_group (dev,  
                        id);
```

```
struct device * dev;  
void * id;
```

Arguments

dev

Device to close devres group for

id

ID of target group, can be NULL

Description

Close the group identified by *id*. If *id* is NULL, the latest open group is selected.

Name

devres_remove_group — Remove a devres group

Synopsis

```
void devres_remove_group (dev,  
                        id);
```

```
struct device * dev;  
void * id;
```

Arguments

dev

Device to remove group for

id

ID of target group, can be NULL

Description

Remove the group identified by *id*. If *id* is NULL, the latest open group is selected. Note that removing a group doesn't affect any other resources.

Name

`devres_release_group` — Release resources in a devres group

Synopsis

```
int devres_release_group (dev,
                          id);
```

```
struct device * dev;
void * id;
```

Arguments

dev

Device to release group for

id

ID of target group, can be NULL

Description

Release all resources in the group identified by *id*. If *id* is NULL, the latest open group is selected. The selected group and groups properly nested inside the selected group are removed.

RETURNS

The number of released non-group resources.

Name

`devm_kzalloc` — Resource-managed kzalloc

Synopsis

```
void * devm_kzalloc (dev,
                    size,
                    gfp);
```

```
struct device * dev;
```

```
size_t      size;  
gfp_t      gfp;
```

Arguments

dev

Device to allocate memory for

size

Allocation size

gfp

Allocation gfp flags

Description

Managed kcalloc. Memory allocated with this function is automatically freed on driver detach. Like all other devres resources, guaranteed alignment is unsigned long long.

RETURNS

Pointer to allocated memory on success, NULL on failure.

Name

devm_kfree — Resource-managed kfree

Synopsis

```
void devm_kfree (dev,  
                p);
```

```
struct device * dev;  
void * p;
```

Arguments

dev

Device this memory belongs to

p

Memory to free

Description

Free memory allocated with `dev_kzalloc`.

Chapter 2. Device drivers infrastructure

Table of Contents

[Device Drivers Base](#)

[Device Drivers Power Management](#)

[Device Drivers ACPI Support](#)

[Device drivers PnP support](#)

[Userspace IO devices](#)

Device Drivers Base

Name

`driver_for_each_device` — Iterator for devices bound to a driver.

Synopsis

```
int driver_for_each_device (drv,  
                           start,  
                           data,  
                           fn);  
  
struct device_driver * drv;  
struct device * start;  
void * data;  
int (* fn(struct device *, void *);
```

Arguments

drv

Driver we're iterating.

start

Device to begin with

data

Data to pass to the callback.

fn

Function to call for each device.

Description

Iterate over the *drv*'s list of devices calling *fn* for each one.

Name

`driver_find_device` — device iterator for locating a particular device.

Synopsis

```
struct device * driver_find_device (drv,
                                     start,
                                     data,
                                     match);

struct device_driver * drv;
struct device *       start;
void *                data;
int (*                match(struct device *dev, void *data);
```

Arguments

drv

The device's driver

start

Device to begin with

data

Data to pass to match function

match

Callback function to check device

Description

This is similar to the `driver_for_each_device` function above, but it returns a reference to a device that is 'found' for later use, as determined by the *match* callback.

The callback should return 0 if the device doesn't match and non-zero if it does. If the callback returns non-zero, this function will return to the caller and not iterate over any more devices.

Name

`driver_create_file` — create sysfs file for driver.

Synopsis

```
int driver_create_file (drv,  
                        attr);  
  
struct device_driver *    drv;  
struct driver_attribute * attr;
```

Arguments

drv

driver.

attr

driver attribute descriptor.

Name

`driver_remove_file` — remove sysfs file for driver.

Synopsis

```
void driver_remove_file (drv,  
                         attr);  
  
struct device_driver *    drv;  
struct driver_attribute * attr;
```

Arguments

drv

driver.

attr

driver attribute descriptor.

Name

`driver_add_kobj` — add a kobject below the specified driver

Synopsis

```
int driver_add_kobj (drv,  
                    kobj,  
                    fmt,  
                    ...);  
  
struct device_driver * drv;  
struct kobject *      kobj;  
const char *          fmt;  
...;
```

Arguments

drv

requesting device driver

kobj

kobject to add below this driver

fmt

format string that names the kobject

...

variable arguments

Description

You really don't want to do this, this is only here due to one looney iseries driver, go poke those developers if you are annoyed about this...

Name

`get_driver` — increment driver reference count.

Synopsis

```
struct device_driver * get_driver (drv);  
  
struct device_driver * drv;
```

Arguments

drv

driver.

Name

`put_driver` — decrement driver's refcount.

Synopsis

```
void put_driver (drv);
```

```
struct device_driver * drv;
```

Arguments

drv

driver.

Name

`driver_register` — register driver with bus

Synopsis

```
int driver_register (drv);
```

```
struct device_driver * drv;
```

Arguments

drv

driver to register

Description

We pass off most of the work to the `bus_add_driver` call, since most of the things we have to do deal with the bus structures.

Name

`driver_unregister` — remove driver from system.

Synopsis

```
void driver_unregister (drv);  
  
struct device_driver * drv;
```

Arguments

drv
driver.

Description

Again, we pass off most of the work to the bus-level call.

Name

`driver_find` — locate driver on a bus by its name.

Synopsis

```
struct device_driver * driver_find (name,  
                                     bus);  
  
const char *      name;  
struct bus_type * bus;
```

Arguments

name
name of the driver.

bus
bus to scan for the driver.

Description

Call `kset_find_obj` to iterate over list of drivers on a bus to find driver by name. Return driver if found.

Note that `kset_find_obj` increments driver's reference count.

Name

`dev_driver_string` — Return a device's driver name, if at all possible

Synopsis

```
const char * dev_driver_string (dev);
```

```
const struct device * dev;
```

Arguments

dev

struct device to get the name of

Description

Will return the device's driver's name if it is bound to a device. If the device is not bound to a device, it will return the name of the bus it is attached to. If it is not attached to a bus either, an empty string will be returned.

Name

`device_create_file` — create sysfs attribute file for device.

Synopsis

```
int device_create_file (dev,  
                        attr);
```

```
struct device * dev;  
struct device_attribute * attr;
```

Arguments

dev

device.

attr

device attribute descriptor.

Name

`device_remove_file` — remove sysfs attribute file.

Synopsis

```
void device_remove_file (dev,  
                        attr);  
  
struct device *      dev;  
struct device_attribute * attr;
```

Arguments

dev

device.

attr

device attribute descriptor.

Name

`device_create_bin_file` — create sysfs binary attribute file for device.

Synopsis

```
int device_create_bin_file (dev,  
                           attr);  
  
struct device *      dev;  
struct bin_attribute * attr;
```

Arguments

dev

device.

attr

device binary attribute descriptor.

Name

`device_remove_bin_file` — remove sysfs binary attribute file

Synopsis

```
void device_remove_bin_file (dev,  
                             attr);
```

```
struct device *      dev;  
struct bin_attribute * attr;
```

Arguments

dev

device.

attr

device binary attribute descriptor.

Name

`device_schedule_callback_owner` — helper to schedule a callback for a device

Synopsis

```
int device_schedule_callback_owner (dev,  
                                     func,  
                                     owner);
```

```
struct device * dev;  
void (*        func(struct device *));  
struct module * owner;
```

Arguments

dev

device.

func

callback function to invoke later.

owner

module owning the callback routine

Description

Attribute methods must not unregister themselves or their parent device (which would amount to the same thing). Attempts to do so will deadlock, since unregistration is mutually exclusive with driver callbacks.

Instead methods can call this routine, which will attempt to allocate and schedule a workqueue request to call back *func* with *dev* as its argument in the workqueue's process context. *dev* will be pinned until *func* returns.

This routine is usually called via the inline `device_schedule_callback`, which automatically sets *owner* to `THIS_MODULE`.

Returns 0 if the request was submitted, `-ENOMEM` if storage could not be allocated, `-ENODEV` if a reference to *owner* isn't available.

NOTE

This routine won't work if `CONFIG_SYSFS` isn't set! It uses an underlying sysfs routine (since it is intended for use by attribute methods), and if sysfs isn't available you'll get nothing but `-ENOSYS`.

Name

`device_initialize` — init device structure.

Synopsis

```
void device_initialize (dev);
```

```
struct device * dev;
```

Arguments

dev

device.

Description

This prepares the device for use by other layers by initializing its fields. It is the first half of `device_register`, if called by that function, though it can also be called separately, so one may use *dev*'s fields. In particular, `get_device`/`put_device` may be used for reference counting of *dev* after calling this function.

NOTE

Use `put_device` to give up your reference instead of freeing `dev` directly once you have called this function.

Name

`dev_set_name` — set a device name

Synopsis

```
int dev_set_name (dev,  
                  fmt,  
                  ...);
```

```
struct device * dev;  
const char *   fmt;  
...;
```

Arguments

dev

device

fmt

format string for the device's name

...

variable arguments

Name

`device_add` — add device to device hierarchy.

Synopsis

```
int device_add (dev);
```

```
struct device * dev;
```

Arguments

dev

device.

Description

This is part 2 of `device_register`, though may be called separately if `_device_initialize` has been called separately.

This adds *dev* to the kobject hierarchy via `kobject_add`, adds it to the global and sibling lists for the device, then adds it to the other relevant subsystems of the driver model.

NOTE

Never directly free *dev* after calling this function, even if it returned an error! Always use `put_device` to give up your reference instead.

Name

`device_register` — register a device with the system.

Synopsis

```
int device_register (dev);
```

```
struct device * dev;
```

Arguments

dev

pointer to the device structure

Description

This happens in two clean steps - initialize the device and add it to the system. The two steps can be called separately, but this is the easiest and most common. I.e. you should only call the two helpers separately if have a clearly defined need to use and refcount the device before it is added to the hierarchy.

NOTE

Never directly free *dev* after calling this function, even if it returned an error! Always use `put_device` to give up the reference initialized in this function instead.

Name

`get_device` — increment reference count for device.

Synopsis

```
struct device * get_device (dev);
```

```
struct device * dev;
```

Arguments

dev

device.

Description

This simply forwards the call to `kobject_get`, though we do take care to provide for the case that we get a NULL pointer passed in.

Name

`put_device` — decrement reference count.

Synopsis

```
void put_device (dev);
```

```
struct device * dev;
```

Arguments

dev

device in question.

Name

`device_del` — delete device from system.

Synopsis


```
void device_del (dev);
```

```
struct device * dev;
```

Arguments

dev

device.

Description

This is the first part of the device unregistration sequence. This removes the device from the lists we control from here, has it removed from the other driver model subsystems it was added to in `device_add`, and removes it from the kobject hierarchy.

NOTE

this should be called manually _iff_ `device_add` was also called manually.

Name

`device_unregister` — unregister device from system.

Synopsis

```
void device_unregister (dev);
```

```
struct device * dev;
```

Arguments

dev

device going away.

Description

We do this in two parts, like we do `device_register`. First, we remove it from all the subsystems with `device_del`, then we decrement the reference count via `put_device`. If that is the final reference count, the device will be cleaned up via `device_release` above. Otherwise, the structure will stick around until the final reference to the device is dropped.

Name

`device_for_each_child` — device child iterator.

Synopsis

```
int device_for_each_child (parent,  
                           data,  
                           fn);  
  
struct device * parent;  
void * data;  
int (* fn(struct device *dev, void *data));
```

Arguments

parent

parent struct device.

data

data for the callback.

fn

function to be called for each device.

Description

Iterate over *parent*'s child devices, and call *fn* for each, passing it *data*.

We check the return of *fn* each time. If it returns anything other than 0, we break out and return that value.

Name

`device_find_child` — device iterator for locating a particular device.

Synopsis

```
struct device * device_find_child (parent,  
                                     data,  
                                     match);  
  
struct device * parent;  
void * data;  
int (* match(struct device *dev, void *data));
```

Arguments

parent

parent struct device

data

Data to pass to match function

match

Callback function to check device

Description

This is similar to the `device_for_each_child` function above, but it returns a reference to a device that is 'found' for later use, as determined by the *match* callback.

The callback should return 0 if the device doesn't match and non-zero if it does. If the callback returns non-zero and a reference to the current device can be obtained, this function will return to the caller and not iterate over any more devices.

Name

`__root_device_register` — allocate and register a root device

Synopsis

```
struct device * __root_device_register (name,  
                                         owner);
```

```
const char *    name;  
struct module * owner;
```

Arguments

name

root device name

owner

owner module of the root device, usually `THIS_MODULE`

Description

This function allocates a root device and registers it using `device_register`. In order to free the

returned device, use `root_device_unregister`.

Root devices are dummy devices which allow other devices to be grouped under `/sys/devices`. Use this function to allocate a root device and then use it as the parent of any device which should appear under `/sys/devices/{name}`

The `/sys/devices/{name}` directory will also contain a 'module' symlink which points to the *owner* directory in `sysfs`.

Note

You probably want to use `root_device_register`.

Name

`root_device_unregister` — unregister and free a root device

Synopsis

```
void root_device_unregister (dev);
```

```
struct device * dev;
```

Arguments

dev

device going away

Description

This function unregisters and cleans up a device that was created by `root_device_register`.

Name

`device_create_vargs` — creates a device and registers it with `sysfs`

Synopsis

```
struct device * device_create_vargs (class,  
                                     parent,  
                                     devt,  
                                     drvdata,  
                                     fmt,  
                                     args);
```

```
struct class *   class;  
struct device * parent;  
dev_t           devt;  
void *          drvdata;  
const char *    fmt;  
va_list         args;
```

Arguments

class

pointer to the struct class that this device should be registered to

parent

pointer to the parent struct device of this new device, if any

devt

the dev_t for the char device to be added

drvdata

the data to be added to the device for callbacks

fmt

string for the device's name

args

va_list for the device's name

Description

This function can be used by char device classes. A struct device will be created in sysfs, registered to the specified class.

A “dev” file will be created, showing the dev_t for the device, if the dev_t is not 0,0. If a pointer to a parent struct device is passed in, the newly created struct device will be a child of that device in sysfs. The pointer to the struct device will be returned from the call. Any further sysfs files that might be required can be created using this pointer.

Note

the struct class passed to this function must have previously been created with a call to `class_create`.

Name

`device_create` — creates a device and registers it with sysfs

Synopsis

```
struct device * device_create (class,
                                parent,
                                devt,
                                drvdata,
                                fmt,
                                ...);
```

```
struct class *   class;
struct device * parent;
dev_t           devt;
void *          drvdata;
const char *    fmt;
               ...;
```

Arguments

class

pointer to the struct class that this device should be registered to

parent

pointer to the parent struct device of this new device, if any

devt

the dev_t for the char device to be added

drvdata

the data to be added to the device for callbacks

fmt

string for the device's name

...

variable arguments

Description

This function can be used by char device classes. A struct device will be created in sysfs, registered to the specified class.

A “dev” file will be created, showing the dev_t for the device, if the dev_t is not 0,0. If a pointer to a parent struct device is passed in, the newly created struct device will be a child of that device in sysfs.

The pointer to the struct device will be returned from the call. Any further sysfs files that might be required can be created using this pointer.

Note

the struct class passed to this function must have previously been created with a call to `class_create`.

Name

`device_destroy` — removes a device that was created with `device_create`

Synopsis

```
void device_destroy (class,  
                    devt);
```

```
struct class * class;  
dev_t         devt;
```

Arguments

class

pointer to the struct class that this device was registered with

devt

the dev_t of the device that was previously registered

Description

This call unregisters and cleans up a device that was created with a call to `device_create`.

Name

`device_rename` — renames a device

Synopsis

```
int device_rename (dev,  
                  new_name);
```

```
struct device * dev;  
char *         new_name;
```

Arguments

dev

the pointer to the struct device to be renamed

new_name

the new name of the device

Description

It is the responsibility of the caller to provide mutual exclusion between two different calls of `device_rename` on the same device to ensure that `new_name` is valid and won't conflict with other devices.

Name

`device_move` — moves a device to a new parent

Synopsis

```
int device_move (dev,  
                 new_parent,  
                 dpm_order);
```

```
struct device * dev;  
struct device * new_parent;  
enum dpm_order dpm_order;
```

Arguments

dev

the pointer to the struct device to be moved

new_parent

the new parent of the device (can be NULL)

dpm_order

how to reorder the `dpm_list`

Name

`__class_create` — create a struct class structure

Synopsis

```
struct class * __class_create (owner,  
                                name,  
                                key);
```

```
struct module *      owner;  
const char *         name;  
struct lock_class_key * key;
```

Arguments

owner

pointer to the module that is to “own” this struct class

name

pointer to a string for the name of this class.

key

the lock_class_key for this class; used by mutex lock debugging

Description

This is used to create a struct class pointer that can then be used in calls to `device_create`.

Note, the pointer created here is to be destroyed when finished by making a call to `class_destroy`.

Name

`class_destroy` — destroys a struct class structure

Synopsis

```
void class_destroy (cls);
```

```
struct class * cls;
```

Arguments

cls

pointer to the struct class that is to be destroyed

Description

Note, the pointer to be destroyed must have been created with a call to `class_create`.

Name

`class_dev_iter_init` — initialize class device iterator

Synopsis

```
void class_dev_iter_init (iter,  
                          class,  
                          start,  
                          type);  
  
struct class_dev_iter *   iter;  
struct class *            class;  
struct device *           start;  
const struct device_type * type;
```

Arguments

iter

class iterator to initialize

class

the class we wanna iterate over

start

the device to start iterating from, if any

type

device_type of the devices to iterate over, NULL for all

Description

Initialize class iterator *iter* such that it iterates over devices of *class*. If *start* is set, the list iteration will start there, otherwise if it is NULL, the iteration starts at the beginning of the list.

Name

`class_dev_iter_next` — iterate to the next device

Synopsis

```
struct device * class_dev_iter_next (iter);  
  
struct class_dev_iter * iter;
```

Arguments

iter

class iterator to proceed

Description

Proceed *iter* to the next device and return it. Returns NULL if iteration is complete.

The returned device is referenced and won't be released till iterator is proceed to the next device or exited. The caller is free to do whatever it wants to do with the device including calling back into class code.

Name

`class_dev_iter_exit` — finish iteration

Synopsis

```
void class_dev_iter_exit (iter);  
  
struct class_dev_iter * iter;
```

Arguments

iter

class iterator to finish

Description

Finish an iteration. Always call this function after iteration is complete whether the iteration ran till the end or not.

Name

`class_for_each_device` — device iterator

Synopsis

```
int class_for_each_device (class,
                          start,
                          data,
                          fn);

struct class *   class;
struct device * start;
void *          data;
int (*          fn(struct device *, void *);
```

Arguments

class

the class we're iterating

start

the device to start with in the list, if any.

data

data for the callback

fn

function to be called for each device

Description

Iterate over *class*'s list of devices, and call *fn* for each, passing it *data*. If *start* is set, the list iteration will start there, otherwise if it is NULL, the iteration starts at the beginning of the list.

We check the return of *fn* each time. If it returns anything other than 0, we break out and return that value.

fn is allowed to do anything including calling back into class code. There's no locking restriction.

Name

`class_find_device` — device iterator for locating a particular device

Synopsis

```

struct device * class_find_device (class,
                                     start,
                                     data,
                                     match);

struct class *   class;
struct device * start;
void *          data;
int (*          match(struct device *, void *));

```

Arguments

class

the class we're iterating

start

Device to begin with

data

data for the match function

match

function to check device

Description

This is similar to the `class_for_each_dev` function above, but it returns a reference to a device that is 'found' for later use, as determined by the *match* callback.

The callback should return 0 if the device doesn't match and non-zero if it does. If the callback returns non-zero, this function will return to the caller and not iterate over any more devices.

Note, you will need to drop the reference with `put_device` after use.

fn is allowed to do anything including calling back into class code. There's no locking restriction.

Name

`class_compat_register` — register a compatibility class

Synopsis

```

struct class_compat * class_compat_register (name);

const char * name;

```

Arguments

name

the name of the class

Description

Compatibility class are meant as a temporary user-space compatibility workaround when converting a family of class devices to a bus devices.

Name

`class_compat_unregister` — unregister a compatibility class

Synopsis

```
void class_compat_unregister (cls);
```

```
struct class_compat * cls;
```

Arguments

cls

the class to unregister

Name

`class_compat_create_link` — create a compatibility class device link to a bus device

Synopsis

```
int class_compat_create_link (cls,  
                               dev,  
                               device_link);
```

```
struct class_compat * cls;  
struct device *      dev;  
struct device *      device_link;
```

Arguments

cls

the compatibility class

dev

the target bus device

device_link

an optional device to which a “device” link should be created

Name

`class_compat_remove_link` — remove a compatibility class device link to a bus device

Synopsis

```
void class_compat_remove_link (cls,  
                                dev,  
                                device_link);
```

```
struct class_compat * cls;  
struct device *      dev;  
struct device *      device_link;
```

Arguments

cls

the compatibility class

dev

the target bus device

device_link

an optional device to which a “device” link was previously created

Name

`request_firmware` — send firmware request and wait for it

Synopsis

```
int request_firmware (firmware_p,  
                      name,
```

```
device);
```

```
const struct firmware ** firmware_p;  
const char *           name;  
struct device *         device;
```

Arguments

firmware_p

pointer to firmware image

name

name of firmware file

device

device for which firmware is being loaded

Description

firmware_p will be used to return a firmware image by the name of *name* for device *device*.

Should be called from user context where sleeping is allowed.

name will be used as \$FIRMWARE in the uevent environment and should be distinctive enough not to be confused with any other firmware image for this or any other device.

Name

release_firmware — release the resource associated with a firmware image

Synopsis

```
void release_firmware (fw);  
  
const struct firmware * fw;
```

Arguments

fw

firmware resource to release

Name

request_firmware_nowait —

Synopsis

```
int request_firmware_nowait (module,  
                             uevent,  
                             name,  
                             device,  
                             context,  
                             cont);  
  
struct module * module;  
int uevent;  
const char * name;  
struct device * device;  
void * context;  
void (* cont)(const struct firmware *fw, void *context);
```

Arguments

module

module requesting the firmware

uevent

sends uevent to copy the firmware image if this flag is non-zero else the firmware copy must be done manually.

name

name of firmware file

device

device for which firmware is being loaded

context

will be passed over to *cont*, and *fw* may be NULL if firmware request fails.

cont

function will be called asynchronously when the firmware request is over.

Description

Asynchronous variant of `request_firmware` for user contexts where it is not possible to sleep for long time. It can't be called in atomic contexts.

Name

`transport_class_register` — register an initial transport class

Synopsis

```
int transport_class_register (tclass);  
  
struct transport_class * tclass;
```

Arguments

tclass

a pointer to the transport class structure to be initialised

Description

The transport class contains an embedded class which is used to identify it. The caller should initialise this structure with zeros and then generic class must have been initialised with the actual transport class unique name. There's a macro `DECLARE_TRANSPORT_CLASS` to do this (declared classes still must be registered).

Returns 0 on success or error on failure.

Name

`transport_class_unregister` — unregister a previously registered class

Synopsis

```
void transport_class_unregister (tclass);  
  
struct transport_class * tclass;
```

Arguments

tclass

The transport class to unregister

Description

Must be called prior to deallocating the memory for the transport class.

Name

`anon_transport_class_register` — register an anonymous class

Synopsis

```
int anon_transport_class_register (atc);  
  
struct anon_transport_class * atc;
```

Arguments

atc

The anon transport class to register

Description

The anonymous transport class contains both a transport class and a container. The idea of an anonymous class is that it never actually has any device attributes associated with it (and thus saves on container storage). So it can only be used for triggering events. Use `prezero` and then use `DECLARE_ANON_TRANSPORT_CLASS` to initialise the anon transport class storage.

Name

`anon_transport_class_unregister` — unregister an anon class

Synopsis

```
void anon_transport_class_unregister (atc);  
  
struct anon_transport_class * atc;
```

Arguments

atc

Pointer to the anon transport class to unregister

Description

Must be called prior to deallocating the memory for the anon transport class.

Name

`transport_setup_device` — declare a new dev for transport class association but don't make it visible yet.

Synopsis

```
void transport_setup_device (dev);  
  
struct device * dev;
```

Arguments

dev

the generic device representing the entity being added

Description

Usually, *dev* represents some component in the HBA system (either the HBA itself or a device remote across the HBA bus). This routine is simply a trigger point to see if any set of transport classes wishes to associate with the added device. This allocates storage for the class device and initialises it, but does not yet add it to the system or add attributes to it (you do this with `transport_add_device`). If you have no need for a separate setup and add operations, use `transport_register_device` (see `transport_class.h`).

Name

`transport_add_device` — declare a new dev for transport class association

Synopsis

```
void transport_add_device (dev);  
  
struct device * dev;
```

Arguments

dev

the generic device representing the entity being added

Description

Usually, *dev* represents some component in the HBA system (either the HBA itself or a device remote across the HBA bus). This routine is simply a trigger point used to add the device to the system and register attributes for it.

Name

transport_configure_device — configure an already set up device

Synopsis

```
void transport_configure_device (dev);  
  
struct device * dev;
```

Arguments

dev

generic device representing device to be configured

Description

The idea of configure is simply to provide a point within the setup process to allow the transport class to extract information from a device after it has been setup. This is used in SCSI because we have to have a setup device to begin using the HBA, but after we send the initial inquiry, we use configure to extract the device parameters. The device need not have been added to be configured.

Name

transport_remove_device — remove the visibility of a device

Synopsis

```
void transport_remove_device (dev);  
  
struct device * dev;
```

Arguments

dev

generic device to remove

Description

This call removes the visibility of the device (to the user from sysfs), but does not destroy it. To eliminate a device entirely you must also call transport_destroy_device. If you don't need to do remove and destroy as separate operations, use transport_unregister_device (see transport_class.h) which

will perform both calls for you.

Name

`transport_destroy_device` — destroy a removed device

Synopsis

```
void transport_destroy_device (dev);
```

```
struct device * dev;
```

Arguments

dev

device to eliminate from the transport class.

Description

This call triggers the elimination of storage associated with the transport classdev. Note: all it really does is relinquish a reference to the classdev. The memory will not be freed until the last reference goes to zero. Note also that the classdev retains a reference count on *dev*, so *dev* too will remain for as long as the transport class device remains around.

Name

`sysdev_driver_register` — Register auxillary driver

Synopsis

```
int sysdev_driver_register (cls,  
                             drv);
```

```
struct sysdev_class * cls;  
struct sysdev_driver * drv;
```

Arguments

cls

Device class driver belongs to.

drv

Driver.

Description

drv is inserted into *cls->drivers* to be called on each operation on devices of that class. The refcount of *cls* is incremented.

Name

`sysdev_driver_unregister` — Remove an auxillary driver.

Synopsis

```
void sysdev_driver_unregister (cls,  
                               drv);
```

```
struct sysdev_class *   cls;  
struct sysdev_driver * drv;
```

Arguments

cls

Class driver belongs to.

drv

Driver.

Name

`sysdev_register` — add a system device to the tree

Synopsis

```
int sysdev_register (sysdev);
```

```
struct sys_device * sysdev;
```

Arguments

sysdev

device in question

Name

`sysdev_suspend` — Suspend all system devices.

Synopsis

```
int sysdev_suspend (state);  
  
pm_message_t state;
```

Arguments

state

Power state to enter.

Description

We perform an almost identical operation as `sysdev_shutdown` above, though calling `->suspend` instead. Interrupts are disabled when this called. Devices are responsible for both saving state and quiescing or powering down the device.

This is only called by the device PM core, so we let them handle all synchronization.

Name

`sysdev_resume` — Bring system devices back to life.

Synopsis

```
int sysdev_resume (void);  
  
void;
```

Arguments

void

no arguments

Description

Similar to `sysdev_suspend`, but we iterate the list forwards to guarantee that parent devices are resumed before their children.

Note

Interrupts are disabled when called.

Name

`platform_get_resource` — get a resource for a device

Synopsis

```
struct resource * platform_get_resource (dev,  
                                           type,  
                                           num);
```

```
struct platform_device * dev;  
unsigned int             type;  
unsigned int             num;
```

Arguments

dev

platform device

type

resource type

num

resource index

Name

`platform_get_irq` — get an IRQ for a device

Synopsis

```
int platform_get_irq (dev,  
                      num);
```

```
struct platform_device * dev;  
unsigned int             num;
```

Arguments

dev

platform device

num

IRQ number index

Name

`platform_get_resource_byname` — get a resource for a device by name

Synopsis

```
struct resource * platform_get_resource_byname (dev,
                                                type,
                                                name);
```

```
struct platform_device * dev;
unsigned int             type;
const char *             name;
```

Arguments

dev

platform device

type

resource type

name

resource name

Name

`platform_get_irq_byname` — get an IRQ for a device

Synopsis

```
int platform_get_irq_byname (dev,
                             name);
```

```
struct platform_device * dev;
const char *             name;
```

Arguments

dev

platform device

name

IRQ name

Name

platform_add_devices — add a numbers of platform devices

Synopsis

```
int platform_add_devices (devs,  
                           num);
```

```
struct platform_device ** devs;  
int num;
```

Arguments

devs

array of platform devices to add

num

number of platform devices in array

Name

platform_device_put —

Synopsis

```
void platform_device_put (pdev);
```

```
struct platform_device * pdev;
```

Arguments

pdev

platform device to free

Description

Free all memory associated with a platform device. This function must *_only_* be externally called in error cases. All other usage is a bug.

Name

platform_device_alloc —

Synopsis

```
struct platform_device * platform_device_alloc (name,  
                                                id);
```

```
const char * name;  
int         id;
```

Arguments

name

base name of the device we're adding

id

instance id

Description

Create a platform device object which can have other objects attached to it, and which will have attached objects freed when it is released.

Name

platform_device_add_resources —

Synopsis

```
int platform_device_add_resources (pdev,  
                                   res,  
                                   num);
```

```
struct platform_device * pdev;  
struct resource *      res;  
unsigned int           num;
```

Arguments

pdev

platform device allocated by `platform_device_alloc` to add resources to

res

set of resources that needs to be allocated for the device

num

number of resources

Description

Add a copy of the resources to the platform device. The memory associated with the resources will be freed when the platform device is released.

Name

`platform_device_add_data` —

Synopsis

```
int platform_device_add_data (pdev,  
                             data,  
                             size);
```

```
struct platform_device * pdev;  
const void *            data;  
size_t                  size;
```

Arguments

pdev

platform device allocated by `platform_device_alloc` to add resources to

data

platform specific data for this platform device

size

size of platform specific data

Description

Add a copy of platform specific data to the platform device's `platform_data` pointer. The memory associated with the platform data will be freed when the platform device is released.

Name

`platform_device_add` — add a platform device to device hierarchy

Synopsis

```
int platform_device_add (pdev);  
  
struct platform_device * pdev;
```

Arguments

pdev

platform device we're adding

Description

This is part 2 of `platform_device_register`, though may be called separately `_iff_` `pdev` was allocated by `platform_device_alloc`.

Name

`platform_device_del` — remove a platform-level device

Synopsis

```
void platform_device_del (pdev);  
  
struct platform_device * pdev;
```

Arguments

pdev

platform device we're removing

Description

Note that this function will also release all memory- and port-based resources owned by the device (*dev->resource*). This function must `_only_` be externally called in error cases. All other usage is a bug.

Name

`platform_device_register` — add a platform-level device

Synopsis

```
int platform_device_register (pdev);
```

```
struct platform_device * pdev;
```

Arguments

pdev

platform device we're adding

Name

`platform_device_unregister` — unregister a platform-level device

Synopsis

```
void platform_device_unregister (pdev);
```

```
struct platform_device * pdev;
```

Arguments

pdev

platform device we're unregistering

Description

Unregistration is done in 2 steps. First we release all resources and remove it from the subsystem, then we drop reference count by calling `platform_device_put`.

Name

platform_device_register_simple —

Synopsis

```
struct platform_device * platform_device_register_simple (name,
                                                         id,
                                                         res,
                                                         num);
```

```
const char *      name;
int               id;
struct resource * res;
unsigned int      num;
```

Arguments

name

base name of the device we're adding

id

instance id

res

set of resources that needs to be allocated for the device

num

number of resources

Description

This function creates a simple platform device that requires minimal resource and memory management. Canned release function freeing memory allocated for the device allows drivers using such devices to be unloaded without waiting for the last reference to the device to be dropped.

This interface is primarily intended for use with legacy drivers which probe hardware directly. Because such drivers create sysfs device nodes themselves, rather than letting system infrastructure handle such device enumeration tasks, they don't fully conform to the Linux driver model. In particular, when such drivers are built as modules, they can't be “hotplugged”.

Name

platform_driver_register —

Synopsis

```
int platform_driver_register (drv);
```

```
struct platform_driver * drv;
```

Arguments

drv

platform driver structure

Name

platform_driver_unregister —

Synopsis

```
void platform_driver_unregister (drv);
```

```
struct platform_driver * drv;
```

Arguments

drv

platform driver structure

Name

platform_driver_probe — register driver for non-hotpluggable device

Synopsis

```
int __init_or_module platform_driver_probe (drv,  
                                              probe);
```

```
struct platform_driver * drv;
```

```
int (* probe)(struct platform_device *);
```

Arguments

drv

platform driver structure

probe

the driver probe routine, probably from an `__init` section

Description

Use this instead of `platform_driver_register` when you know the device is not hotpluggable and has already been registered, and you want to remove its run-once probe infrastructure from memory after the driver has bound to the device.

One typical use for this would be with drivers for controllers integrated into system-on-chip processors, where the controller devices have been configured as part of board setup.

Returns zero if the driver registered and bound to a device, else returns a negative error code and with the driver not registered.

Name

`bus_for_each_dev` — device iterator.

Synopsis

```
int bus_for_each_dev (bus,
                    start,
                    data,
                    fn);

struct bus_type * bus;
struct device * start;
void * data;
int (* fn)(struct device *, void *);
```

Arguments

bus

bus type.

start

device to start iterating from.

data

data for the callback.

fn

function to be called for each device.

Description

Iterate over *bus*'s list of devices, and call *fn* for each, passing it *data*. If *start* is not NULL, we use that device to begin iterating from.

We check the return of *fn* each time. If it returns anything other than 0, we break out and return that value.

NOTE

The device that returns a non-zero value is not retained in any way, nor is its refcount incremented. If the caller needs to retain this data, it should do so, and increment the reference count in the supplied callback.

Name

`bus_find_device` — device iterator for locating a particular device.

Synopsis

```
struct device * bus_find_device (bus,
                                   start,
                                   data,
                                   match);

struct bus_type * bus;
struct device *   start;
void *            data;
int (*            match(struct device *dev, void *data);
```

Arguments

bus

bus type

start

Device to begin with

data

Data to pass to match function

match

Callback function to check device

Description

This is similar to the `bus_for_each_dev` function above, but it returns a reference to a device that is 'found' for later use, as determined by the `match` callback.

The callback should return 0 if the device doesn't match and non-zero if it does. If the callback returns non-zero, this function will return to the caller and not iterate over any more devices.

Name

`bus_find_device_by_name` — device iterator for locating a particular device of a specific name

Synopsis

```
struct device * bus_find_device_by_name (bus,
                                           start,
                                           name);
```

```
struct bus_type * bus;
struct device *   start;
const char *      name;
```

Arguments

bus

bus type

start

Device to begin with

name

name of the device to match

Description

This is similar to the `bus_find_device` function above, but it handles searching by a name automatically, no need to write another `strcmp` matching function.

Name

`bus_for_each_drv` — driver iterator

Synopsis

```
int bus_for_each_drv (bus,
                      start,
                      data,
                      fn);

struct bus_type *      bus;
struct device_driver * start;
void *                 data;
int (*                 fn(struct device_driver *, void *));
```

Arguments

bus

bus we're dealing with.

start

driver to start iterating on.

data

data to pass to the callback.

fn

function to call for each driver.

Description

This is nearly identical to the device iterator above. We iterate over each driver that belongs to *bus*, and call *fn* for each. If *fn* returns anything but 0, we break out and return it. If *start* is not NULL, we use it as the head of the list.

NOTE

we don't return the driver that returns a non-zero value, nor do we leave the reference count incremented for that driver. If the caller needs to know that info, it must set it in the callback. It must also be sure to increment the refcount so it doesn't disappear before returning to the caller.

Name

bus_rescan_devices — rescan devices on the bus for possible drivers

Synopsis

```
int bus_rescan_devices (bus);

struct bus_type * bus;
```

Arguments

bus

the bus to scan.

Description

This function will look for devices on the bus with no driver attached and rescan it against existing drivers to see if it matches any by calling `device_attach` for the unbound devices.

Name

`device_reprobe` — remove driver for a device and probe for a new driver

Synopsis

```
int device_reprobe (dev);

struct device * dev;
```

Arguments

dev

the device to reprobe

Description

This function detaches the attached driver (if any) for the given device and restarts the driver probing process. It is intended to use if probing criteria changed during a devices lifetime and driver attachment should change accordingly.

Name

`bus_register` — register a bus with the system.

Synopsis

```
int bus_register (bus);

struct bus_type * bus;
```

Arguments

bus

bus.

Description

Once we have that, we registered the bus with the kobject infrastructure, then register the children subsystems it has: the devices and drivers that belong to the bus.

Name

bus_unregister — remove a bus from the system

Synopsis

```
void bus_unregister (bus);

struct bus_type * bus;
```

Arguments

bus

bus.

Description

Unregister the child subsystems and the bus itself. Finally, we call bus_put to release the refcount

Device Drivers Power Management

Name

dpm_resume_noirq — Execute “early resume” callbacks for non-sysdev devices.

Synopsis

```
void dpm_resume_noirq (state);  
  
pm_message_t state;
```

Arguments

state

PM transition of the system being carried out.

Description

Call the “noirq” resume handlers for all devices marked as DPM_OFF_IRQ and enable device drivers to receive interrupts.

Name

dpm_resume_end — Execute “resume” callbacks and complete system transition.

Synopsis

```
void dpm_resume_end (state);  
  
pm_message_t state;
```

Arguments

state

PM transition of the system being carried out.

Description

Execute “resume” callbacks for all devices and complete the PM transition of the system.

Name

dpm_suspend_noirq — Execute “late suspend” callbacks for non-sysdev devices.

Synopsis

```
int dpm_suspend_noirq (state);
```



```
pm_message_t  state;
```

Arguments

state

PM transition of the system being carried out.

Description

Prevent device drivers from receiving interrupts and call the “noirq” suspend handlers for all non-sysdev devices.

Name

dpm_suspend_start — Prepare devices for PM transition and suspend them.

Synopsis

```
int dpm_suspend_start (state);
```

```
pm_message_t  state;
```

Arguments

state

PM transition of the system being carried out.

Description

Prepare all non-sysdev devices for system PM transition and execute “suspend” callbacks for them.

Device Drivers ACPI Support

Name

acpi_bus_register_driver — register a driver with the ACPI bus

Synopsis

```
int acpi_bus_register_driver (driver);
```

```
struct acpi_driver * driver;
```

Arguments

driver

driver being registered

Description

Registers a driver with the ACPI bus. Searches the namespace for all devices that match the driver's criteria and binds. Returns zero for success or a negative error status for failure.

Name

`acpi_bus_unregister_driver` — unregisters a driver with the APIC bus

Synopsis

```
void acpi_bus_unregister_driver (driver);
```

```
struct acpi_driver * driver;
```

Arguments

driver

driver to unregister

Description

Unregisters a driver with the ACPI bus. Searches the namespace for all devices that match the driver's criteria and unbinds.

Name

`acpi_bus_driver_init` — add a device to a driver

Synopsis

```
int acpi_bus_driver_init (device,  
                           driver);
```

```
struct acpi_device * device;
```

```
struct acpi_driver * driver;
```

Arguments

device

the device to add and initialize

driver

driver for the device

Description

Used to initialize a device via its device driver. Called whenever a driver is bound to a device. Invokes the driver's add ops.

Device drivers PnP support

Name

`pnp_register_protocol` — adds a pnp protocol to the pnp layer

Synopsis

```
int pnp_register_protocol (protocol);
```

```
struct pnp_protocol * protocol;
```

Arguments

protocol

pointer to the corresponding `pnp_protocol` structure

Ex protocols

ISAPNP, PNPBIOS, etc

Name

`pnp_unregister_protocol` — removes a pnp protocol from the pnp layer

Synopsis

```
void pnp_unregister_protocol (protocol);
```

```
struct pnp_protocol * protocol;
```

Arguments

protocol

pointer to the corresponding pnp_protocol structure

Name

pnp_request_card_device — Searches for a PnP device under the specified card

Synopsis

```
struct pnp_dev * pnp_request_card_device (clink,  
                                           id,  
                                           from);
```

```
struct pnp_card_link * clink;  
const char *          id;  
struct pnp_dev *      from;
```

Arguments

clink

pointer to the card link, cannot be NULL

id

pointer to a PnP ID structure that explains the rules for finding the device

from

Starting place to search from. If NULL it will start from the beginning.

Name

pnp_release_card_device — call this when the driver no longer needs the device

Synopsis

```
void pnp_release_card_device (dev);  
  
struct pnp_dev * dev;
```

Arguments

dev

pointer to the PnP device stucture

Name

`pnnp_register_card_driver` — registers a PnP card driver with the PnP Layer

Synopsis

```
int pnnp_register_card_driver (drv);
```

```
struct pnp_card_driver * drv;
```

Arguments

drv

pointer to the driver to register

Name

`pnnp_unregister_card_driver` — unregisters a PnP card driver from the PnP Layer

Synopsis

```
void pnnp_unregister_card_driver (drv);
```

```
struct pnp_card_driver * drv;
```

Arguments

drv

pointer to the driver to unregister

Name

`pnnp_add_id` — adds an EISA id to the specified device

Synopsis

```
struct pnp_id * pnp_add_id (dev,  
                             id);  
  
struct pnp_dev * dev;  
char * id;
```

Arguments

dev

pointer to the desired device

id

pointer to an EISA id string

Name

`pnp_start_dev` — low-level start of the PnP device

Synopsis

```
int pnp_start_dev (dev);  
  
struct pnp_dev * dev;
```

Arguments

dev

pointer to the desired device

Description

assumes that resources have already been allocated

Name

`pnp_stop_dev` — low-level disable of the PnP device

Synopsis

```
int pnp_stop_dev (dev);
```

```
struct pnp_dev * dev;
```

Arguments

dev

pointer to the desired device

Description

does not free resources

Name

`pnp_activate_dev` — activates a PnP device for use

Synopsis

```
int pnp_activate_dev (dev);
```

```
struct pnp_dev * dev;
```

Arguments

dev

pointer to the desired device

Description

does not validate or set resources so be careful.

Name

`pnp_disable_dev` — disables device

Synopsis

```
int pnp_disable_dev (dev);
```

```
struct pnp_dev * dev;
```

Arguments

dev

pointer to the desired device

Description

inform the correct pnp protocol so that resources can be used by other devices

Name

`pnp_is_active` — Determines if a device is active based on its current resources

Synopsis

```
int pnp_is_active (dev);
```

```
struct pnp_dev * dev;
```

Arguments

dev

pointer to the desired PnP device

Userspace IO devices

Name

`uio_event_notify` — trigger an interrupt event

Synopsis

```
void uio_event_notify (info);
```

```
struct uio_info * info;
```

Arguments

info

UIO device capabilities

Name

`__uio_register_device` — register a new userspace IO device

Synopsis

```
int __uio_register_device (owner,  
                           parent,  
                           info);
```

```
struct module *   owner;  
struct device *   parent;  
struct uio_info * info;
```

Arguments

owner

module that creates the new device

parent

parent device

info

UIO device capabilities

Description

returns zero on success or a negative error code.

Name

`uio_unregister_device` — unregister a industrial IO device

Synopsis

```
void uio_unregister_device (info);
```

```
struct uio_info * info;
```

Arguments

info

UIO device capabilities

Name

struct uio_mem — description of a UIO memory region

Synopsis

```
struct uio_mem {  
    const char * name;  
    unsigned long addr;  
    unsigned long size;  
    int memtype;  
    void __iomem * internal_addr;  
    struct uio_map * map;  
};
```

Members

name

name of the memory region for identification

addr

address of the device's memory

size

size of IO

memtype

type of memory addr points to

internal_addr

ioremap-ped version of addr, for driver internal use

map

for use by the UIO core only.

Name

struct uio_port — description of a UIO port region

Synopsis

```
struct uio_port {  
    const char * name;
```

```

    unsigned long start;
    unsigned long size;
    int porttype;
    struct uio_portio * portio;
};

```

Members

name

name of the port region for identification

start

start of port region

size

size of port region

porttype

type of port (see `UIO_PORT_*` below)

portio

for use by the UIO core only.

Name

struct uio_info — UIO device capabilities

Synopsis

```

struct uio_info {
    struct uio_device * uio_dev;
    const char * name;
    const char * version;
    struct uio_mem mem[MAX_UIO_MAPS];
    struct uio_port port[MAX_UIO_PORT_REGIONS];
    long irq;
    unsigned long irq_flags;
    void * priv;
    irqreturn_t (* handler) (int irq, struct uio_info *dev_info);
    int (* mmap) (struct uio_info *info, struct vm_area_struct *vma);
    int (* open) (struct uio_info *info, struct inode *inode);
    int (* release) (struct uio_info *info, struct inode *inode);
    int (* irqcontrol) (struct uio_info *info, s32 irq_on);
};

```

Members

`uio_dev`

the UIO device this info belongs to

`name`

device name

`version`

device driver version

`mem[MAX_UIO_MAPS]`

list of mappable memory regions, size==0 for end of list

`port[MAX_UIO_PORT_REGIONS]`

list of port regions, size==0 for end of list

`irq`

interrupt number or `UIO_IRQ_CUSTOM`

`irq_flags`

flags for `request_irq`

`priv`

optional private data

`handler`

the device's irq handler

`mmap`

mmap operation for this uio device

`open`

open operation for this uio device

`release`

release operation for this uio device

`irqcontrol`

disable/enable irqs when 0/1 is written to `/dev/uioX`

Chapter 3. Parallel Port Devices

Table of Contents

[parport_yield](#) — relinquish a parallel port temporarily
[parport_yield_blocking](#) — relinquish a parallel port temporarily
[parport_wait_event](#) — wait for an event on a parallel port
[parport_wait_peripheral](#) — wait for status lines to change in 35ms
[parport_negotiate](#) — negotiate an IEEE 1284 mode
[parport_write](#) — write a block of data to a parallel port
[parport_read](#) — read a block of data from a parallel port
[parport_set_timeout](#) — set the inactivity timeout for a device
[parport_register_driver](#) — register a parallel port device driver
[parport_unregister_driver](#) — deregister a parallel port device driver
[parport_get_port](#) — increment a port's reference count
[parport_put_port](#) — decrement a port's reference count
[parport_register_port](#) — register a parallel port
[parport_announce_port](#) — tell device drivers about a parallel port
[parport_remove_port](#) — deregister a parallel port
[parport_register_device](#) — register a device on a parallel port
[parport_unregister_device](#) — deregister a device on a parallel port
[parport_find_number](#) — find a parallel port by number
[parport_find_base](#) — find a parallel port by base address
[parport_claim](#) — claim access to a parallel port device
[parport_claim_or_block](#) — claim access to a parallel port device
[parport_release](#) — give up access to a parallel port device
[parport_open](#) — find a device by canonical device number
[parport_close](#) — close a device opened with `parport_open`

Name

`parport_yield` — relinquish a parallel port temporarily

Synopsis

```
int parport_yield (dev);
```

```
struct pardevice * dev;
```

Arguments

dev

a device on the parallel port

Description

This function relinquishes the port if it would be helpful to other drivers to do so. Afterwards it tries to

reclaim the port using `parport_claim`, and the return value is the same as for `parport_claim`. If it fails, the port is left unclaimed and it is the driver's responsibility to reclaim the port.

The `parport_yield` and `parport_yield_blocking` functions are for marking points in the driver at which other drivers may claim the port and use their devices. Yielding the port is similar to releasing it and reclaiming it, but is more efficient because no action is taken if there are no other devices needing the port. In fact, nothing is done even if there are other devices waiting but the current device is still within its “timeslice”. The default timeslice is half a second, but it can be adjusted via the `/proc` interface.

Name

`parport_yield_blocking` — relinquish a parallel port temporarily

Synopsis

```
int parport_yield_blocking (dev);
```

```
struct pardevice * dev;
```

Arguments

dev

a device on the parallel port

Description

This function relinquishes the port if it would be helpful to other drivers to do so. Afterwards it tries to reclaim the port using `parport_claim_or_block`, and the return value is the same as for `parport_claim_or_block`.

Name

`parport_wait_event` — wait for an event on a parallel port

Synopsis

```
int parport_wait_event (port,  
                        timeout);
```

```
struct parport * port;  
signed long      timeout;
```

Arguments

port

port to wait on

timeout

time to wait (in jiffies)

Description

This function waits for up to *timeout* jiffies for an interrupt to occur on a parallel port. If the port timeout is set to zero, it returns immediately.

If an interrupt occurs before the timeout period elapses, this function returns zero immediately. If it times out, it returns one. An error code less than zero indicates an error (most likely a pending signal), and the calling code should finish what it's doing as soon as it can.

Name

parport_wait_peripheral — wait for status lines to change in 35ms

Synopsis

```
int parport_wait_peripheral (port,  
                             mask,  
                             result);
```

```
struct parport * port;  
unsigned char    mask;  
unsigned char    result;
```

Arguments

port

port to watch

mask

status lines to watch

result

desired values of chosen status lines

Description

This function waits until the masked status lines have the desired values, or until 35ms have elapsed (see IEEE 1284-1994 page 24 to 25 for why this value in particular is hardcoded). The *mask* and *result* parameters are bitmasks, with the bits defined by the constants in `parport.h`: `PARPORT_STATUS_BUSY`, and so on.

The port is polled quickly to start off with, in anticipation of a fast response from the peripheral. This fast polling time is configurable (using `/proc`), and defaults to 500usec. If the timeout for this port (see `parport_set_timeout`) is zero, the fast polling time is 35ms, and this function does not call `schedule`.

If the timeout for this port is non-zero, after the fast polling fails it uses `parport_wait_event` to wait for up to 10ms, waking up if an interrupt occurs.

Name

`parport_negotiate` — negotiate an IEEE 1284 mode

Synopsis

```
int parport_negotiate (port,  
                      mode);
```

```
struct parport * port;  
int             mode;
```

Arguments

port

port to use

mode

mode to negotiate to

Description

Use this to negotiate to a particular IEEE 1284 transfer mode. The *mode* parameter should be one of the constants in `parport.h` starting `IEEE1284_MODE_XXX`.

The return value is 0 if the peripheral has accepted the negotiation to the mode specified, -1 if the peripheral is not IEEE 1284 compliant (or not present), or 1 if the peripheral has rejected the negotiation.

Name

`parport_write` — write a block of data to a parallel port

Synopsis

```
ssize_t parport_write (port,
                      buffer,
                      len);
```

```
struct parport * port;
const void *    buffer;
size_t          len;
```

Arguments

port

port to write to

buffer

data buffer (in kernel space)

len

number of bytes of data to transfer

Description

This will write up to *len* bytes of *buffer* to the port specified, using the IEEE 1284 transfer mode most recently negotiated to (using `parport_negotiate`), as long as that mode supports forward transfers (host to peripheral).

It is the caller's responsibility to ensure that the first *len* bytes of *buffer* are valid.

This function returns the number of bytes transferred (if zero or positive), or else an error code.

Name

`parport_read` — read a block of data from a parallel port

Synopsis

```
ssize_t parport_read (port,
                     buffer,
                     len);
```

```
struct parport * port;
void *          buffer;
size_t          len;
```

Arguments

port

port to read from

buffer

data buffer (in kernel space)

len

number of bytes of data to transfer

Description

This will read up to *len* bytes of *buffer* to the port specified, using the IEEE 1284 transfer mode most recently negotiated to (using `parport_negotiate`), as long as that mode supports reverse transfers (peripheral to host).

It is the caller's responsibility to ensure that the first *len* bytes of *buffer* are available to write to.

This function returns the number of bytes transferred (if zero or positive), or else an error code.

Name

`parport_set_timeout` — set the inactivity timeout for a device

Synopsis

```
long parport_set_timeout (dev,  
                           inactivity);
```

```
struct pardevice * dev;  
long               inactivity;
```

Arguments

dev

device on a port

inactivity

inactivity timeout (in jiffies)

Description

This sets the inactivity timeout for a particular device on a port. This affects functions like `parport_wait_peripheral`. The special value 0 means not to call `schedule` while dealing with this device.

The return value is the previous inactivity timeout.

Any callers of `parport_wait_event` for this device are woken up.

Name

`parport_register_driver` — register a parallel port device driver

Synopsis

```
int parport_register_driver (drv);  
  
struct parport_driver * drv;
```

Arguments

drv
structure describing the driver

Description

This can be called by a parallel port device driver in order to receive notifications about ports being found in the system, as well as ports no longer available.

The *drv* structure is allocated by the caller and must not be deallocated until after calling `parport_unregister_driver`.

The driver's `attach` function may block. The port that `attach` is given will be valid for the duration of the callback, but if the driver wants to take a copy of the pointer it must call `parport_get_port` to do so. Calling `parport_register_device` on that port will do this for you.

The driver's `detach` function may block. The port that `detach` is given will be valid for the duration of the callback, but if the driver wants to take a copy of the pointer it must call `parport_get_port` to do so.

Returns 0 on success. Currently it always succeeds.

Name

`parport_unregister_driver` — deregister a parallel port device driver

Synopsis

```
void parport_unregister_driver (drv);

struct parport_driver * drv;
```

Arguments

drv

structure describing the driver that was given to `parport_register_driver`

Description

This should be called by a parallel port device driver that has registered itself using `parport_register_driver` when it is about to be unloaded.

When it returns, the driver's `attach` routine will no longer be called, and for each port that `attach` was called for, the `detach` routine will have been called.

All the driver's `attach` and `detach` calls are guaranteed to have finished by the time this function returns.

Name

`parport_get_port` — increment a port's reference count

Synopsis

```
struct parport * parport_get_port (port);

struct parport * port;
```

Arguments

port

the port

Description

This ensures that a struct `parport` pointer remains valid until the matching `parport_put_port` call.

Name

`parport_put_port` — decrement a port's reference count

Synopsis

```
void parport_put_port (port);
```

```
struct parport * port;
```

Arguments

port

the port

Description

This should be called once for each call to `parport_get_port`, once the port is no longer needed.

Name

`parport_register_port` — register a parallel port

Synopsis

```
struct parport * parport_register_port (base,  
                                         irq,  
                                         dma,  
                                         ops);
```

```
unsigned long          base;  
int                   irq;  
int                   dma;  
struct parport_operations * ops;
```

Arguments

base

base I/O address

irq

IRQ line

dma

DMA channel

ops

pointer to the port driver's port operations structure

Description

When a parallel port (lowlevel) driver finds a port that should be made available to parallel port device drivers, it should call `parport_register_port`. The *base*, *irq*, and *dma* parameters are for the convenience of port drivers, and for ports where they aren't meaningful needn't be set to anything special. They can be altered afterwards by adjusting the relevant members of the `parport` structure that is returned and represents the port. They should not be tampered with after calling `parport_announce_port`, however.

If there are parallel port device drivers in the system that have registered themselves using `parport_register_driver`, they are not told about the port at this time; that is done by `parport_announce_port`.

The *ops* structure is allocated by the caller, and must not be deallocated before calling `parport_remove_port`.

If there is no memory to allocate a new `parport` structure, this function will return `NULL`.

Name

`parport_announce_port` — tell device drivers about a parallel port

Synopsis

```
void parport_announce_port (port);
```

```
struct parport * port;
```

Arguments

port

parallel port to announce

Description

After a port driver has registered a parallel port with `parport_register_port`, and performed any necessary initialisation or adjustments, it should call `parport_announce_port` in order to notify all device drivers that have called `parport_register_driver`. Their attach functions will be called, with *port* as the parameter.

Name

`parport_remove_port` — deregister a parallel port

Synopsis

```
void parport_remove_port (port);
```

```
struct parport * port;
```

Arguments

port

parallel port to deregister

Description

When a parallel port driver is forcibly unloaded, or a parallel port becomes inaccessible, the port driver must call this function in order to deal with device drivers that still want to use it.

The `parport` structure associated with the port has its operations structure replaced with one containing 'null' operations that return errors or just don't do anything.

Any drivers that have registered themselves using `parport_register_driver` are notified that the port is no longer accessible by having their `detach` routines called with *port* as the parameter.

Name

`parport_register_device` — register a device on a parallel port

Synopsis

```
struct pardevice * parport_register_device (port,
                                              name,
                                              pf,
                                              kf,
                                              irq_func,
                                              flags,
                                              handle);
```

```
struct parport * port;
const char *    name;
int (*          pf(void *);
void (*         kf(void *);
void (*         irq_func(void *);
int            flags;
void *         handle;
```

Arguments

port

port to which the device is attached

name

a name to refer to the device

pf

preemption callback

kf

kick callback (wake-up)

irq_func

interrupt handler

flags

registration flags

handle

data for callback functions

Description

This function, called by parallel port device drivers, declares that a device is connected to a port, and tells the system all it needs to know.

The *name* is allocated by the caller and must not be deallocated until the caller calls *parport_unregister_device* for that device.

The preemption callback function, *pf*, is called when this device driver has claimed access to the port but another device driver wants to use it. It is given *handle* as its parameter, and should return zero if it is willing for the system to release the port to another driver on its behalf. If it wants to keep control of the port it should return non-zero, and no action will be taken. It is good manners for the driver to try to release the port at the earliest opportunity after its preemption callback rejects a preemption attempt. Note that if a preemption callback is happy for preemption to go ahead, there is no need to release the port; it is done automatically. This function may not block, as it may be called from interrupt context. If the device driver does not support preemption, *pf* can be `NULL`.

The wake-up (“kick”) callback function, *kf*, is called when the port is available to be claimed for exclusive access; that is, *parport_claim* is guaranteed to succeed when called from inside the wake-up callback function. If the driver wants to claim the port it should do so; otherwise, it need not take any action. This function may not block, as it may be called from interrupt context. If the device driver does not want to be explicitly invited to claim the port in this way, *kf* can be `NULL`.

The interrupt handler, *irq_func*, is called when an interrupt arrives from the parallel port. Note that if a device driver wants to use interrupts it should use *parport_enable_irq*, and can also check the *irq* member of the *parport* structure representing the port.

The parallel port (lowlevel) driver is the one that has called *request_irq* and whose interrupt handler is called first. This handler does whatever needs to be done to the hardware to acknowledge the interrupt (for PC-style ports there is nothing special to be done). It then tells the IEEE 1284 code about the interrupt, which may involve reacting to an IEEE 1284 event depending on the current IEEE 1284 phase. After this, it calls *irq_func*. Needless to say, *irq_func* will be called from interrupt context, and may not block.

The *PARPORT_DEV_EXCL* flag is for preventing port sharing, and so should only be used when sharing the port with other device drivers is impossible and would lead to incorrect behaviour. Use it sparingly! Normally, *flags* will be zero.

This function returns a pointer to a structure that represents the device on the port, or *NULL* if there is not enough memory to allocate space for that structure.

Name

parport_unregister_device — deregister a device on a parallel port

Synopsis

```
void parport_unregister_device (dev);
```

```
struct pardevice * dev;
```

Arguments

dev

pointer to structure representing device

Description

This undoes the effect of *parport_register_device*.

Name

parport_find_number — find a parallel port by number

Synopsis

```
struct parport * parport_find_number (number);
```

```
int  number;
```

Arguments

number

parallel port number

Description

This returns the parallel port with the specified number, or `NULL` if there is none.

There is an implicit `parport_get_port` done already; to throw away the reference to the port that `parport_find_number` gives you, use `parport_put_port`.

Name

`parport_find_base` — find a parallel port by base address

Synopsis

```
struct parport * parport_find_base (base);
```

```
unsigned long base;
```

Arguments

base

base I/O address

Description

This returns the parallel port with the specified base address, or `NULL` if there is none.

There is an implicit `parport_get_port` done already; to throw away the reference to the port that `parport_find_base` gives you, use `parport_put_port`.

Name

`parport_claim` — claim access to a parallel port device

Synopsis

```
int parport_claim (dev);
```

```
struct pardevice * dev;
```

Arguments

dev

pointer to structure representing a device on the port

Description

This function will not block and so can be used from interrupt context. If `parport_claim` succeeds in claiming access to the port it returns zero and the port is available to use. It may fail (returning non-zero) if the port is in use by another driver and that driver is not willing to relinquish control of the port.

Name

`parport_claim_or_block` — claim access to a parallel port device

Synopsis

```
int parport_claim_or_block (dev);
```

```
struct pardevice * dev;
```

Arguments

dev

pointer to structure representing a device on the port

Description

This behaves like `parport_claim`, but will block if necessary to wait for the port to be free. A return value of 1 indicates that it slept; 0 means that it succeeded without needing to sleep. A negative error code indicates failure.

Name

`parport_release` — give up access to a parallel port device

Synopsis

```
void parport_release (dev);
```

```
struct pardevice * dev;
```

Arguments

dev

pointer to structure representing parallel port device

Description

This function cannot fail, but it should not be called without the port claimed. Similarly, if the port is already claimed you should not try claiming it again.

Name

`parport_open` — find a device by canonical device number

Synopsis

```
struct pardevice * parport_open (devnum,  
                                   name);
```

```
int                devnum;  
const char * name;
```

Arguments

devnum

canonical device number

name

name to associate with the device

Description

This function is similar to `parport_register_device`, except that it locates a device by its number rather than by the port it is attached to.

All parameters except for *devnum* are the same as for `parport_register_device`. The return value is the same as for `parport_register_device`.

Name

`parport_close` — close a device opened with `parport_open`

Synopsis

```
void parport_close (dev);
```

```
struct pardevice * dev;
```

Arguments

dev

device to close

Description

This is to `parport_open` as `parport_unregister_device` is to `parport_register_device`.

Chapter 4. Message-based devices

Table of Contents

[Fusion message devices](#)

[I2O message devices](#)

Fusion message devices

Name

`mpt_register` — Register protocol-specific main callback handler.

Synopsis

```
u8 mpt_register (cbfunc,  
                  dclass);
```

```
MPT_CALLBACK      cbfunc;
```

```
MPT_DRIVER_CLASS  dclass;
```

Arguments

cbfunc

callback function pointer

dclass

Protocol driver's class (`MPT_DRIVER_CLASS` enum value)

Description

This routine is called by a protocol-specific driver (SCSI host, LAN, SCSI target) to register its reply callback routine. Each protocol-specific driver must do this before it will be able to use any IOC resources, such as obtaining request frames.

NOTES

The SCSI protocol driver currently calls this routine thrice in order to register separate callbacks; one for “normal” SCSI IO; one for `MptScsiTaskMgmt` requests; one for Scan/DV requests.

Returns `u8` valued “handle” in the range (and S.O.D. order) `{N,...,7,6,5,...,1}` if successful. A return value of `MPT_MAX_PROTOCOL_DRIVERS` (including zero!) should be considered an error by the caller.

Name

`mpt_deregister` — Deregister a protocol drivers resources.

Synopsis

```
void mpt_deregister (cb_idx);
```

```
u8 cb_idx;
```

Arguments

cb_idx

previously registered callback handle

Description

Each protocol-specific driver should call this routine when its module is unloaded.

Name

`mpt_event_register` — Register protocol-specific event callback handler.

Synopsis

```
int mpt_event_register (cb_idx,  
                        ev_cbfunc);
```

```
u8      cb_idx;  
MPT_EVHANDLER ev_cbfunc;
```

Arguments

cb_idx

previously registered (via `mpt_register`) callback handle

ev_cbfunc

callback function

Description

This routine can be called by one or more protocol-specific drivers if/when they choose to be notified of MPT events.

Returns 0 for success.

Name

`mpt_event_deregister` — Deregister protocol-specific event callback handler

Synopsis

```
void mpt_event_deregister (cb_idx);
```

```
u8  cb_idx;
```

Arguments

cb_idx

previously registered callback handle

Description

Each protocol-specific driver should call this routine when it does not (or can no longer) handle events, or when its module is unloaded.

Name

`mpt_reset_register` — Register protocol-specific IOC reset handler.

Synopsis

```
int mpt_reset_register (cb_idx,  
                        reset_func);
```

```
u8          cb_idx;  
MPT_RESETHANDLER reset_func;
```

Arguments

cb_idx

previously registered (via `mpt_register`) callback handle

reset_func

reset function

Description

This routine can be called by one or more protocol-specific drivers if/when they choose to be notified of IOC resets.

Returns 0 for success.

Name

`mpt_reset_deregister` — Deregister protocol-specific IOC reset handler.

Synopsis

```
void mpt_reset_deregister (cb_idx);
```

```
u8 cb_idx;
```

Arguments

cb_idx

previously registered callback handle

Description

Each protocol-specific driver should call this routine when it does not (or can no longer) handle IOC reset handling, or when its module is unloaded.

Name

`mpt_device_driver_register` — Register device driver hooks

Synopsis

```
int mpt_device_driver_register (dd_cbfunc,
                                cb_idx);
```

```
struct mpt_pci_driver * dd_cbfunc;
u8                      cb_idx;
```

Arguments

dd_cbfunc

driver callbacks struct

cb_idx

MPT protocol driver index

Name

`mpt_device_driver_deregister` — DeRegister device driver hooks

Synopsis

```
void mpt_device_driver_deregister (cb_idx);
```

```
u8 cb_idx;
```

Arguments

cb_idx

MPT protocol driver index

Name

`mpt_get_msg_frame` — Obtain an MPT request frame from the pool

Synopsis

```
MPT_FRAME_HDR* mpt_get_msg_frame (cb_idx,
                                     ioc);
```

```
u8                cb_idx;
MPT_ADAPTER *    ioc;
```

Arguments

cb_idx

Handle of registered MPT protocol driver

ioc

Pointer to MPT adapter structure

Description

Obtain an MPT request frame from the pool (of 1024) that are allocated per MPT adapter.

Returns pointer to a MPT request frame or NULL if none are available or IOC is not active.

Name

`mpt_put_msg_frame` — Send a protocol-specific MPT request frame to an IOC

Synopsis

```
void mpt_put_msg_frame (cb_idx,
                        ioc,
                        mf);
```

```
u8                cb_idx;
MPT_ADAPTER *    ioc;
MPT_FRAME_HDR *  mf;
```

Arguments

cb_idx

Handle of registered MPT protocol driver

ioc

Pointer to MPT adapter structure

mf

Pointer to MPT request frame

Description

This routine posts an MPT request frame to the request post FIFO of a specific MPT adapter.

Name

`mpt_put_msg_frame_hi_pri` — Send a hi-pri protocol-specific MPT request frame

Synopsis

```
void mpt_put_msg_frame_hi_pri (cb_idx,  
                               ioc,  
                               mf);
```

```
u8          cb_idx;  
MPT_ADAPTER * ioc;  
MPT_FRAME_HDR * mf;
```

Arguments

cb_idx

Handle of registered MPT protocol driver

ioc

Pointer to MPT adapter structure

mf

Pointer to MPT request frame

Description

Send a protocol-specific MPT request frame to an IOC using hi-priority request queue.

This routine posts an MPT request frame to the request post FIFO of a specific MPT adapter.

Name

`mpt_free_msg_frame` — Place MPT request frame back on FreeQ.

Synopsis

```
void mpt_free_msg_frame (ioc,
                        mf);
```

```
MPT_ADAPTER *    ioc;
MPT_FRAME_HDR *  mf;
```

Arguments

ioc

Pointer to MPT adapter structure

mf

Pointer to MPT request frame

Description

This routine places a MPT request frame back on the MPT adapter's FreeQ.

Name

`mpt_send_handshake_request` — Send MPT request via doorbell handshake method.

Synopsis

```
int mpt_send_handshake_request (cb_idx,
                                ioc,
                                reqBytes,
                                req,
                                sleepFlag);
```

```
u8      cb_idx;
MPT_ADAPTER * ioc;
int      reqBytes;
u32 *    req;
int      sleepFlag;
```

Arguments

cb_idx

Handle of registered MPT protocol driver

ioc

Pointer to MPT adapter structure

reqBytes

Size of the request in bytes

req

Pointer to MPT request frame

sleepFlag

Use schedule if CAN_SLEEP else use udelay.

Description

This routine is used exclusively to send MptScsiTaskMgmt requests since they are required to be sent via doorbell handshake.

NOTE

It is the callers responsibility to byte-swap fields in the request which are greater than 1 byte in size.

Returns 0 for success, non-zero for failure.

Name

mpt_verify_adapter — Given IOC identifier, set pointer to its adapter structure.

Synopsis

```
int mpt_verify_adapter (iocid,  
                        iocpp);
```

```
int iocid;  
MPT_ADAPTER ** iocpp;
```

Arguments

iocid

IOC unique identifier (integer)

iocpp

Pointer to pointer to IOC adapter

Description

Given a unique IOC identifier, set pointer to the associated MPT adapter structure.

Returns iocid and sets iocpp if iocid is found. Returns -1 if iocid is not found.

Name

mpt_attach — Install a PCI intelligent MPT adapter.

Synopsis

```
int mpt_attach (pdev,
                id);

struct pci_dev *          pdev;
const struct pci_device_id * id;
```

Arguments

pdev

Pointer to pci_dev structure

id

PCI device ID information

Description

This routine performs all the steps necessary to bring the IOC of a MPT adapter to a OPERATIONAL state. This includes registering memory regions, registering the interrupt, and allocating request and reply memory pools.

This routine also pre-fetches the LAN MAC address of a Fibre Channel MPT adapter.

Returns 0 for success, non-zero for failure.

TODO

Add support for polled controllers

Name

`mpt_detach` — Remove a PCI intelligent MPT adapter.

Synopsis

```
void mpt_detach (pdev);
```

```
struct pci_dev * pdev;
```

Arguments

pdev

Pointer to `pci_dev` structure

Name

`mpt_suspend` — Fusion MPT base driver suspend routine.

Synopsis

```
int mpt_suspend (pdev,  
                 state);
```

```
struct pci_dev * pdev;  
pm_message_t    state;
```

Arguments

pdev

Pointer to `pci_dev` structure

state

new state to enter

Name

`mpt_resume` — Fusion MPT base driver resume routine.

Synopsis

```
int mpt_resume (pdev);
```

```
struct pci_dev * pdev;
```

Arguments

pdev

Pointer to pci_dev structure

Name

mpt_GetIocState — Get the current state of a MPT adapter.

Synopsis

```
u32 mpt_GetIocState (ioc,
                    cooked);
```

```
MPT_ADAPTER * ioc;
int          cooked;
```

Arguments

ioc

Pointer to MPT_ADAPTER structure

cooked

Request raw or cooked IOC state

Description

Returns all IOC Doorbell register bits if cooked==0, else just the Doorbell bits in MPI_IOC_STATE_MASK.

Name

mpt_alloc_fw_memory — allocate firmware memory

Synopsis

```
int mpt_alloc_fw_memory (ioc,
                        size);
```

```
MPT_ADAPTER * ioc;
int          size;
```


Arguments

ioc

Pointer to MPT_ADAPTER structure

size

total FW bytes

Description

If memory has already been allocated, the same (cached) value is returned.

Return 0 if successfull, or non-zero for failure

Name

mpt_free_fw_memory — free firmware memory

Synopsis

```
void mpt_free_fw_memory (ioc);
```

```
MPT_ADAPTER * ioc;
```

Arguments

ioc

Pointer to MPT_ADAPTER structure

Description

If alt_img is NULL, delete from ioc structure. Else, delete a secondary image in same format.