

# Linux Filesystems API

This documentation is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

This program is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with this program; if not, write to the Free Software Foundation, Inc., 59 Temple Place, Suite 330, Boston, MA 02111-1307 USA

For more details see the file COPYING in the source distribution of Linux.

---

## Table of Contents

### [1. The Linux VFS](#)

[The Filesystem types](#)

[The Directory Cache](#)

[Inode Handling](#)

[Registration and Superblocks](#)

[File Locks](#)

[Other Functions](#)

### [2. The proc filesystem](#)

[sysctl interface](#)

[proc filesystem interface](#)

### [3. The Filesystem for Exporting Kernel Objects](#)

[sysfs create file](#) — create an attribute file for an object.

[sysfs add file to group](#) — add an attribute file to a pre-existing group.

[sysfs chmod file](#) — update the modified mode value on an object attribute.

[sysfs remove file](#) — remove an object attribute.

[sysfs remove file from group](#) — remove an attribute file from a group.

[sysfs schedule callback](#) — helper to schedule a callback for a kobject

[sysfs create link](#) — create symlink between two objects.

[sysfs remove link](#) — remove symlink in object's directory.

[sysfs create bin file](#) — create binary file for object.

[sysfs remove bin file](#) — remove binary file for object.

### [4. The debugfs filesystem](#)

[debugfs interface](#)

### [5. The Linux Journalling API](#)

[Overview](#)[Details](#)[Summary](#)[Data Types](#)[Structures](#)[Functions](#)[Journal Level](#)[Transaction Level](#)[See also](#)[6. splice API](#)

[splice to pipe](#) — fill passed data into a pipe

[generic file splice read](#) — splice data from file to a pipe

[splice from pipe feed](#) — feed available data from a pipe to a file

[splice from pipe next](#) — wait for some data to splice from

[splice from pipe begin](#) — start splicing from pipe

[splice from pipe end](#) — finish splicing from pipe

[\\_\\_splice from pipe](#) — splice data from a pipe to given actor

[splice from pipe](#) — splice data from a pipe to a file

[generic file splice write](#) — splice data from a pipe to a file

[generic splice sendpage](#) — splice data from a pipe to a socket

[splice direct to actor](#) — splices data directly between two non-pipes

[do splice direct](#) — splices data directly between two files

[7. pipes API](#)

[struct pipe\\_buffer](#) — a linux kernel pipe buffer

[struct pipe\\_inode\\_info](#) — a linux kernel pipe

[generic pipe buf map](#) — virtually map a pipe buffer

[generic pipe buf unmap](#) — unmap a previously mapped pipe buffer

[generic pipe buf steal](#) — attempt to take ownership of a pipe\_buffer

[generic pipe buf get](#) — get a reference to a struct pipe\_buffer

[generic pipe buf confirm](#) — verify contents of the pipe buffer

[generic pipe buf release](#) — put a reference to a struct pipe\_buffer

## Chapter 1. The Linux VFS

### Table of Contents

[The Filesystem types](#)

[The Directory Cache](#)

[Inode Handling](#)

[Registration and Superblocks](#)

[File Locks](#)  
[Other Functions](#)

## The Filesystem types

### Name

enum positive\_aop\_returns — aop return codes with specific semantics

### Synopsis

```
enum positive_aop_returns {
    AOP_WRITEPAGE_ACTIVATE,
    AOP_TRUNCATED_PAGE
};
```

### Constants

AOP\_WRITEPAGE\_ACTIVATE

Notifies the caller that page writeback has completed, that the page is still locked, and should be considered active. The VM uses this hint to return the page to the active list -- it won't be a candidate for writeback again in the near future. Other callers must be careful to unlock the page if they get this return. Returned by `writepage`;

AOP\_TRUNCATED\_PAGE

The AOP method that was handed a locked page has unlocked it and the page might have been truncated. The caller should back up to acquiring a new page and trying again. The aop will be taking reasonable precautions not to livelock. If the caller held a page reference, it should drop it before retrying. Returned by `readpage`.

### Description

`address_space_operation` functions return these large constants to indicate special semantics to the caller. These are much larger than the bytes in a page to allow for functions that return the number of bytes operated on in a given page.

### Name

inc\_nlink — directly increment an inode's link count

### Synopsis

```
void inc_nlink (inode);
```

```
struct inode * inode;
```

## Arguments

*inode*

inode

## Description

This is a low-level filesystem helper to replace any direct filesystem manipulation of `i_nlink`. Currently, it is only here for parity with `dec_nlink`.

---

## Name

`drop_nlink` — directly drop an inode's link count

## Synopsis

```
void drop_nlink (inode);
```

```
struct inode * inode;
```

## Arguments

*inode*

inode

## Description

This is a low-level filesystem helper to replace any direct filesystem manipulation of `i_nlink`. In cases where we are attempting to track writes to the filesystem, a decrement to zero means an imminent write when the file is truncated and actually unlinked on the filesystem.

---

## Name

`clear_nlink` — directly zero an inode's link count

## Synopsis

```
void clear_nlink (inode);
```

```
struct inode * inode;
```

## Arguments

*inode*

inode

## Description

This is a low-level filesystem helper to replace any direct filesystem manipulation of `i_nlink`. See `drop_nlink` for why we care about `i_nlink` hitting zero.

---

## Name

`inode_inc_iversion` — increments `i_version`

## Synopsis

```
void inode_inc_iversion (inode);
```

```
struct inode * inode;
```

## Arguments

*inode*

inode that need to be updated

## Description

Every time the inode is modified, the `i_version` field will be incremented. The filesystem has to be mounted with `i_version` flag

## The Directory Cache

## Name

`d_invalidate` — invalidate a dentry

## Synopsis

```
int d_invalidate (dentry);
```

```
struct dentry * dentry;
```

## Arguments

*dentry*

dentry to invalidate

## Description

Try to invalidate the dentry if it turns out to be possible. If there are other dentries that can be reached through this one we can't delete it and we return -EBUSY. On success we return 0.

no dcache lock.

## Name

shrink\_dcache\_sb — shrink dcache for a superblock

## Synopsis

```
void shrink_dcache_sb (sb);
```

```
struct super_block * sb;
```

## Arguments

*sb*

superblock

## Description

Shrink the dcache for the specified super block. This is used to free the dcache before unmounting a file system

## Name

have\_submounts — check for mounts over a dentry

## Synopsis

```
int have_submounts (parent);
```

```
struct dentry * parent;
```

## Arguments

*parent*

dentry to check.

## Description

Return true if the parent or its subdirectories contain a mount point

---

## Name

shrink\_dcache\_parent — prune dcache

## Synopsis

```
void shrink_dcache_parent (parent);
```

```
struct dentry * parent;
```

## Arguments

*parent*

parent of entries to prune

## Description

Prune the dcache to remove unused children of the parent dentry.

---

## Name

d\_alloc — allocate a dcache entry

## Synopsis

```
struct dentry * d_alloc (parent,  
                        name);
```

```
struct dentry *      parent;  
const struct qstr * name;
```

## Arguments

*parent*

parent of entry to allocate

*name*

qstr of the name

## Description

Allocates a dentry. It returns `NULL` if there is insufficient memory available. On a success the dentry is returned. The name passed in is copied and the copy passed in may be reused after this call.

---

## Name

`d_instantiate` — fill in inode information for a dentry

## Synopsis

```
void d_instantiate (entry,
                  inode);
```

```
struct dentry * entry;
struct inode *  inode;
```

## Arguments

*entry*

dentry to complete

*inode*

inode to attach to this dentry

## Description

Fill in inode information in the entry.

This turns negative dentries into productive full members of society.

NOTE! This assumes that the inode count has been incremented (or otherwise set) by the caller to indicate that it is now in use by the dcache.

---

## Name



`d_alloc_root` — allocate root dentry

## Synopsis

```
struct dentry * d_alloc_root (root_inode);  
  
struct inode * root_inode;
```

## Arguments

*root\_inode*

inode to allocate the root for

## Description

Allocate a root (“/”) dentry for the inode given. The inode is instantiated and returned. `NULL` is returned if there is insufficient memory or the inode passed is `NULL`.

---

## Name

`d_obtain_alias` — find or allocate a dentry for a given inode

## Synopsis

```
struct dentry * d_obtain_alias (inode);  
  
struct inode * inode;
```

## Arguments

*inode*

inode to allocate the dentry for

## Description

Obtain a dentry for an inode resulting from NFS filehandle conversion or similar open by handle operations. The returned dentry may be anonymous, or may have a full name (if the inode was already in the cache).

When called on a directory inode, we must ensure that the inode only ever has one dentry. If a dentry is found, that is returned instead of allocating a new one.

On successful return, the reference to the inode has been transferred to the dentry. In case of an error the reference on the inode is released. To make it easier to use in export operations a `NULL` or `IS_ERR` inode

may be passed in and will be the error will be propagate to the return value, with a `NULL inode` replaced by `ERR_PTR(-ESTALE)`.

---

## Name

`d_splice_alias` — splice a disconnected dentry into the tree if one exists

## Synopsis

```
struct dentry * d_splice_alias (inode,  
                                dentry);
```

```
struct inode *   inode;  
struct dentry * dentry;
```

## Arguments

*inode*

the inode which may have a disconnected dentry

*dentry*

a negative dentry which we want to point to the inode.

## Description

If *inode* is a directory and has a 'disconnected' dentry (i.e. `IS_ROOT` and `DCACHE_DISCONNECTED`), then `d_move` that in place of the given dentry and return it, else simply `d_add` the inode to the dentry and return `NULL`.

This is needed in the lookup routine of any filesystem that is exportable (via `knfsd`) so that we can build dcache paths to directories effectively.

If a dentry was found and moved, then it is returned. Otherwise `NULL` is returned. This matches the expected return value of `->lookup`.

---

## Name

`d_add_ci` — lookup or allocate new dentry with case-exact name

## Synopsis

```
struct dentry * d_add_ci (dentry,  
                           inode,  
                           name);
```

```
struct dentry * dentry;  
struct inode * inode;  
struct qstr * name;
```

## Arguments

*dentry*

the negative dentry that was passed to the parent's lookup func

*inode*

the inode case-insensitive lookup has found

*name*

the case-exact name to be associated with the returned dentry

## Description

This is to avoid filling the dcache with case-insensitive names to the same inode, only the actual correct case is stored in the dcache for case-insensitive filesystems.

For a case-insensitive lookup match and if the the case-exact dentry already exists in in the dcache, use it and return it.

If no entry exists with the exact case name, allocate new dentry with the exact case, and return the spliced entry.

---

## Name

d\_lookup — search for a dentry

## Synopsis

```
struct dentry * d_lookup (parent,  
                           name);
```

```
struct dentry * parent;  
struct qstr * name;
```

## Arguments

*parent*

parent dentry

*name*

qstr of name we wish to find

## Description

Searches the children of the parent dentry for the name in question. If the dentry is found its reference count is incremented and the dentry is returned. The caller must use dput to free the entry when it has finished using it. `NULL` is returned on failure.

`__d_lookup` is `dcache_lock` free. The hash list is protected using RCU. Memory barriers are used while updating and doing lockless traversal. To avoid races with `d_move` while rename is happening, `d_lock` is used.

Overflows in `memcmp`, while `d_move`, are avoided by keeping the length and name pointer in one structure pointed by `d_qstr`.

`rcu_read_lock` and `rcu_read_unlock` are used to disable preemption while lookup is going on.

The dentry unused LRU is not updated even if lookup finds the required dentry in there. It is updated in places such as `prune_dcache`, `shrink_dcache_sb`, `select_parent` and `__dget_locked`. This laziness saves lookup from `dcache_lock` acquisition.

`d_lookup` is protected against the concurrent renames in some unrelated directory using the `seqlock_t` `rename_lock`.

## Name

`d_validate` — verify dentry provided from insecure source

## Synopsis

```
int d_validate (dentry,
                dparent);
```

```
struct dentry * dentry;
struct dentry * dparent;
```

## Arguments

*dentry*

The dentry alleged to be valid child of *dparent*

*dparent*

The parent dentry (known to be valid)

## Description

An insecure source has sent us a dentry, here we verify it and `dget` it. This is used by `ncpfs` in its `readdir` implementation. Zero is returned in the dentry is invalid.

---

## Name

`d_delete` — delete a dentry

## Synopsis

```
void d_delete (dentry);  
  
struct dentry * dentry;
```

## Arguments

*dentry*

The dentry to delete

## Description

Turn the dentry into a negative dentry if possible, otherwise remove it from the hash queues so it can be deleted later

---

## Name

`d_rehash` — add an entry back to the hash

## Synopsis

```
void d_rehash (entry);  
  
struct dentry * entry;
```

## Arguments

*entry*

dentry to add to the hash

## Description

Adds a dentry to the hash according to its name.

---

## Name

`d_move` — move a dentry

## Synopsis

```
void d_move (dentry,  
             target);
```

```
struct dentry * dentry;  
struct dentry * target;
```

## Arguments

*dentry*

entry to move

*target*

new dentry

## Description

Update the dcache to reflect the move of a file name. Negative dcache entries should not be moved in this way.

---

## Name

`d_materialise_unique` — introduce an inode into the tree

## Synopsis

```
struct dentry * d_materialise_unique (dentry,  
                                       inode);
```

```
struct dentry * dentry;  
struct inode * inode;
```

## Arguments

*dentry*

candidate dentry

*inode*

inode to bind to the dentry, to which aliases may be attached

## Description

Introduces an dentry into the tree, substituting an extant disconnected root directory alias in its place if there is one

---

## Name

d\_path — return the path of a dentry

## Synopsis

```
char * d_path (path,  
               buf,  
               buflen);
```

```
const struct path * path;  
char *              buf;  
int                 buflen;
```

## Arguments

*path*

path to report

*buf*

buffer to return value in

*buflen*

buffer length

## Description

Convert a dentry into an ASCII path name. If the entry has been deleted the string “ (deleted)” is appended. Note that this is ambiguous.

Returns a pointer into the buffer or an error code if the path was too long. Note: Callers should use the returned pointer, not the passed in buffer, to use the name! The implementation often starts at an offset into the buffer, and may leave 0 bytes at the start.

“buflen” should be positive.

---

## Name

`find_inode_number` — check for dentry with name

## Synopsis

```
ino_t find_inode_number (dir,  
                           name);
```

```
struct dentry * dir;  
struct qstr * name;
```

## Arguments

*dir*

directory to check

*name*

Name to find.

## Description

Check whether a dentry already exists for the given name, and return the inode number if it has an inode. Otherwise 0 is returned.

This routine is used to post-process directory listings for filesystems using synthetic inode numbers, and is necessary to keep `getcwd` working.

---

## Name

`__d_drop` — drop a dentry

## Synopsis

```
void __d_drop (dentry);
```

```
struct dentry * dentry;
```

## Arguments

*dentry*



dentry to drop

## Description

`d_drop` unhashes the entry from the parent dentry hashes, so that it won't be found through a VFS lookup any more. Note that this is different from deleting the dentry - `d_delete` will try to mark the dentry negative if possible, giving a successful `_negative_` lookup, while `d_drop` will just make the cache lookup fail.

`d_drop` is used mainly for stuff that wants to invalidate a dentry for some reason (NFS timeouts or autofs deletes).

`__d_drop` requires `dentry->d_lock`.

---

## Name

`d_add` — add dentry to hash queues

## Synopsis

```
void d_add (entry,
            inode);
```

```
struct dentry * entry;
struct inode * inode;
```

## Arguments

*entry*

dentry to add

*inode*

The inode to attach to this dentry

## Description

This adds the entry to the hash queues and initializes *inode*. The entry was actually filled in earlier during `d_alloc`.

---

## Name

`d_add_unique` — add dentry to hash queues without aliasing

## Synopsis

```
struct dentry * d_add_unique (entry,  
                               inode);
```

```
struct dentry * entry;  
struct inode *  inode;
```

## Arguments

*entry*

dentry to add

*inode*

The inode to attach to this dentry

## Description

This adds the entry to the hash queues and initializes *inode*. The entry was actually filled in earlier during `d_alloc`.

---

## Name

`dget` — get a reference to a dentry

## Synopsis

```
struct dentry * dget (dentry);
```

```
struct dentry * dentry;
```

## Arguments

*dentry*

dentry to get a reference to

## Description

Given a dentry or NULL pointer increment the reference count if appropriate and return the dentry. A dentry will not be destroyed when it has references. `dget` should never be called for dentries with zero reference counter. For these cases (preferably none, functions in `dcache.c` are sufficient for normal needs and they take necessary precautions) you should hold `dcache_lock` and call `dget_locked` instead of `dget`.

---

## Name

`d_unhashed` — is dentry hashed

## Synopsis

```
int d_unhashed (dentry);  
  
struct dentry * dentry;
```

## Arguments

*dentry*  
entry to check

## Description

Returns true if the dentry passed is not currently hashed.

## Inode Handling

## Name

`inode_init_always` — perform inode structure intialisation

## Synopsis

```
int inode_init_always (sb,  
                        inode);  
  
struct super_block * sb;  
struct inode *      inode;
```

## Arguments

*sb*  
superblock inode belongs to  
  
*inode*  
inode to initialise

## Description

These are initializations that need to be done on every inode allocation as the fields are not initialised by slab allocation.

---

## Name

`clear_inode` — clear an inode

## Synopsis

```
void clear_inode (inode);
```

```
struct inode * inode;
```

## Arguments

*inode*

inode to clear

## Description

This is called by the filesystem to tell us that the inode is no longer useful. We just terminate it with extreme prejudice.

---

## Name

`invalidate_inodes` — discard the inodes on a device

## Synopsis

```
int invalidate_inodes (sb);
```

```
struct super_block * sb;
```

## Arguments

*sb*

superblock

## Description

Discard all of the inodes for a given superblock. If the discard fails because there are busy inodes then a non zero value is returned. If the discard is successful all the inodes have been discarded.

---

## Name

`inode_add_to_lists` — add a new inode to relevant lists

## Synopsis

```
void inode_add_to_lists (sb,  
                        inode);
```

```
struct super_block * sb;  
struct inode *      inode;
```

## Arguments

*sb*

superblock inode belongs to

*inode*

inode to mark in use

## Description

When an inode is allocated it needs to be accounted for, added to the in use list, the owning superblock and the inode hash. This needs to be done under the `inode_lock`, so export a function to do this rather than the inode lock itself. We calculate the hash list to add to here so it is all internal which requires the caller to have already set up the inode number in the inode to add.

---

## Name

`new_inode` — obtain an inode

## Synopsis

```
struct inode * new_inode (sb);
```

```
struct super_block * sb;
```

## Arguments

*sb*

superblock

## Description

Allocates a new inode for given superblock. The default `gfp_mask` for allocations related to `inode->i_mapping` is `GFP_HIGHUSER_MOVABLE`. If `HIGHMEM` pages are unsuitable or it is known that pages allocated for the page cache are not reclaimable or migratable, `mapping_set_gfp_mask` must be called with suitable flags on the newly created inode's mapping

---

## Name

`inunique` — get a unique inode number

## Synopsis

```
ino_t inunique (sb,
                 max_reserved);

struct super_block * sb;
ino_t                max_reserved;
```

## Arguments

*sb*

superblock

*max\_reserved*

highest reserved inode number

## Description

Obtain an inode number that is unique on the system for a given superblock. This is used by file systems that have no natural permanent inode numbering system. An inode number is returned that is higher than the reserved limit but unique.

## BUGS

With a large number of inodes live on the file system this function currently becomes quite slow.

---

## Name

`ilookup5_nowait` — search for an inode in the inode cache

# Synopsis

```
struct inode * ilookup5_nowait (sb,
                                hashval,
                                test,
                                data);

struct super_block * sb;
unsigned long        hashval;
int (*               test(struct inode *, void *);
void *               data;
```

# Arguments

*sb*

super block of file system to search

*hashval*

hash value (usually inode number) to search for

*test*

callback used for comparisons between inodes

*data*

opaque data pointer to pass to *test*

# Description

`ilookup5` uses `ifind` to search for the inode specified by *hashval* and *data* in the inode cache. This is a generalized version of `ilookup` for file systems where the inode number is not sufficient for unique identification of an inode.

If the inode is in the cache, the inode is returned with an incremented reference count. Note, the inode lock is not waited upon so you have to be very careful what you do with the returned inode. You probably should be using `ilookup5` instead.

Otherwise NULL is returned.

Note, *test* is called with the `inode_lock` held, so can't sleep.

---

# Name

`ilookup5` — search for an inode in the inode cache

## Synopsis

```
struct inode * ilookup5 (sb,  
                        hashval,  
                        test,  
                        data);  
  
struct super_block * sb;  
unsigned long        hashval;  
int (*               test(struct inode *, void *);  
void *               data;
```

## Arguments

*sb*

super block of file system to search

*hashval*

hash value (usually inode number) to search for

*test*

callback used for comparisons between inodes

*data*

opaque data pointer to pass to *test*

## Description

`ilookup5` uses `ifind` to search for the inode specified by *hashval* and *data* in the inode cache. This is a generalized version of `ilookup` for file systems where the inode number is not sufficient for unique identification of an inode.

If the inode is in the cache, the inode lock is waited upon and the inode is returned with an incremented reference count.

Otherwise NULL is returned.

Note, *test* is called with the `inode_lock` held, so can't sleep.

---

## Name

`ilookup` — search for an inode in the inode cache

## Synopsis



```
struct inode * ilookup (sb,  
                        ino);  
  
struct super_block * sb;  
unsigned long        ino;
```

## Arguments

*sb*

super block of file system to search

*ino*

inode number to search for

## Description

`ilookup` uses `ifind_fast` to search for the inode *ino* in the inode cache. This is for file systems where the inode number is sufficient for unique identification of an inode.

If the inode is in the cache, the inode is returned with an incremented reference count.

Otherwise NULL is returned.

---

## Name

`iget5_locked` — obtain an inode from a mounted file system

## Synopsis

```
struct inode * iget5_locked (sb,  
                             hashval,  
                             test,  
                             set,  
                             data);  
  
struct super_block * sb;  
unsigned long        hashval;  
int (*               test(struct inode *, void *);  
int (*               set(struct inode *, void *);  
void *               data;
```

## Arguments

*sb*

super block of file system

*hashval*

hash value (usually inode number) to get

*test*

callback used for comparisons between inodes

*set*

callback used to initialize a new struct inode

*data*

opaque data pointer to pass to *test* and *set*

## Description

`iget5_locked` uses `ifind` to search for the inode specified by *hashval* and *data* in the inode cache and if present it is returned with an increased reference count. This is a generalized version of `iget_locked` for file systems where the inode number is not sufficient for unique identification of an inode.

If the inode is not in cache, `get_new_inode` is called to allocate a new inode and this is returned locked, hashed, and with the `I_NEW` flag set. The file system gets to fill it in before unlocking it via `unlock_new_inode`.

Note both *test* and *set* are called with the `inode_lock` held, so can't sleep.

## Name

`iget_locked` — obtain an inode from a mounted file system

## Synopsis

```
struct inode * iget_locked (sb,
                             ino);
```

```
struct super_block * sb;
unsigned long        ino;
```

## Arguments

*sb*

super block of file system

*ino*

inode number to get

## Description

`iget_locked` uses `ifind_fast` to search for the inode specified by `ino` in the inode cache and if present it is returned with an increased reference count. This is for file systems where the inode number is sufficient for unique identification of an inode.

If the inode is not in cache, `get_new_inode_fast` is called to allocate a new inode and this is returned locked, hashed, and with the `I_NEW` flag set. The file system gets to fill it in before unlocking it via `unlock_new_inode`.

## Name

`__insert_inode_hash` — hash an inode

## Synopsis

```
void __insert_inode_hash (inode,
                          hashval);
```

```
struct inode * inode;
unsigned long hashval;
```

## Arguments

*inode*

unhashed inode

*hashval*

unsigned long value used to locate this object in the `inode_hashtable`.

## Description

Add an inode to the inode hash for this superblock.

## Name

`remove_inode_hash` — remove an inode from the hash

## Synopsis

```
void remove_inode_hash (inode);
```

```
struct inode * inode;
```

## Arguments

*inode*

inode to unhash

## Description

Remove an inode from the superblock.

---

## Name

`generic_detach_inode` — remove inode from inode lists

## Synopsis

```
int generic_detach_inode (inode);
```

```
struct inode * inode;
```

## Arguments

*inode*

inode to remove

## Description

Remove inode from inode lists, write it if it's dirty. This is just an internal VFS helper exported for hugetlbfs. Do not use!

Returns 1 if inode should be completely destroyed.

---

## Name

`iput` — put an inode

## Synopsis

```
void iput (inode);
```

```
struct inode * inode;
```

## Arguments

*inode*

inode to put

## Description

Puts an inode, dropping its usage count. If the inode use count hits zero, the inode is then freed and may also be destroyed.

Consequently, `iput` can sleep.

---

## Name

`bmap` — find a block number in a file

## Synopsis

```
sector_t bmap (inode,  
               block);
```

```
struct inode * inode;  
sector_t      block;
```

## Arguments

*inode*

inode of file

*block*

block to find

## Description

Returns the block number on the device holding the inode that is the disk block number for the block of the file requested. That is, asked for block 4 of inode 1 the function will return the disk block relative to the disk start that holds that block of the file.

---

## Name

`touch_atime` — update the access time

## Synopsis

```
void touch_atime (mnt,  
                  dentry);
```

```
struct vfsmount * mnt;  
struct dentry *   dentry;
```

## Arguments

*mnt*

mount the inode is accessed on

*dentry*

dentry accessed

## Description

Update the accessed time on an inode and mark it for writeback. This function automatically handles read only file systems and media, as well as the “noatime” flag and inode specific “noatime” markers.

---

## Name

`file_update_time` — update mtime and ctime time

## Synopsis

```
void file_update_time (file);
```

```
struct file * file;
```

## Arguments

*file*

file accessed

## Description

Update the mtime and ctime members of an inode and mark the inode for writeback. Note that this function is meant exclusively for usage in the file write path of filesystems, and filesystems may choose to explicitly ignore update via this function with the `S_NOCMTIME` inode flag, e.g. for network

filesystem where these timestamps are handled by the server.

---

## Name

`make_bad_inode` — mark an inode bad due to an I/O error

## Synopsis

```
void make_bad_inode (inode);
```

```
struct inode * inode;
```

## Arguments

*inode*

Inode to mark bad

## Description

When an inode cannot be read due to a media or remote network failure this function makes the inode “bad” and causes I/O operations on it to fail from this point on.

---

## Name

`is_bad_inode` — is an inode errored

## Synopsis

```
int is_bad_inode (inode);
```

```
struct inode * inode;
```

## Arguments

*inode*

inode to test

## Description

Returns true if the inode in question has been marked as bad.

---

## Name

`iget_failed` — Mark an under-construction inode as dead and release it

## Synopsis

```
void iget_failed (inode);  
  
struct inode * inode;
```

## Arguments

*inode*

The inode to discard

## Description

Mark an under-construction inode as dead and release it.

## Registration and Superblocks

## Name

`deactivate_super` — drop an active reference to superblock

## Synopsis

```
void deactivate_super (s);  
  
struct super_block * s;
```

## Arguments

*s*

superblock to deactivate

## Description

Drops an active reference to superblock, acquiring a temporary one if there is no active references left. In that case we lock superblock, tell fs driver to shut it down and drop the temporary reference we had just acquired.



## Name

`deactivate_locked_super` — drop an active reference to superblock

## Synopsis

```
void deactivate_locked_super (s);
```

```
struct super_block * s;
```

## Arguments

*s*

superblock to deactivate

## Description

Equivalent of `up_write(s->s_umount); deactivate_super(s);`, except that it does not unlock it until it's all over. As the result, it's safe to use to dispose of new superblock on `->get_sb` failure exits - nobody will see the sucker until it's all over. Equivalent using `up_write + deactivate_super` is safe for that purpose only if superblock is either safe to use or has `NULL ->s_root` when we unlock.

---

## Name

`generic_shutdown_super` — common helper for `->kill_sb`

## Synopsis

```
void generic_shutdown_super (sb);
```

```
struct super_block * sb;
```

## Arguments

*sb*

superblock to kill

## Description

`generic_shutdown_super` does all fs-independent work on superblock shutdown. Typical `->kill_sb` should pick all fs-specific objects that need destruction out of superblock, call `generic_shutdown_super` and release aforementioned objects. Note: dentries and inodes `_are_` taken care of and do not need specific handling.

Upon calling this function, the filesystem may no longer alter or rearrange the set of dentries belonging to this `super_block`, nor may it change the attachments of dentries to inodes.

---

## Name

`sget` — find or create a superblock

## Synopsis

```
struct super_block * sget (type,
                          test,
                          set,
                          data);

struct file_system_type * type;
int (* test(struct super_block *,void *);
int (* set(struct super_block *,void *);
void * data;
```

## Arguments

*type*

filesystem type superblock should belong to

*test*

comparison callback

*set*

setup callback

*data*

argument to each of them

---

## Name

`get_super` — get the superblock of a device

## Synopsis

```
struct super_block * get_super (bdev);

struct block_device * bdev;
```

## Arguments

*bdev*

device to get the superblock for

## Description

Scans the superblock list and finds the superblock of the file system mounted on the device given. `NULL` is returned if no match is found.

## File Locks

### Name

`posix_lock_file` — Apply a POSIX-style lock to a file

### Synopsis

```
int posix_lock_file (filp,  
                     fl,  
                     conflock);
```

```
struct file *      filp;  
struct file_lock * fl;  
struct file_lock * conflock;
```

### Arguments

*filp*

The file to apply the lock to

*fl*

The lock to be applied

*conflock*

Place to return a copy of the conflicting lock, if found.

### Description

Add a POSIX style lock to a file. We merge adjacent & overlapping locks whenever possible. POSIX locks are sorted by owner task, then by starting address

Note that if called with an `FL_EXISTS` argument, the caller may determine whether or not a lock was successfully freed by testing the return value for `-ENOENT`.

## Name

`posix_lock_file_wait` — Apply a POSIX-style lock to a file

## Synopsis

```
int posix_lock_file_wait (filp,
                          fl);
```

```
struct file *      filp;
struct file_lock * fl;
```

## Arguments

*filp*

The file to apply the lock to

*fl*

The lock to be applied

## Description

Add a POSIX style lock to a file. We merge adjacent & overlapping locks whenever possible. POSIX locks are sorted by owner task, then by starting address

## Name

`locks_mandatory_area` — Check for a conflicting lock

## Synopsis

```
int locks_mandatory_area (read_write,
                          inode,
                          filp,
                          offset,
                          count);
```

```
int      read_write;
struct inode * inode;
struct file * filp;
loff_t    offset;
size_t    count;
```

# Arguments

*read\_write*

FLOCK\_VERIFY\_WRITE for exclusive access, FLOCK\_VERIFY\_READ for shared

*inode*

the file to check

*filp*

how the file was opened (if it was)

*offset*

start of area to check

*count*

length of area to check

# Description

Searches the inode's list of locks to find any POSIX locks which conflict. This function is called from `rw_verify_area` and `locks_verify_truncate`.

---

# Name

`__break_lease` — revoke all outstanding leases on file

# Synopsis

```
int __break_lease (inode,
                  mode);
```

```
struct inode * inode;
unsigned int   mode;
```

# Arguments

*inode*

the inode of the file to return

*mode*

the open mode (read or write)

## Description

`break_lease` (inlined for speed) has checked there already is at least some kind of lock (maybe a lease) on this file. Leases are broken on a call to `open` or `truncate`. This function can sleep unless you specified `O_NONBLOCK` to your `open`.

---

## Name

`lease_get_mtime` — get the last modified time of an inode

## Synopsis

```
void lease_get_mtime (inode,
                      time);
```

```
struct inode *      inode;
struct timespec *   time;
```

## Arguments

*inode*

the inode

*time*

pointer to a `timespec` which will contain the last modified time

## Description

This is to force NFS clients to flush their caches for files with exclusive leases. The justification is that if someone has an exclusive lease, then they could be modifying it.

---

## Name

`generic_setlease` — sets a lease on an open file

## Synopsis

```
int generic_setlease (filp,
                      arg,
                      flp);
```

```
struct file *      filp;
long               arg;
```

```
struct file_lock ** flp;
```

## Arguments

*filp*

file pointer

*arg*

type of lease to obtain

*flp*

input - file\_lock to use, output - file\_lock inserted

## Description

The (input) flp->fl\_lmops->fl\_break function is required by break\_lease.

Called with kernel lock held.

---

## Name

flock\_lock\_file\_wait — Apply a FLOCK-style lock to a file

## Synopsis

```
int flock_lock_file_wait (filp,
                          fl);
```

```
struct file *      filp;
struct file_lock * fl;
```

## Arguments

*filp*

The file to apply the lock to

*fl*

The lock to be applied

## Description

Add a FLOCK style lock to a file.

---

## Name

`vfs_test_lock` — test file byte range lock

## Synopsis

```
int vfs_test_lock (filp,  
                  fl);  
  
struct file *      filp;  
struct file_lock * fl;
```

## Arguments

*filp*

The file to test lock for

*fl*

The lock to test; also used to hold result

## Description

Returns -ERRNO on failure. Indicates presence of conflicting lock by setting `conf->fl_type` to something other than `F_UNLCK`.

---

## Name

`vfs_lock_file` — file byte range lock

## Synopsis

```
int vfs_lock_file (filp,  
                  cmd,  
                  fl,  
                  conf);  
  
struct file *      filp;  
unsigned int       cmd;  
struct file_lock * fl;  
struct file_lock * conf;
```

## Arguments



*filp*

The file to apply the lock to

*cmd*

type of locking operation (F\_SETLK, F\_GETLK, etc.)

*fl*

The lock to be applied

*conf*

Place to return a copy of the conflicting lock, if found.

## Description

A caller that doesn't care about the conflicting lock may pass NULL as the final argument.

If the filesystem defines a private `->lock` method, then *conf* will be left unchanged; so a caller that cares should initialize it to some acceptable default.

To avoid blocking kernel daemons, such as `lockd`, that need to acquire POSIX locks, the `->lock` interface may return asynchronously, before the lock has been granted or denied by the underlying filesystem, if (and only if) `fl_grant` is set. Callers expecting `->lock` to return asynchronously will only use `F_SETLK`, not `F_SETLKW`; they will set `FL_SLEEP` if (and only if) the request is for a blocking lock. When `->lock` does return asynchronously, it must return `FILE_LOCK_DEFERRED`, and call `->fl_grant` when the lock request completes. If the request is for non-blocking lock the file system should return `FILE_LOCK_DEFERRED` then try to get the lock and call the callback routine with the result. If the request timed out the callback routine will return a nonzero return code and the file system should release the lock. The file system is also responsible to keep a corresponding posix lock when it grants a lock so the VFS can find out which locks are locally held and do the correct lock cleanup when required. The underlying filesystem must not drop the kernel lock or call `->fl_grant` before returning to the caller with a `FILE_LOCK_DEFERRED` return code.

---

## Name

`posix_unblock_lock` — stop waiting for a file lock

## Synopsis

```
int posix_unblock_lock (filp,
                        waiter);
```

```
struct file *      filp;
struct file_lock * waiter;
```

## Arguments

*filp*

how the file was opened

*waiter*

the lock which was waiting

## Description

lockd needs to block waiting for locks.

---

## Name

vfs\_cancel\_lock — file byte range unblock lock

## Synopsis

```
int vfs_cancel_lock (filp,
                     fl);
```

```
struct file *      filp;
struct file_lock * fl;
```

## Arguments

*filp*

The file to apply the unblock to

*fl*

The lock to be unblocked

## Description

Used by lock managers to cancel blocked requests

---

## Name

lock\_may\_read — checks that the region is free of locks

## Synopsis

```
int lock_may_read (inode,
```

```
    start,  
    len);
```

```
struct inode * inode;  
loff_t        start;  
unsigned long len;
```

## Arguments

*inode*

the inode that is being read

*start*

the first byte to read

*len*

the number of bytes to read

## Description

Emulates Windows locking requirements. Whole-file mandatory locks (share modes) can prohibit a read and byte-range POSIX locks can prohibit a read if they overlap.

N.B. this function is only ever called from knfsd and ownership of locks is never checked.

---

## Name

`lock_may_write` — checks that the region is free of locks

## Synopsis

```
int lock_may_write (inode,  
                    start,  
                    len);
```

```
struct inode * inode;  
loff_t        start;  
unsigned long len;
```

## Arguments

*inode*

the inode that is being written

*start*

the first byte to write

*len*

the number of bytes to write

## Description

Emulates Windows locking requirements. Whole-file mandatory locks (share modes) can prohibit a write and byte-range POSIX locks can prohibit a write if they overlap.

N.B. this function is only ever called from knfsd and ownership of locks is never checked.

---

## Name

`locks_mandatory_locked` — Check for an active lock

## Synopsis

```
int locks_mandatory_locked (inode);
```

```
struct inode * inode;
```

## Arguments

*inode*

the file to check

## Description

Searches the inode's list of locks to find any POSIX locks which conflict. This function is called from `locks_verify_locked` only.

---

## Name

`fcntl_getlease` — Enquire what lease is currently active

## Synopsis

```
int fcntl_getlease (filp);
```

```
struct file * filp;
```

## Arguments

*filp*

the file

## Description

The value returned by this function will be one of (if no lease break is pending):

`F_RDLCK` to indicate a shared lease is held.

`F_WRLCK` to indicate an exclusive lease is held.

`F_UNLCK` to indicate no lease is held.

(if a lease break is pending):

`F_RDLCK` to indicate an exclusive lease needs to be changed to a shared lease (or removed).

`F_UNLCK` to indicate the lease needs to be removed.

## XXX

sfr & willy disagree over whether `F_INPROGRESS` should be returned to userspace.

---

## Name

`fcntl_setlease` — sets a lease on an open file

## Synopsis

```
int fcntl_setlease (fd,
                  filp,
                  arg);
```

```
unsigned int  fd;
struct file * filp;
long         arg;
```

## Arguments

*fd*

open file descriptor

*filp*

file pointer

*arg*

type of lease to obtain

## Description

Call this `fcntl` to establish a lease on the file. Note that you also need to call `F_SETSIG` to receive a signal when the lease is broken.

---

## Name

`sys_flock` — `flock` system call.

## Synopsis

```
long sys_flock (fd,  
                cmd);
```

```
unsigned int fd;  
unsigned int cmd;
```

## Arguments

*fd*

the file descriptor to lock.

*cmd*

the type of lock to apply.

## Description

Apply a `FL_FLOCK` style lock to an open file descriptor. The *cmd* can be one of

`LOCK_SH` -- a shared lock.

`LOCK_EX` -- an exclusive lock.

`LOCK_UN` -- remove an existing lock.

`LOCK_MAND` -- a 'mandatory' flock. This exists to emulate Windows Share Modes.

`LOCK_MAND` can be combined with `LOCK_READ` or `LOCK_WRITE` to allow other processes read and write access respectively.

# Other Functions

## Name

`mpage_readpages` — populate an address space with some pages & start reads against them

## Synopsis

```
int mpage_readpages (mapping,
                    pages,
                    nr_pages,
                    get_block);
```

```
struct address_space * mapping;
struct list_head *    pages;
unsigned              nr_pages;
get_block_t           get_block;
```

## Arguments

*mapping*

the `address_space`

*pages*

The address of a `list_head` which contains the target pages. These pages have their `->index` populated and are otherwise uninitialised. The page at `pages->prev` has the lowest file offset, and reads should be issued in `pages->prev` to `pages->next` order.

*nr\_pages*

The number of pages at `*pages`

*get\_block*

The filesystem's block mapper function.

## Description

This function walks the pages and the blocks within each page, building and emitting large BIOs.

If anything unusual happens, such as:

- encountering a page which has buffers
- encountering a page which has a non-hole after a hole
- encountering a page with non-contiguous blocks

then this code just gives up and calls the `buffer_head`-based read function. It does handle a page which has holes at the end - that is a common case: the end-of-file on `blocksize < PAGE_CACHE_SIZE`

setups.

## BH\_Boundary explanation

There is a problem. The mpage read code assembles several pages, gets all their disk mappings, and then submits them all. That's fine, but obtaining the disk mappings may require I/O. Reads of indirect blocks, for example.

So an mpage read of the first 16 blocks of an ext2 file will cause I/O to be

## submitted in the following order

12 0 1 2 3 4 5 6 7 8 9 10 11 13 14 15 16

because the indirect block has to be read to get the mappings of blocks 13,14,15,16. Obviously, this impacts performance.

So what we do it to allow the filesystem's `get_block` function to set `BH_Boundary` when it maps block 11. `BH_Boundary` says: mapping of the block after this one will require I/O against a block which is probably close to this one. So you should push what I/O you have currently accumulated.

This all causes the disk requests to be issued in the correct order.

## Name

`mpage_writepages` — walk the list of dirty pages of the given address space & `writepage` all of them

## Synopsis

```
int mpage_writepages (mapping,
                     wbc,
                     get_block);

struct address_space *      mapping;
struct writeback_control * wbc;
get_block_t                get_block;
```

## Arguments

*mapping*

address space structure to write

*wbc*

subtract the number of written pages from `*wbc->nr_to_write`

*get\_block*



the filesystem's block mapper function. If this is NULL then use `a_ops->writepage`. Otherwise, go direct-to-BIO.

## Description

This is a library function, which implements the `writepages` `address_space_operation`.

If a page is already under I/O, `generic_writepages` skips it, even if it's dirty. This is desirable behaviour for memory-cleaning writeback, but it is INCORRECT for data-integrity system calls such as `fsync`. `fsync` and `msync` need to guarantee that all the data which was dirty at the time the call was made get new I/O started against them. If `wbc->sync_mode` is `WB_SYNC_ALL` then we were called for data integrity and we must wait for existing IO to complete.

## Name

`generic_permission` — check for access rights on a Posix-like filesystem

## Synopsis

```
int generic_permission (inode,
                       mask,
                       check_acl);

struct inode * inode;
int mask;
int (* check_acl(struct inode *inode, int mask);
```

## Arguments

*inode*

inode to check access rights for

*mask*

right to check for (`MAY_READ`, `MAY_WRITE`, `MAY_EXEC`)

*check\_acl*

optional callback to check for Posix ACLs

## Description

Used to check for read/write/execute permissions on a file. We use “fsuid” for this, letting us set arbitrary permissions for filesystem access without changing the “normal” uids which are used for other things..

## Name

`inode_permission` — check for access rights to a given inode

## Synopsis

```
int inode_permission (inode,  
                      mask);
```

```
struct inode * inode;  
int           mask;
```

## Arguments

*inode*

inode to check permission on

*mask*

right to check for (`MAY_READ`, `MAY_WRITE`, `MAY_EXEC`)

## Description

Used to check for read/write/execute permissions on an inode. We use “fsuid” for this, letting us set arbitrary permissions for filesystem access without changing the “normal” uids which are used for other things.

---

## Name

`file_permission` — check for additional access rights to a given file

## Synopsis

```
int file_permission (file,  
                    mask);
```

```
struct file * file;  
int          mask;
```

## Arguments

*file*

file to check access rights for

*mask*

right to check for (MAY\_READ, MAY\_WRITE, MAY\_EXEC)

## Description

Used to check for read/write/execute permissions on an already opened file.

## Note

Do not use this function in new code. All access checks should be done using `inode_permission`.

---

## Name

`path_get` — get a reference to a path

## Synopsis

```
void path_get (path);
```

```
struct path * path;
```

## Arguments

*path*

path to get the reference to

## Description

Given a path increment the reference count to the dentry and the vfsmount.

---

## Name

`path_put` — put a reference to a path

## Synopsis

```
void path_put (path);
```

```
struct path * path;
```

## Arguments

*path*

path to put the reference to

## Description

Given a path decrement the reference count to the dentry and the vfsmount.

---

## Name

`vfs_path_lookup` — lookup a file path relative to a dentry-vfsmount pair

## Synopsis

```
int vfs_path_lookup (dentry,  
                     mnt,  
                     name,  
                     flags,  
                     nd);
```

```
struct dentry *      dentry;  
struct vfsmount *    mnt;  
const char *         name;  
unsigned int         flags;  
struct nameidata *   nd;
```

## Arguments

*dentry*

pointer to dentry of the base directory

*mnt*

pointer to vfs mount of the base directory

*name*

pointer to file name

*flags*

lookup flags

*nd*

pointer to nameidata

---

## Name

`lookup_one_len` — filesystem helper to lookup single pathname component

## Synopsis

```
struct dentry * lookup_one_len (name,
                                base,
                                len);
```

```
const char *      name;
struct dentry *   base;
int               len;
```

## Arguments

*name*

pathname component to lookup

*base*

base directory to lookup from

*len*

maximum length *len* should be interpreted to

## Description

Note that this routine is purely a helper for filesystem usage and should not be called by generic code. Also note that by using this function the nameidata argument is passed to the filesystem methods and a filesystem using this helper needs to be prepared for that.

---

## Name

`filp_open` — open file and return file pointer

## Synopsis

```
struct file * filp_open (filename,
                        flags,
                        mode);
```

```
const char *   filename;
int            flags;
int            mode;
```

## Arguments

*filename*

path to open

*flags*

open flags as per the open(2) second argument

*mode*

mode for the new file if O\_CREAT is set, else ignored

## Description

This is the helper to open a file from kernelspace if you really have to. But in generally you should not do this, so please move along, nothing to see here..

---

## Name

lookup\_create — lookup a dentry, creating it if it doesn't exist

## Synopsis

```
struct dentry * lookup_create (nd,  
                                is_dir);
```

```
struct nameidata * nd;  
int                is_dir;
```

## Arguments

*nd*

nameidata info

*is\_dir*

directory flag

## Description

Simple function to lookup and return a dentry and create it if it doesn't exist. Is SMP-safe.

Returns with nd->path.dentry->d\_inode->i\_mutex locked.

---

## Name

`sync_mapping_buffers` — write out & wait upon a mapping's “associated” buffers

## Synopsis

```
int sync_mapping_buffers (mapping);  
  
struct address_space * mapping;
```

## Arguments

*mapping*

the mapping which wants those buffers written

## Description

Starts I/O against the buffers at `mapping->private_list`, and waits upon that I/O.

Basically, this is a convenience function for `fsync`. *mapping* is a file or directory which needs those buffers to be written for a successful `fsync`.

---

## Name

`mark_buffer_dirty` — mark a `buffer_head` as needing writeout

## Synopsis

```
void mark_buffer_dirty (bh);  
  
struct buffer_head * bh;
```

## Arguments

*bh*

the `buffer_head` to mark dirty

## Description

`mark_buffer_dirty` will set the dirty bit against the buffer, then set its backing page dirty, then tag the page as dirty in its `address_space`'s radix tree and then attach the `address_space`'s inode to its superblock's dirty inode list.

`mark_buffer_dirty` is atomic. It takes `bh->b_page->mapping->private_lock`, `mapping->tree_lock` and the global `inode_lock`.

---

## Name

`__bread` — reads a specified block and returns the `bh`

## Synopsis

```
struct buffer_head * __bread (bdev,
                             block,
                             size);
```

```
struct block_device * bdev;
sector_t             block;
unsigned             size;
```

## Arguments

*bdev*

the `block_device` to read from

*block*

number of block

*size*

size (in bytes) to read

## Description

Reads a specified block, and returns buffer head that contains it. It returns `NULL` if the block was unreadable.

---

## Name

`block_invalidatepage` — invalidate part of all of a buffer-backed page

## Synopsis

```
void block_invalidatepage (page,
                          offset);
```

```
struct page * page;
```



```
unsigned long offset;
```

## Arguments

*page*

the page which is affected

*offset*

the index of the truncation point

## Description

`block_invalidatepage` is called when all or part of the page has become invalidated by a truncate operation.

`block_invalidatepage` does not have to release all buffers, but it must ensure that no dirty buffer is left outside *offset* and that no I/O is underway against any of the blocks which are outside the truncation point. Because the caller is about to free (and possibly reuse) those blocks on-disk.

## Name

`ll_rw_block` — level access to block devices (DEPRECATED)

## Synopsis

```
void ll_rw_block (rw,
                  nr,
                  bhs[]);

int          rw;
int          nr;
struct buffer_head * bhs[];
```

## Arguments

*rw*

whether to READ or WRITE or SWRITE or maybe READA (readahead)

*nr*

number of struct `buffer_heads` in the array

*bhs[]*

array of pointers to struct `buffer_head`

## Description

`ll_rw_block` takes an array of pointers to `struct buffer_heads`, and requests an I/O operation on them, either a `READ` or a `WRITE`. The third `SWRITE` is like `WRITE` only we make sure that the `*current*` data in buffers are sent to disk. The fourth `READA` option is described in the documentation for `generic_make_request` which `ll_rw_block` calls.

This function drops any buffer that it cannot get a lock on (with the `BH_Lock` state bit) unless `SWRITE` is required, any buffer that appears to be clean when doing a write request, and any buffer that appears to be up-to-date when doing read request. Further it marks as clean buffers that are processed for writing (the buffer cache won't assume that they are actually clean until the buffer gets unlocked).

`ll_rw_block` sets `b_end_io` to simple completion handler that marks the buffer up-to-date (if appropriate), unlocks the buffer and wakes any waiters.

All of the buffers must be for the same device, and must also be a multiple of the current approved size for the device.

---

## Name

`bh_uptodate_or_lock` — Test whether the buffer is uptodate

## Synopsis

```
int bh_uptodate_or_lock (bh);
```

```
struct buffer_head * bh;
```

## Arguments

*bh*

`struct buffer_head`

## Description

Return true if the buffer is up-to-date and false, with the buffer locked, if not.

---

## Name

`bh_submit_read` — Submit a locked buffer for reading

## Synopsis

```
int bh_submit_read (bh);

struct buffer_head * bh;
```

## Arguments

*bh*

struct buffer\_head

## Description

Returns zero on success and -EIO on error.

---

## Name

bio\_alloc\_bioset — allocate a bio for I/O

## Synopsis

```
struct bio * bio_alloc_bioset (gfp_mask,
                                nr_iovecs,
                                bs);

gfp_t          gfp_mask;
int            nr_iovecs;
struct bio_set * bs;
```

## Arguments

*gfp\_mask*

the GFP\_ mask given to the slab allocator

*nr\_iovecs*

number of iovecs to pre-allocate

*bs*

the bio\_set to allocate from. If NULL, just use kmalloc

## Description

bio\_alloc\_bioset will first try its own mempool to satisfy the allocation. If \_\_GFP\_WAIT is set then we will block on the internal pool waiting for a struct bio to become free. If a NULL *bs* is passed in, we will fall back to just using *kmalloc* to allocate the required memory.

Note that the caller must set `->bi_destructor` on succesful return of a bio, to do the appropriate freeing of the bio once the reference count drops to zero.

---

## Name

`bio_alloc` — allocate a new bio, memory pool backed

## Synopsis

```
struct bio * bio_alloc (gfp_mask,
                        nr_iovecs);
```

```
gfp_t  gfp_mask;
int     nr_iovecs;
```

## Arguments

*gfp\_mask*

allocation mask to use

*nr\_iovecs*

number of iovecs

## Description

Allocate a new bio with *nr\_iovecs* bvecs. If *gfp\_mask* contains `__GFP_WAIT`, the allocation is guaranteed to succeed.

## RETURNS

Pointer to new bio on success, NULL on failure.

---

## Name

`bio_kmalloc` — allocate a bio for I/O

## Synopsis

```
struct bio * bio_kmalloc (gfp_mask,
                        nr_iovecs);
```

```
gfp_t  gfp_mask;
int     nr_iovecs;
```

## Arguments

*gfp\_mask*

the GFP\_ mask given to the slab allocator

*nr\_iovecs*

number of iovecs to pre-allocate

## Description

`bio_alloc` will allocate a bio and associated `bio_vec` array that can hold at least *nr\_iovecs* entries. Allocations will be done from the `fs_bio_set`. Also see *bio\_alloc\_bioset*.

If `__GFP_WAIT` is set, then `bio_alloc` will always be able to allocate a bio. This is due to the mempool guarantees. To make this work, callers must never allocate more than 1 bio at a time from this pool. Callers that need to allocate more than 1 bio must always submit the previously allocated bio for IO before attempting to allocate a new one. Failure to do so can cause livelocks under memory pressure.

---

## Name

`bio_put` — release a reference to a bio

## Synopsis

```
void bio_put (bio);
```

```
struct bio * bio;
```

## Arguments

*bio*

bio to release reference to

## Description

Put a reference to a struct bio, either one you have gotten with `bio_alloc` or `bio_get`. The last put of a bio will free it.

---

## Name

`__bio_clone` — clone a bio

## Synopsis

```
void __bio_clone (bio,  
                 bio_src);
```

```
struct bio * bio;  
struct bio * bio_src;
```

## Arguments

*bio*

destination bio

*bio\_src*

bio to clone

## Description

Clone a bio. Caller will own the returned bio, but not the actual data it points to. Reference count of returned bio will be one.

---

## Name

bio\_clone — clone a bio

## Synopsis

```
struct bio * bio_clone (bio,  
                        gfp_mask);
```

```
struct bio * bio;  
gfp_t gfp_mask;
```

## Arguments

*bio*

bio to clone

*gfp\_mask*

allocation priority

## Description

Like `__bio_clone`, only also allocates the returned bio

---

## Name

`bio_get_nr_vecs` — return approx number of vecs

## Synopsis

```
int bio_get_nr_vecs (bdev);

struct block_device * bdev;
```

## Arguments

*bdev*

I/O target

## Description

Return the approximate number of pages we can send to this target. There's no guarantee that you will be able to fit this number of pages into a bio, it does not account for dynamic restrictions that vary on offset.

---

## Name

`bio_add_pc_page` — attempt to add page to bio

## Synopsis

```
int bio_add_pc_page (q,
                    bio,
                    page,
                    len,
                    offset);

struct request_queue * q;
struct bio *          bio;
struct page *         page;
unsigned int          len;
unsigned int          offset;
```

## Arguments

*q*

the target queue

*bio*

destination bio

*page*

page to add

*len*

vec entry length

*offset*

vec entry offset

## Description

Attempt to add a page to the `bio_vec` maplist. This can fail for a number of reasons, such as the bio being full or target block device limitations. The target block device must allow bios smaller than `PAGE_SIZE`, so it is always possible to add a single page to an empty bio. This should only be used by `REQ_PC` bios.

---

## Name

`bio_add_page` — attempt to add page to bio

## Synopsis

```
int bio_add_page (bio,  
                  page,  
                  len,  
                  offset);
```

```
struct bio *   bio;  
struct page *  page;  
unsigned int   len;  
unsigned int   offset;
```

## Arguments

*bio*

destination bio



*page*

page to add

*len*

vec entry length

*offset*

vec entry offset

## Description

Attempt to add a page to the `bio_vec` maplist. This can fail for a number of reasons, such as the bio being full or target block device limitations. The target block device must allow bio's smaller than `PAGE_SIZE`, so it is always possible to add a single page to an empty bio.

---

## Name

`bio_uncopy_user` — finish previously mapped bio

## Synopsis

```
int bio_uncopy_user (bio);
```

```
struct bio * bio;
```

## Arguments

*bio*

bio being terminated

## Description

Free pages allocated from `bio_copy_user` and write back data to user space in case of a read.

---

## Name

`bio_copy_user` — copy user data to bio

## Synopsis

```
struct bio * bio_copy_user (q,
```

```

    map_data,
    uaddr,
    len,
    write_to_vm,
    gfp_mask);

```

```

struct request_queue * q;
struct rq_map_data * map_data;
unsigned long uaddr;
unsigned int len;
int write_to_vm;
gfp_t gfp_mask;

```

## Arguments

*q*

destination block queue

*map\_data*

pointer to the `rq_map_data` holding pages (if necessary)

*uaddr*

start of user address

*len*

length in bytes

*write\_to\_vm*

bool indicating writing to pages or not

*gfp\_mask*

memory allocation flags

## Description

Prepares and returns a bio for indirect user io, bouncing data to/from kernel pages as necessary. Must be paired with call `bio_uncopy_user` on io completion.

---

## Name

`bio_map_user` — map user address into bio

## Synopsis

```
struct bio * bio_map_user (q,  
                             bdev,  
                             uaddr,  
                             len,  
                             write_to_vm,  
                             gfp_mask);  
  
struct request_queue * q;  
struct block_device * bdev;  
unsigned long uaddr;  
unsigned int len;  
int write_to_vm;  
gfp_t gfp_mask;
```

## Arguments

*q*

the struct request\_queue for the bio

*bdev*

destination block device

*uaddr*

start of user address

*len*

length in bytes

*write\_to\_vm*

bool indicating writing to pages or not

*gfp\_mask*

memory allocation flags

## Description

Map the user space address into a bio suitable for io to a block device. Returns an error pointer in case of error.

---

## Name

bio\_unmap\_user — unmap a bio

## Synopsis

```
void bio_unmap_user (bio);
```

```
struct bio * bio;
```

## Arguments

*bio*

the bio being unmapped

## Description

Unmap a bio previously mapped by `bio_map_user`. Must be called with a process context.

`bio_unmap_user` may sleep.

---

## Name

`bio_map_kern` — map kernel address into bio

## Synopsis

```
struct bio * bio_map_kern (q,  
                           data,  
                           len,  
                           gfp_mask);
```

```
struct request_queue * q;  
void * data;  
unsigned int len;  
gfp_t gfp_mask;
```

## Arguments

*q*

the struct `request_queue` for the bio

*data*

pointer to buffer to map

*len*

length in bytes

*gfp\_mask*

allocation flags for bio allocation

## Description

Map the kernel address into a bio suitable for io to a block device. Returns an error pointer in case of error.

---

## Name

bio\_copy\_kern — copy kernel address into bio

## Synopsis

```
struct bio * bio_copy_kern (q,  
                             data,  
                             len,  
                             gfp_mask,  
                             reading);
```

```
struct request_queue * q;  
void * data;  
unsigned int len;  
gfp_t gfp_mask;  
int reading;
```

## Arguments

*q*

the struct request\_queue for the bio

*data*

pointer to buffer to copy

*len*

length in bytes

*gfp\_mask*

allocation flags for bio and page allocation

*reading*

data direction is READ

## Description

copy the kernel address into a bio suitable for io to a block device. Returns an error pointer in case of error.

---

## Name

`bio_endio` — end I/O on a bio

## Synopsis

```
void bio_endio (bio,
                error);
```

```
struct bio * bio;
int          error;
```

## Arguments

*bio*

bio

*error*

error, if any

## Description

`bio_endio` will end I/O on the whole bio. `bio_endio` is the preferred way to end I/O on a bio, it takes care of clearing `BIO_UPTODATE` on error. *error* is 0 on success, and one of the established -Exxxx (-EIO, for instance) error values in case something went wrong. Noone should call `bi_end_io` directly on a bio unless they own it and thus know that it has an `end_io` function.

---

## Name

`bio_sector_offset` — Find hardware sector offset in bio

## Synopsis

```
sector_t bio_sector_offset (bio,
                             index,
                             offset);
```

```
struct bio * bio;
```

```
unsigned short  index;  
unsigned int    offset;
```

## Arguments

*bio*

bio to inspect

*index*

bio\_vec index

*offset*

offset in bv\_page

## Description

Return the number of hardware sectors between beginning of bio and an end point indicated by a bio\_vec index and an offset within that vector's page.

---

## Name

biovec\_create — Create a bio\_vec

## Synopsis

```
struct bio_vec * biovec_create (pool_size,  
                                front_pad);
```

```
unsigned int  pool_size;  
unsigned int  front_pad;
```

## Arguments

*pool\_size*

Number of bio and bio\_vecs to cache in the mempool

*front\_pad*

Number of bytes to allocate in front of the returned bio

## Description

Set up a bio\_vec to be used with *bio\_alloc\_biovec*. Allows the caller to ask for a number of bytes to be

allocated in front of the bio. Front pad allocation is useful for embedding the bio inside another structure, to avoid allocating extra data to go with the bio. Note that the bio must be embedded at the END of that structure always, or things will break badly.

---

## Name

`seq_open` — initialize sequential file

## Synopsis

```
int seq_open (file,
              op);

struct file *                               file;
const struct seq_operations * op;
```

## Arguments

*file*

file we initialize

*op*

method table describing the sequence

## Description

`seq_open` sets *file*, associating it with a sequence described by *op*. *op*->`start` sets the iterator up and returns the first element of sequence. *op*->`stop` shuts it down. *op*->`next` returns the next element of sequence. *op*->`show` prints element into the buffer. In case of error *op*->`start` and *op*->`next` return `ERR_PTR(error)`. In the end of sequence they return `NULL`. *op*->`show` returns 0 in case of success and negative number in case of error. Returning `SEQ_SKIP` means “discard this element and move on”.

---

## Name

`seq_read` — *op*->`read` method for sequential files.

## Synopsis

```
ssize_t seq_read (file,
                  buf,
                  size,
                  ppos);
```



```
struct file *  file;  
char __user *  buf;  
size_t        size;  
loff_t *      ppos;
```

## Arguments

*file*

the file to read from

*buf*

the buffer to read to

*size*

the maximum number of bytes to read

*ppos*

the current position in the file

## Description

Ready-made ->f\_op->read

---

## Name

seq\_lseek — ->llseek method for sequential files.

## Synopsis

```
loff_t seq_lseek (file,  
                  offset,  
                  origin);
```

```
struct file *  file;  
loff_t        offset;  
int           origin;
```

## Arguments

*file*

the file in question

*offset*

new position

*origin*

0 for absolute, 1 for relative position

## Description

Ready-made ->f\_op->llseek

---

## Name

seq\_release — free the structures associated with sequential file.

## Synopsis

```
int seq_release (inode,
                 file);
```

```
struct inode * inode;
struct file * file;
```

## Arguments

*inode*

file->f\_path.dentry->d\_inode

*file*

file in question

## Description

Frees the structures associated with sequential file; can be used as ->f\_op->release if you don't have private data to destroy.

---

## Name

seq\_escape — print string into buffer, escaping some characters

## Synopsis

```
int seq_escape (m,
                s,
```

```
    esc);
```

```
struct seq_file * m;
const char *      s;
const char *      esc;
```

## Arguments

*m*

target buffer

*s*

string

*esc*

set of characters that need escaping

## Description

Puts string into buffer, replacing each occurrence of character from *esc* with usual octal escape. Returns 0 in case of success, -1 - in case of overflow.

---

## Name

`mangle_path` — mangle and copy path to buffer beginning

## Synopsis

```
char * mangle_path (s,
                    p,
                    esc);
```

```
char * s;
char * p;
char * esc;
```

## Arguments

*s*

buffer start

*p*

beginning of path in above buffer

*esc*

set of characters that need escaping

## Description

Copy the path from *p* to *s*, replacing each occurrence of character from *esc* with usual octal escape. Returns pointer past last written character in *s*, or NULL in case of failure.

---

## Name

seq\_path — seq\_file interface to print a pathname

## Synopsis

```
int seq_path (m,
              path,
              esc);

struct seq_file * m;
struct path * path;
char * esc;
```

## Arguments

*m*

the seq\_file handle

*path*

the struct path to print

*esc*

set of characters to escape in the output

## Description

return the absolute path of 'path', as represented by the dentry / mnt pair in the path parameter.

---

## Name

seq\_write — write arbitrary data to buffer

## Synopsis

```
int seq_write (seq,
               data,
               len);

struct seq_file * seq;
const void *      data;
size_t           len;
```

## Arguments

*seq*

seq\_file identifying the buffer to which data should be written

*data*

data address

*len*

number of bytes

## Description

Return 0 on success, non-zero otherwise.

---

## Name

register\_filesystem — register a new filesystem

## Synopsis

```
int register_filesystem (fs);

struct file_system_type * fs;
```

## Arguments

*fs*

the file system structure

## Description

Adds the file system passed to the list of file systems the kernel is aware of for mount and other syscalls.

Returns 0 on success, or a negative errno code on an error.

The struct `file_system_type` that is passed is linked into the kernel structures and must not be freed until the file system has been unregistered.

---

## Name

`unregister_filesystem` — unregister a file system

## Synopsis

```
int unregister_filesystem (fs);
```

```
struct file_system_type * fs;
```

## Arguments

*fs*

filesystem to unregister

## Description

Remove a file system that was previously successfully registered with the kernel. An error is returned if the file system is not found. Zero is returned on a success.

Once this function has returned the struct `file_system_type` structure may be freed or reused.

---

## Name

`__mark_inode_dirty` — internal function

## Synopsis

```
void __mark_inode_dirty (inode,  
                        flags);
```

```
struct inode * inode;  
int           flags;
```

## Arguments

*inode*

inode to mark

*flags*

what kind of dirty (i.e. I\_DIRTY\_SYNC) Mark an inode as dirty. Callers should use mark\_inode\_dirty or mark\_inode\_dirty\_sync.

## Description

Put the inode on the super block's dirty list.

CAREFUL! We mark it dirty unconditionally, but move it onto the dirty list only if it is hashed or if it refers to a blockdev. If it was not hashed, it will never be added to the dirty list even if it is later hashed, as it will have been marked dirty already.

In short, make sure you hash any inodes *\_before\_* you start marking them dirty.

This function *\*must\** be atomic for the I\_DIRTY\_PAGES case - set\_page\_dirty is called under spinlock in several places.

Note that for blockdevs, inode->dirtyed\_when represents the dirtying time of the block-special inode (/dev/hda1) itself. And the ->dirtyed\_when field of the kernel-internal blockdev inode represents the dirtying time of the blockdev's pages. This is why for I\_DIRTY\_PAGES we always use page->mapping->host, so the page-dirtying time is recorded in the internal blockdev inode.

---

## Name

writeback\_inodes\_sb — writeback dirty inodes from given super\_block

## Synopsis

```
void writeback_inodes_sb (sb);
```

```
struct super_block * sb;
```

## Arguments

*sb*

the superblock

## Description

Start writeback on some inodes on this super\_block. No guarantees are made on how many (if any) will be written, and this function does not wait for IO completion of submitted IO. The number of pages submitted is returned.

## Name

sync\_inodes\_sb — sync sb inode pages

## Synopsis

```
void sync_inodes_sb (sb);  
  
struct super_block * sb;
```

## Arguments

*sb*  
the superblock

## Description

This function writes and waits on any dirty inode belonging to this super\_block. The number of pages synced is returned.

---

## Name

write\_inode\_now — write an inode to disk

## Synopsis

```
int write_inode_now (inode,  
                     sync);  
  
struct inode * inode;  
int           sync;
```

## Arguments

*inode*  
inode to write to disk

*sync*  
whether the write should be synchronous or not

## Description



This function commits an inode to disk immediately if it is dirty. This is primarily needed by knfsd.

The caller must either have a ref on the inode or must have set I\_WILL\_FREE.

---

## Name

`sync_inode` — write an inode and its pages to disk.

## Synopsis

```
int sync_inode (inode,
                wbc);
```

```
struct inode *          inode;
struct writeback_control * wbc;
```

## Arguments

*inode*

the inode to sync

*wbc*

controls the writeback mode

## Description

`sync_inode` will write an inode and its pages to disk. It will also correctly update the inode on its superblock's dirty inode lists and will update `inode->i_state`.

The caller must have a ref on the inode.

---

## Name

`freeze_bdev` — - lock a filesystem and force it into a consistent state

## Synopsis

```
struct super_block * freeze_bdev (bdev);
```

```
struct block_device * bdev;
```

## Arguments

*bdev*

blockdevice to lock

## Description

If a superblock is found on this device, we take the `s_umount` semaphore on it to make sure nobody unmounts until the snapshot creation is done. The reference counter (`bd_fsfreeze_count`) guarantees that only the last unfreeze process can unfreeze the frozen filesystem actually when multiple freeze requests arrive simultaneously. It counts up in `freeze_bdev` and count down in `thaw_bdev`. When it becomes 0, `thaw_bdev` will unfreeze actually.

---

## Name

`thaw_bdev` — - unlock filesystem

## Synopsis

```
int thaw_bdev (bdev,
               sb);
```

```
struct block_device * bdev;
struct super_block *  sb;
```

## Arguments

*bdev*

blockdevice to unlock

*sb*

associated superblock

## Description

Unlocks the filesystem and marks it writeable again after `freeze_bdev`.

---

## Name

`bd_claim_by_disk` — wrapper function for `bd_claim_by_kobject`

## Synopsis

```
int bd_claim_by_disk (bdev,
```

```
holder,  
disk);
```

```
struct block_device * bdev;  
void *                holder;  
struct gendisk *      disk;
```

## Arguments

*bdev*

block device to be claimed

*holder*

holder's signature

*disk*

holder's gendisk

## Description

Call `bd_claim_by_kobject` with getting `disk->slave_dir`.

---

## Name

`bd_release_from_disk` — wrapper function for `bd_release_from_kobject`

## Synopsis

```
void bd_release_from_disk (bdev,  
                           disk);
```

```
struct block_device * bdev;  
struct gendisk *      disk;
```

## Arguments

*bdev*

block device to be claimed

*disk*

holder's gendisk

## Description

Call `bd_release_from_kobject` and put `disk->slave_dir`.

---

## Name

`check_disk_size_change` — checks for disk size change and adjusts bdev size.

## Synopsis

```
void check_disk_size_change (disk,  
                             bdev);
```

```
struct gendisk *      disk;  
struct block_device * bdev;
```

## Arguments

*disk*

struct gendisk to check

*bdev*

struct bdev to adjust.

## Description

This routine checks to see if the bdev size does not match the disk size and adjusts it if it differs.

---

## Name

`revalidate_disk` — wrapper for lower-level driver's `revalidate_disk` call-back

## Synopsis

```
int revalidate_disk (disk);
```

```
struct gendisk * disk;
```

## Arguments

*disk*

struct gendisk to be revalidated

# Description

This routine is a wrapper for lower-level driver's `revalidate_disk` call-backs. It is used to do common pre and post operations needed for all `revalidate_disk` operations.

---

## Name

`lookup_bdev` — lookup a struct `block_device` by name

## Synopsis

```
struct block_device * lookup_bdev (pathname);
```

```
const char * pathname;
```

## Arguments

*pathname*

special file representing the block device